



CSCE 670 - Information Storage and Retrieval

Lecture 9: Learning to Rank

Yu Zhang


yuzhang@tamu.edu

September 23, 2025

Course Website: <https://yuzhang-teaching.github.io/CSCE670-F25.html>

Ranking should consider multiple factors

- **YouTube videos**: view, subscribers, video length, user profile factors (e.g., age, location), title relevance, video quality, recency, ...
- **LinkedIn job postings**: posting popularity, company popularity, number of openings, skill match with the user, nearness, recency, salary, ...
- **Our record store**: record popularity, singer popularity, language, keyword match, ...



Hand-tuning a Ranking Function

- $\text{Score}(q, d) =$
 - $a_1 \times \text{TF-IDF}(q, d) +$
 - $a_2 \times \text{BM25}(q, d) +$
 - $a_3 \times \# \text{ views in the last day}(d) +$
 - $a_4 \times \# \text{ views in the last week}(d) +$
 - $a_5 \times \text{recency}(d) +$
 - $a_6 \times \text{PageRank}(d) +$
 - ...
- After checking some examples, you set a_1 as 0.5, a_2 as 0.8, ...
- Problems with this strategy?

Instead, let's learn a good ranker!

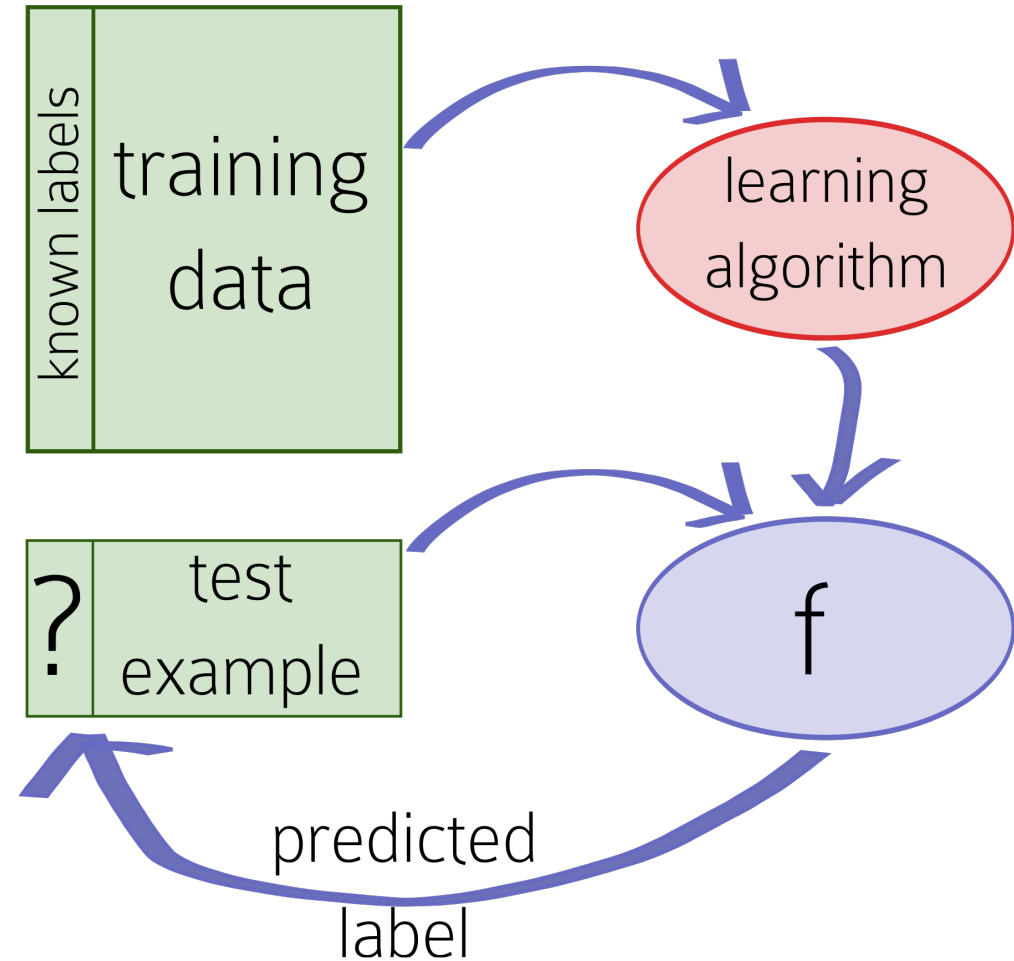
- **Rough Idea** (not 100% accurately framed): Learn the value of a_1, a_2, \dots from data (e.g., relevant query-document pairs according to user clickthrough history)
 - A very natural idea (especially these days)
- But it took a while for ML and IR to be good friends
 - Wong et al., *Linear structure in information retrieval*. SIGIR 1988.
 - Fuhr, *Probabilistic methods in information retrieval*. Computer Journal 1992.
 - Gey, *Inferring probability of relevance using the method of logistic regression*. SIGIR 1994.
 - Herbrich et al., *Large margin rank boundaries for ordinal regression*. Advances in Large Margin Classifiers 2000.

Background: Text Classification

- Given:
 - A document space \mathcal{X}
 - A fixed set of classes $\mathcal{C} = \{c_1, c_2, \dots\}$
 - A training set of labeled documents:
 - E.g., $d_1 \rightarrow c_1, d_2 \rightarrow c_1, d_3 \rightarrow c_2, \dots$
- Use a learning algorithm to learn a classifier f that maps documents to classes $f: \mathcal{X} \rightarrow \mathcal{C}$
- Examples
 - **Paper Topic Classification**: \mathcal{X} = academic papers, $\mathcal{C} = \{\text{math, physics, chemistry, ...}\}$
 - **Review Sentiment Analysis**: \mathcal{X} = food reviews, $\mathcal{C} = \{1\text{-star, 2-star, 3-star, 4-star, 5-star}\}$
 - **Songwriter Prediction**: \mathcal{X} = lyrics, \mathcal{C} = songwriters

Background: Text Classification

- **Training:** Use a learning algorithm to learn a classifier f that maps documents to classes $f: \mathcal{X} \rightarrow \mathcal{C}$
- **Testing/Inference:** Given an unseen document d_{test}
 - Apply our classifier function $f(d_{\text{test}})$ to determine the most appropriate class in \mathcal{C}



A Couple of Simple Text Classifiers: Rocchio

- **Training:** “Learn” class centers for each class by finding the centroid of all the training examples from each class
- **Testing/Inference:** Assign a new example to the class of the nearest class center
- Example:
 - 2-class classification (*chemistry* paper vs. *history* paper)
 - Training samples
 - *chemistry*: $d_1 = (1.0, 0.9)$, $d_2 = (0.9, 1.0)$
 - *history*: $d_3 = (0.2, 0.3)$, $d_4 = (0.3, 0.2)$

A Couple of Simple Text Classifiers: Rocchio


- Example:
 - 2-class classification (*chemistry* paper vs. *history* paper)
 - Training samples
 - *chemistry*: $d_1 = (1.0, 0.9)$, $d_2 = (0.9, 1.0)$
 - *history*: $d_3 = (0.2, 0.3)$, $d_4 = (0.3, 0.2)$
- **Step I**: Compute class centroids
 - *chemistry*: $c_{\text{chemistry}} = \frac{d_1 + d_2}{2} = (0.95, 0.95)$
 - *history*: $c_{\text{history}} = \frac{d_3 + d_4}{2} = (0.25, 0.25)$

A Couple of Simple Text Classifiers: Rocchio

- **Step 1:** Compute class centroids
 - *chemistry*: $c_{\text{chemistry}} = \frac{d_1 + d_2}{2} = (0.95, 0.95)$
 - *history*: $c_{\text{history}} = \frac{d_3 + d_4}{2} = (0.25, 0.25)$
- **Step 2:** Classify a new document
 - New document: $d_5 = (0.8, 0.85)$
 - Compute Euclidean distance:
 - To *chemistry*: $\text{dist}(d_5, c_{\text{chemistry}}) = \sqrt{(0.8 - 0.95)^2 + (0.85 - 0.95)^2} \approx 0.1803$
 - To *history*: $\text{dist}(d_5, c_{\text{history}}) = \sqrt{(0.8 - 0.25)^2 + (0.85 - 0.25)^2} \approx 0.8124$
 - d_5 is closer to *chemistry*, so we classify it as a *chemistry* paper.

A Couple of Simple Text Classifiers: Rocchio

- Can you raise an example where Rocchio does NOT work?




A Couple of Simple Text Classifiers: k Nearest Neighbors (k NN)

- **Training:** None
- **Testing/Inference:** Assign a new example to the majority class of the k -nearest training examples
- **Example:**
 - 2-class classification (*chemistry* paper vs. *history* paper)
 - Training samples
 - *chemistry*: $d_1 = (1.0, 0.9)$, $d_2 = (0.9, 1.0)$, $d_3 = (0.2, 0.3)$, $d_4 = (0.3, 0.2)$
 - *history*: $d_5 = (1.0, 0.3)$, $d_6 = (0.9, 0.2)$, $d_7 = (0.3, 1.0)$, $d_8 = (0.2, 0.9)$
 - New document: $d_9 = (0.4, 1.0)$
 - $k = 3$


A Couple of Simple Text Classifiers: k Nearest Neighbors (k NN)

- Nearest neighbor: *history*
- 2nd nearest neighbor: *history*
- 3rd nearest neighbor: *chemistry*
- **Majority voting:** the new document has more *history* neighbors, so we classify it as a *history* paper.



A Couple of Simple Text Classifiers: k Nearest Neighbors (k NN)

- Can you raise an example where k NN does NOT work?
- How to determine k ? What if $k = 5$?
- Nearest neighbor: *history*
- 2nd nearest neighbor: *history*
- 3rd nearest neighbor: *chemistry*
- 4th nearest neighbor: *chemistry*
- 5th nearest neighbor: *chemistry*




In practice: Which features?

- Very important to select good features to represent our documents
- Features we know about:
 - TF-IDF score of each word (one feature per word)
 - PageRank/Hub/Authority score of the document
 - Popularity, # of clicks, freshness, ...

In practice: Which classifier?

- Many, many ways to learn a good classifier
 - Rocchio
 - k NN
 - Support Vector Machine
 - Naive Bayes
 - Decision Tree
 - Random Forest
 - Gradient-Boosted Decision Tree
 - ...

In practice: How to evaluate?



- Need a way to evaluate how well we do
- Classification accuracy is one way
 - For a held-out test set (for which we know the correct labels), calculate how many labels our classifier correctly predicts
- Many others (some we may talk about later)
- Keep part of the labeled data separate as a validation set
- Train a model over the training data and “test” over the validation set
- Train another model over the training data and “test” over the validation set (and so on and so on)
- Choose model that minimizes error on the validation set


Back to Ranking

- Assume we have a test collection:
 - A benchmark document collection
 - A benchmark suite of queries
 - A binary assessment of either Relevant or Non-relevant for each query and each document
- Sounds like **classification**!
 - **Classification Training**: Given a training set of (*query*, *document* \rightarrow *relevance*) triples, learn a model f that outputs Relevant or Non-relevant
 - **Classification Testing**: Given unseen (*query*, *document*), apply $f(\text{query}, \text{document})$ and output Relevant or Non-relevant
 - **NOTE**: Now our input is not just a *document* but both a *document* and a *query*!

Relevance Classification: Example

				term proximity	
example	docID	query	cosine score	ω	judgment
Φ_1	37	linux operating system	0.032	3	<i>relevant</i>
Φ_2	37	penguin logo	0.02	4	<i>nonrelevant</i>
Φ_3	238	operating system	0.043	2	<i>relevant</i>
Φ_4	238	runtime environment	0.004	2	<i>nonrelevant</i>
Φ_5	1741	kernel layer	0.022	3	<i>relevant</i>
Φ_6	2094	device driver	0.03	2	<i>relevant</i>
Φ_7	3191	device driver	0.027	5	<i>nonrelevant</i>

Relevance Classification: Example



Discriminative Models for Information Retrieval

Ramesh Nallapati
Center for Intelligent Information Retrieval
Department of Computer Science
University of Massachusetts
Amherst, MA 01003
nmramesh@cs.umass.edu

ABSTRACT

Discriminative models have been preferred over generative models in many machine learning problems in the recent past owing to some of their attractive theoretical properties. In this paper, we explore the applicability of discriminative classifiers for IR. We have compared the performance of two popular discriminative models, namely the maximum entropy model and support vector machines with that of language modeling, the state-of-the-art generative model for IR. Our experiments on ad-hoc retrieval indicate that although maximum entropy is significantly worse than language models, support vector machines are on par with language

One of the first theoretically motivated IR models is the binary independence retrieval (BIR) model introduced by Robertson and Sparck Jones [25] in 1976. To the best of our knowledge, this is the first model that viewed IR as a classification problem. They consider retrieval as essentially a process of classifying the entire collection of documents into two classes: relevant and non-relevant. However, instead of doing a hard classification, they estimate the probability of relevance and non-relevance with respect to the query and rank the retrieved documents by their log-likelihood ratio of relevance. Although this was a promising framework, the model did not perform well because of problems in estimation of proba-

Nallapati [SIGIR 2004]

- Experiments:
 - Comparisons with **Lemur (LM)**, a state-of-the-art open-source IR engine
 - **Which classifier?** SVM with linear kernel
 - **What features?** 6 features, all variants of TF, IDF, and TF-IDF scores

	Feature		Feature
1	$\sum_{q_i \in Q \cap D} \log(c(q_i, D))$	4	$\sum_{q_i \in Q \cap D} (\log(\frac{ C }{c(q_i, C)}))$
2	$\sum_{i=1}^n \log(1 + \frac{c(q_i, D)}{ D })$	5	$\sum_{i=1}^n \log(1 + \frac{c(q_i, D)}{ D } idf(q_i))$
3	$\sum_{q_i \in Q \cap D} \log(idf(q_i))$	6	$\sum_{i=1}^n \log(1 + \frac{c(q_i, D)}{ D } \frac{ C }{c(q_i, C)})$

Figure 2: Features in the discriminative models: $c(w, D)$ represents the raw count of word w in document D , C represents the collection, n is the number of terms in the query, $|\cdot|$ is the size-of function and $idf(\cdot)$ is the inverse document frequency.

Experiments on 4 TREC Datasets

- **Metric:** Mean Average Precision (MAP)

Train ↓ Test →		Disks 1-2 (151-200)	Disk 3 (101-150)	Disks 4-5 (401-450)	WT2G (426-450)
Disks 1-2 (101-150)	LM ($\mu^* = 1900$)	0.2561 (6.75e-3)	0.1842	0.2377 (0.80)	0.2665 (0.61)
	SVM	0.2145	0.1877 (0.3)	0.2356	0.2598
	ME	0.1513	0.1240	0.1803	0.1815
Disk 3 (51-100)	LM ($\mu^* = 500$)	0.2605 (1.08e-4)	0.1785 (0.11)	0.2503 (0.21)	0.2666
	SVM	0.2064	0.1728	0.2432	0.2750 (0.55)
	ME	0.1599	0.1221	0.1719	0.1706
Disks 4-5 (301-350)	LM ($\mu^* = 450$)	0.2592 (1.75e-4)	0.1773 (7.9e-3)	0.2516 (0.036)	0.2656
	SVM	0.2078	0.1646	0.2355	0.2675 (0.89)
	ME	0.1413	0.0978	0.1403	0.1355
WT2G (401-425)	LM ($\mu^* = 2400$)	0.2524 (4.6e-3)	0.1838 (0.08)	0.2335	0.2639
	SVM	0.2199	0.1744	0.2487 (0.046)	0.2798 (0.037)
	ME	0.1353	0.0969	0.1441	0.1432
Best TREC runs (Site)		0.4226 (UMass)	N/A	0.3207 (Queen's College)	N/A

Experiments on 4 TREC Datasets

- At best the results are about equal to Lemur
 - Actually a little bit below
- Paper's advertisement: Easy to add more features
- This is illustrated on a homepage finding task on WT10G:

	Success@10
Lemur	0.52
SVM with text features only	0.58
SVM with URL-depth and in-link features	0.78

Questions?

But Boolean \neq Ranking!


- Assigning (*query, document*) to
 - Relevant
 - or
 - Not Relevant
- Is not really what we want when we think about *ranking*

Pointwise Learning

- Assume we have training data like (*query*, *document*, *score*)
- Here the *score* could be a relevance score like
 - 4 Perfect match
 - 3 Very relevant
 - 2 Relevant
 - 1 Somewhat relevant
 - 0 Not relevant at all
- Our goal is to output a score
 - This is *regression* (if we can output any value)
 - Or *ordinal regression* (If we can only output 0, 1, 2, 3, or 4)

Pointwise Learning

- What could be a difference between **classification** and **ordinal regression**?



Pointwise Learning

- **Regression Training:** Given a training set of $(query, document \rightarrow score)$ triples, learn a model f
- **Regression Testing:** Given unseen $(query, document)$, apply $f(query, document)$ and output the $score$
- Challenges?
 - Expensive to collect labels
 - Focuses on scores, not relative ordering (or relationship to other documents)
 - Bias towards frequent queries
 - ...

Pairwise Learning

- Aim is to classify instance pairs as correctly ranked or incorrectly ranked
 - Given the query q and two candidates (c_i, c_k) , predict if c_i should be ranked higher than c_k (denoted as $c_i > c_k$)
- This turns an ordinal regression problem back into a binary classification problem in an expanded space
 - We only need lots of (c_i, c_k) , where we already know $c_i > c_k$, for training
- Formally, we want a ranking function f such that
 - $c_i > c_k \iff f(\psi_i) > f(\psi_k)$
 - ψ_i is the feature vector of c_i given the query q (e.g., one entry can be $\text{TF-IDF}(q, c_i)$)
- To simplify our discussion, let's suppose that f is a linear function: $f(\psi_i) = w^T \psi_i$

Joachims [KDD 2002]

Optimizing Search Engines using Clickthrough Data

Thorsten Joachims
Cornell University
Department of Computer Science
Ithaca, NY 14853 USA
tj@cs.cornell.edu

ABSTRACT

This paper presents an approach to automatically optimizing the retrieval quality of search engines using clickthrough data. Intuitively, a good information retrieval system should present relevant documents high in the ranking, with less relevant documents following below. While previous approaches to learning retrieval functions from examples exist, they typically require training data generated from relevance judgments by experts. This makes them difficult and ex-

ceive millions of queries per day, such data is available in abundance. Compared to explicit feedback data, which is typically elicited in laborious user studies, any information that can be extracted from logfiles is virtually free and substantially more timely.

This paper presents an approach to learning retrieval functions by analyzing which links the users click on in the presented ranking. This leads to a problem of learning with preference examples like "for query q , document d_a should


Training a Linear SVM for Ranking

Ranking function:

$$f(\psi_i) = w^T \psi_i$$

Should be
ranked lower


Should be
ranked higher




Training a Linear SVM for Ranking

- But we don't have pointwise training data!
 - Remember we only have lots of (c_i, c_k) , where we already know $f(\psi_i) > f(\psi_k)$
 - We don't know the value of $f(\psi_i)$ or $f(\psi_k)$
- **Idea:** Create a new instance space from pairwise learning
 - We have $c_i > c_k \iff f(\psi_i) > f(\psi_k)$
 - We also have $f(\psi_i) = w^T \psi_i$ and $f(\psi_k) = w^T \psi_k$
 - So $c_i > c_k \iff w^T \psi_i > w^T \psi_k \iff w^T (\psi_i - \psi_k) > 0$
 - Let's create a new instance $\phi_u = \psi_i - \psi_k$
 - And $z_u = +1, 0, -1$ as $c_i >, =, < c_k$
 - From training data $\mathcal{S} = \{\phi_u\}$, we train an SVM

Two Queries in the Original Space



Two Queries in the Pairwise Space



Performance of Ranking SVM

Comparison	more clicks on learned	less clicks on learned	tie (with clicks)	no clicks	total
Learned vs. Google	29	13	27	19	88
Learned vs. MSNSearch	18	4	7	11	40
Learned vs. Toprank	21	9	11	11	52

Table 2: Pairwise comparison of the learned retrieval function with Google, MSNSearch, and the non-learning meta-search ranking. The counts indicate for how many queries a user clicked on more links from the top of the ranking returned by the respective retrieval function.

weight	feature		
		0.16	top1_hotbot
0.60	query_abstract_cosine	...	
0.48	top10_google	0.14	domain_name_in_query
0.24	query_url_cosine	...	
0.24	top1count_1	-0.13	domain_tu-bs
0.24	top10_msnsearch	-0.15	country_fi
0.22	host_citeseer	-0.16	top50count_4
0.21	domain_nec	-0.17	url_length
0.19	top10count_3	-0.32	top10count_0
0.17	top1_google	-0.38	top1count_0
0.17	country_de		
...			
0.16	abstract_contains_home		

Table 3: Features with largest and smallest weights as learned from the training data in the online experiment.

Questions?

Burges et al. [ICML 2025] (RankNet)

Learning to Rank using Gradient Descent

Keywords: ranking, gradient descent, neural networks, probabilistic cost functions, internet search

Chris Burges

CBURGES@MICROSOFT.COM

Tal Shaked*

TAL.SHAKED@GMAIL.COM

Erin Renshaw

ERINREN@MICROSOFT.COM

Microsoft Research, One Microsoft Way, Redmond, WA 98052-6399

Ari Lazier

ARIEL@MICROSOFT.COM

Matt Deeds

MADEEDS@MICROSOFT.COM

Nicole Hamilton

NICHAM@MICROSOFT.COM

Greg Hullender

GREGHULL@MICROSOFT.COM

Microsoft, One Microsoft Way, Redmond, WA 98052-6399

Abstract

We investigate using gradient descent methods for learning ranking functions; we pro-

that maps to the reals (having the model evaluate on pairs would be prohibitively slow for many applications). However (Herbrich et al., 2000) cast the rank-

Burges et al. [ICML 2015] (RankNet)

- Led to popular and successful variants:
 - LambdaRank
 - LambdaMART: top performer at [the 2010 Yahoo Learning to Rank Challenge](#)

JMLR: Workshop and Conference Proceedings 14 (2011) 1–24

Yahoo! Learning to Rank Challenge

Yahoo! Learning to Rank Challenge Overview

Olivier Chapelle*

Yi Chang

Yahoo! Labs

Sunnyvale, CA

CHAP@YAHOO-INC.COM

YICHANG@YAHOO-INC.COM

Existing Public Datasets

Table 1: Characteristics of publicly available datasets for learning to rank: number of queries, documents, relevance levels, features and year of release. The size of the 6 datasets for the '.gov' collection in LETOR have been added together. Even though this collection has a fairly large number of documents, only 2000 of them are relevant.

	Queries	Doc.	Rel.	Feat.	Year
LETOR 3.0 – Gov	575	568 k	2	64	2008
LETOR 3.0 – Ohsumed	106	16 k	3	45	2008
LETOR 4.0	2,476	85 k	3	46	2009
Yandex	20,267	213 k	5	245	2009
Yahoo!	36,251	883 k	5	700	2010
Microsoft	31,531	3,771 k	5	136	2010

Datasets for the Challenge

Table 2: Statistics of the two datasets released for the challenge.

	SET 1			SET 2		
	Train	Valid.	Test	Train	Valid.	Test
Queries	19,944	2,994	6,983	1,266	1,266	3,798
Documents	473,134	71,083	165,660	34,815	34,881	103,174
Features		519			596	

Table 3: Distribution of relevance labels.

Grade	Label	SET 1	SET 2
Perfect	4	1.67%	1.89%
Excellent	3	3.88%	7.67%
Good	2	22.30%	28.55%
Fair	1	50.22%	35.80%
Bad	0	21.92%	26.09%

Features

- **Web graph**: in-links, out-links, PageRank, ...
- **Doc statistics**: # of words in title, # of words in body, number of slashes in URL, ...
- **Doc classifier**: spam, topic, language, ...
- **Query**: # of terms, frequency of query and its terms, ...
- **Text match**: BM25, counts, ...
- **Clicks**: probability of a click, dwell time, ...
- **External references**: tags
- **Time**: age of doc, age of in-links, ...

Baselines

Table 5: Performance of the 3 baselines methods on the validation and test sets of SET 1: BM25F-SD is a text match feature, RankSVM is linear pairwise learning to rank method and GBDT is a non-linear regression technique.

	Validation		Test	
	ERR	NDCG	ERR	NDCG
BM25F-SD	0.42598	0.73231	0.42853	0.73214
RankSVM	0.43109	0.75156	0.43680	0.75924
GBDT	0.45625	0.78608	0.46201	0.79013

The Winners

Track 1

RankNet

1	C. Burges, K. Svore, O. Dekel, Q. Wu, P. Bennett, A. Pastusiak and J. Platt (Microsoft Research)	0.46861
2	E. Gottschalk (Activision Blizzard) and D. Vogel (Data Mining Solutions)	0.46786
3	M. Parakhin (Microsoft) – <i>Prize declined</i>	0.46695
4	D. Pavlov and C. Brunk (Yandex Labs)	0.46678
5	D. Sorokina (Yandex Labs)	0.46616

- What is RankNet? Next lecture!



Thank You!

Course Website: <https://yuzhang-teaching.github.io/CSCE670-F25.html>