



CSCE 670 - Information Storage and Retrieval

Lecture 18: BERT-Based Ranking, ColBERT

Yu Zhang

yuzhang@tamu.edu

October 28, 2025

Course Website: <https://yuzhang-teaching.github.io/CSCE670-F25.html>

Unfinished Part from Last Lecture

Recap: Transformer as a Black Box

- Transformer is a **neural network**
- It has two types of architecture: **encoder** and **decoder**
- BERT uses the **Transformer encoder** architecture
- Input to a **Transformer encoder** can be a piece of text:
 - A sequence of words w_1, w_2, \dots, w_L
 - Represented by their corresponding embeddings $e_{w_1}, e_{w_2}, \dots, e_{w_L}$
- Then, the output is a sequence of contextualized word vectors $h_{w_1}, h_{w_2}, \dots, h_{w_L}$
 - The output vector h_{w_i} captures the meaning of w_i by considering the entire input sequence as w_i 's context
 - $h_{w_i} = \text{Transformer}(e_{w_i} | e_{w_1}, e_{w_2}, \dots, e_{w_L})$

BERT [Devlin et al., NAACL 2019]

BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Jacob Devlin Ming-Wei Chang Kenton Lee Kristina Toutanova

Google AI Language

{jacobdevlin, mingweichang, kentonl, kristout}@google.com

Abstract

We introduce a new language representation model called **BERT**, which stands for Bidirectional Encoder Representations from Transformers. Unlike recent language representation models (Peters et al., 2018a; Radford et al., 2018), BERT is designed to pre-train deep bidirectional representations from unlabeled text by jointly conditioning on both left and right context in all layers. As a re-

There are two existing pre-trained language representation tasks: *feature-based* and *sequence-based*. The feature-based approach, proposed by (Peters et al., 2018a), uses task-specific features to predict the output of a sequence. The sequence-based approach, proposed by (Lipton et al., 2018), uses task-specific features to predict the output of a sequence. The fine-tuning approach, such as the Generative Pre-trained Transformer (OpenAI GPT) (Radford et al., 2018), introduces minimal changes to the pre-trained model to make it suitable for a specific task.

Bert: Pre-training of deep bidirectional transformers for language understanding

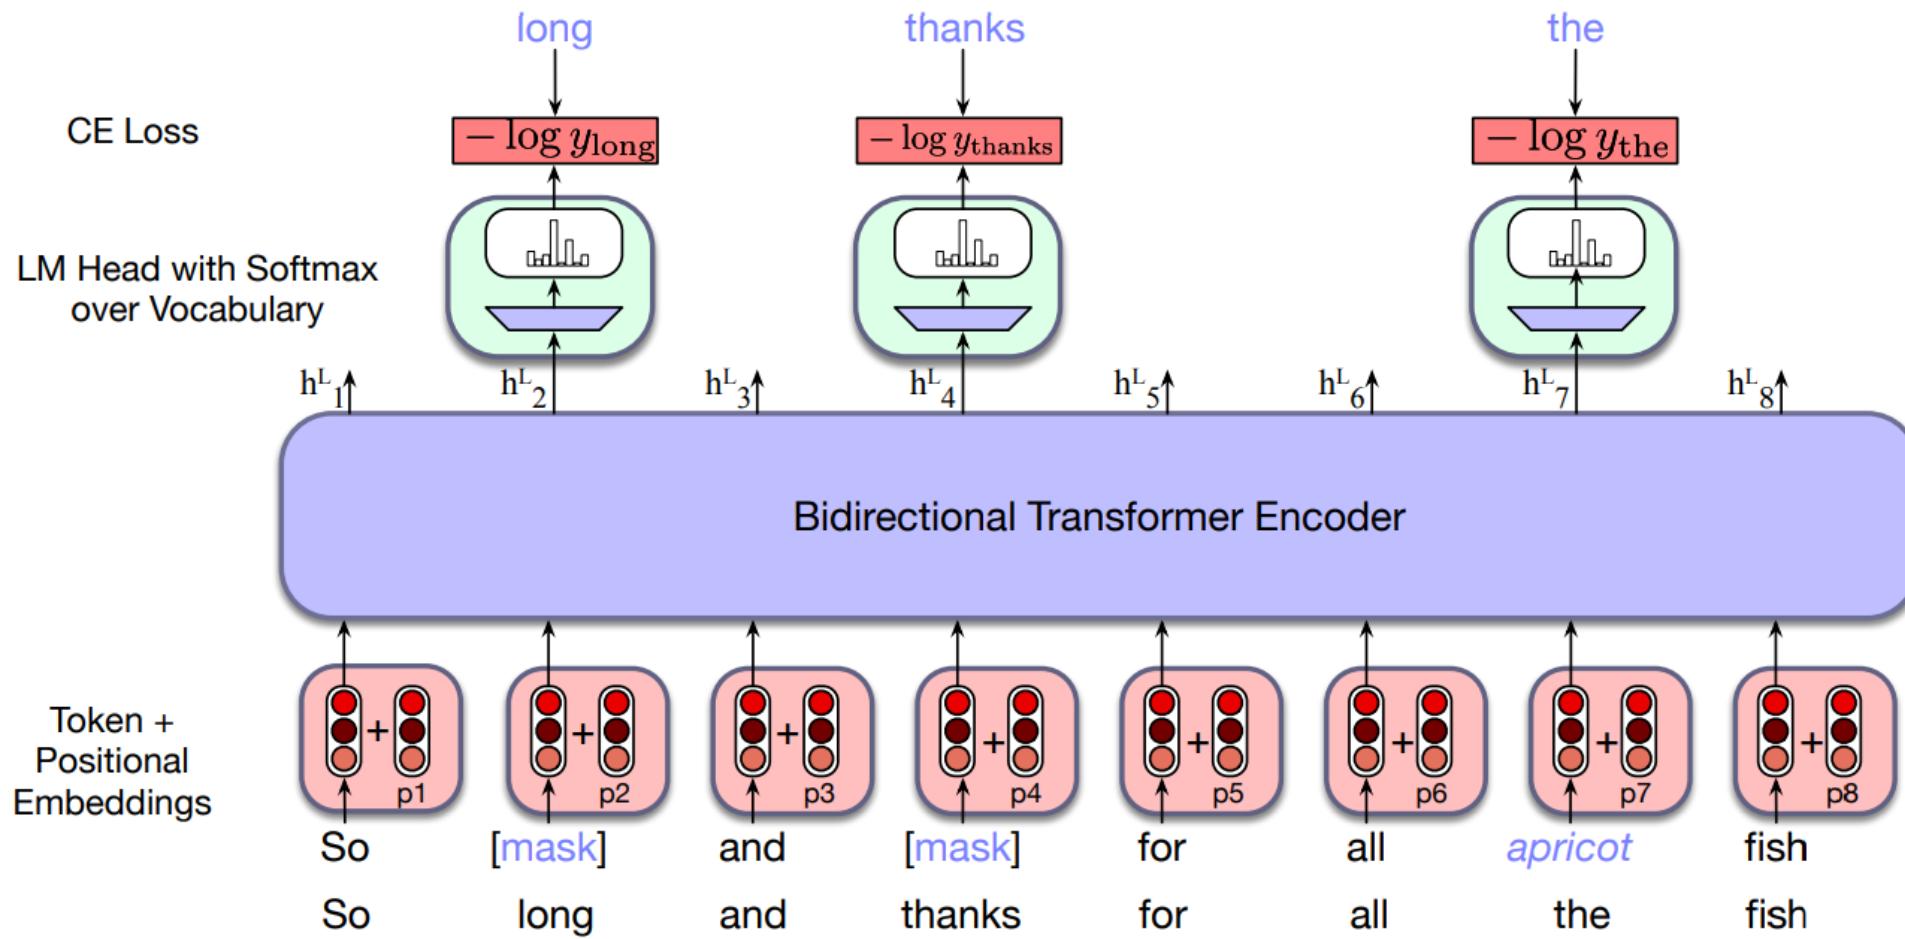
J Devlin, MW Chang, K Lee... - Proceedings of the 2019 ... , 2019 - aclanthology.org

... deep bidirectionality of **BERT** by evaluating two **pretraining** objectives using exactly the same **pretraining** objective. No NSP: A **bidirectional** model is trained using the “masked LM” (MLM) ...

☆ Cited by 146893 Related articles ➔

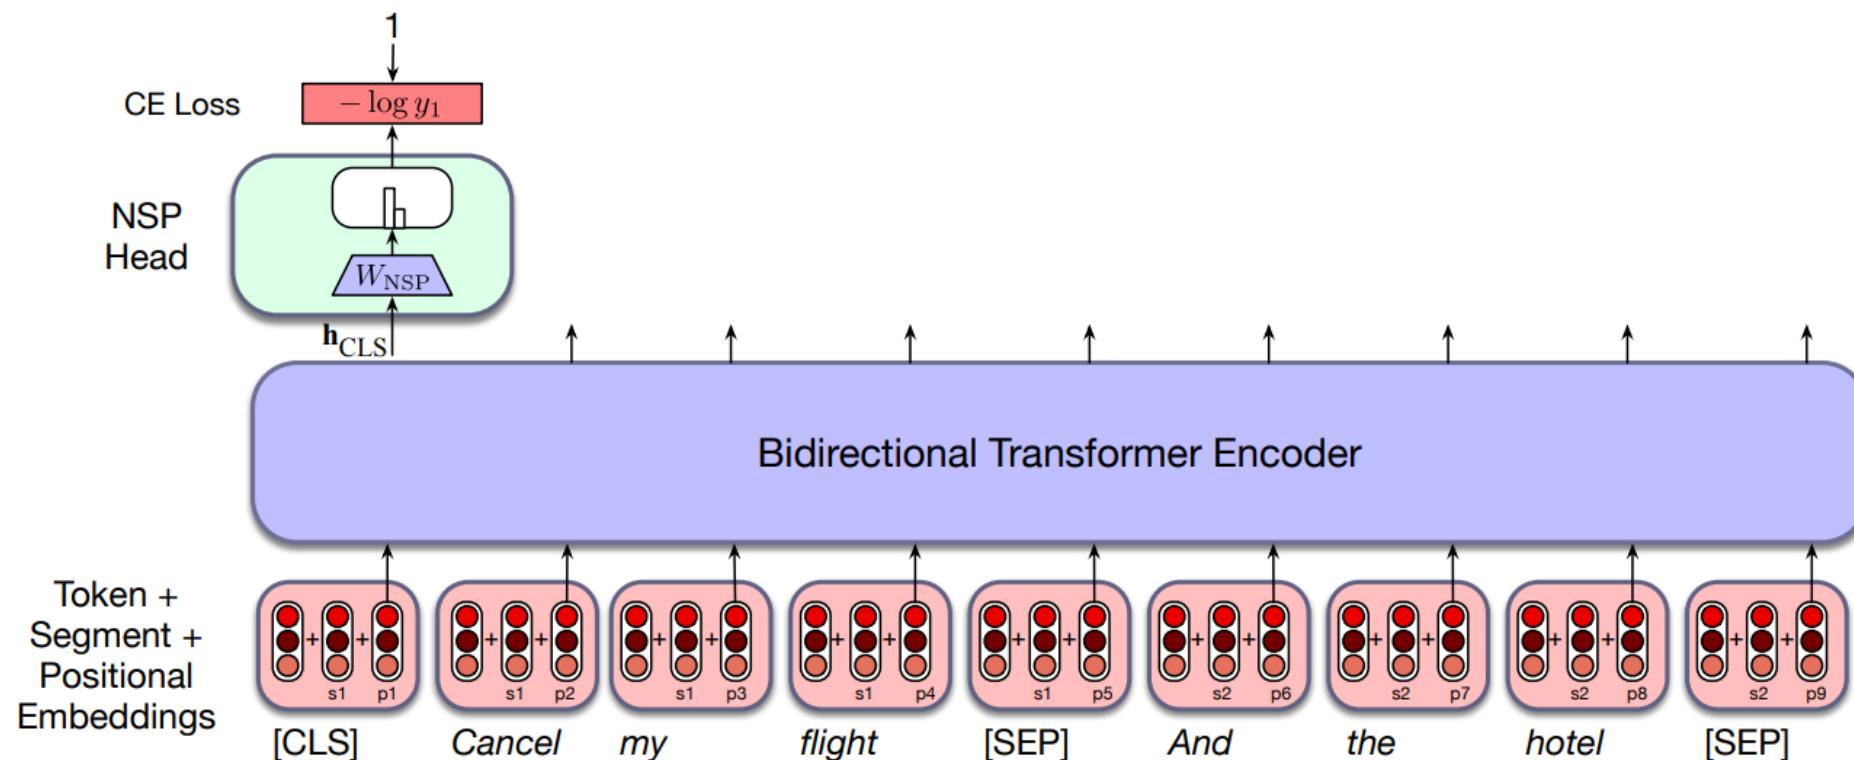
BERT Pre-training

- Task 1 – Masked Language Modeling (MLM): With 15% words randomly masked, the model learns bidirectional contextual information to predict the masked words.



BERT Pre-training

- **Task 2 – Next Sentence Prediction (NSP):** The model is presented with pairs of sentences. It is trained to predict whether each pair consists of an actual pair of adjacent sentences from the training corpus or a pair of unrelated sentence.



Immediate Impact of BERT

- In 2018, BERT came out and largely outperformed most previous methods on common NLP tasks (e.g., sentiment classification, natural language inference, question answering).
- BERT got the best paper award at the NAACL 2019 conference.

System	MNLI-(m/mm)	QQP	QNLI	SST-2	CoLA	STS-B	MRPC	RTE	Average
	392k	363k	108k	67k	8.5k	5.7k	3.5k	2.5k	-
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
BERT _{BASE}	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
BERT _{LARGE}	86.7/85.9	72.1	92.7	94.9	60.5	86.5	89.3	70.1	82.1

- “Open AI GPT”: GPT-2
- “BERT-Base”: 12 Transformer encoder layers; ~110M parameters
- “BERT-Large”: 24 Transformer encoder layers; ~340M parameters

Improving BERT: RoBERTa [Liu et al., arXiv 2019]

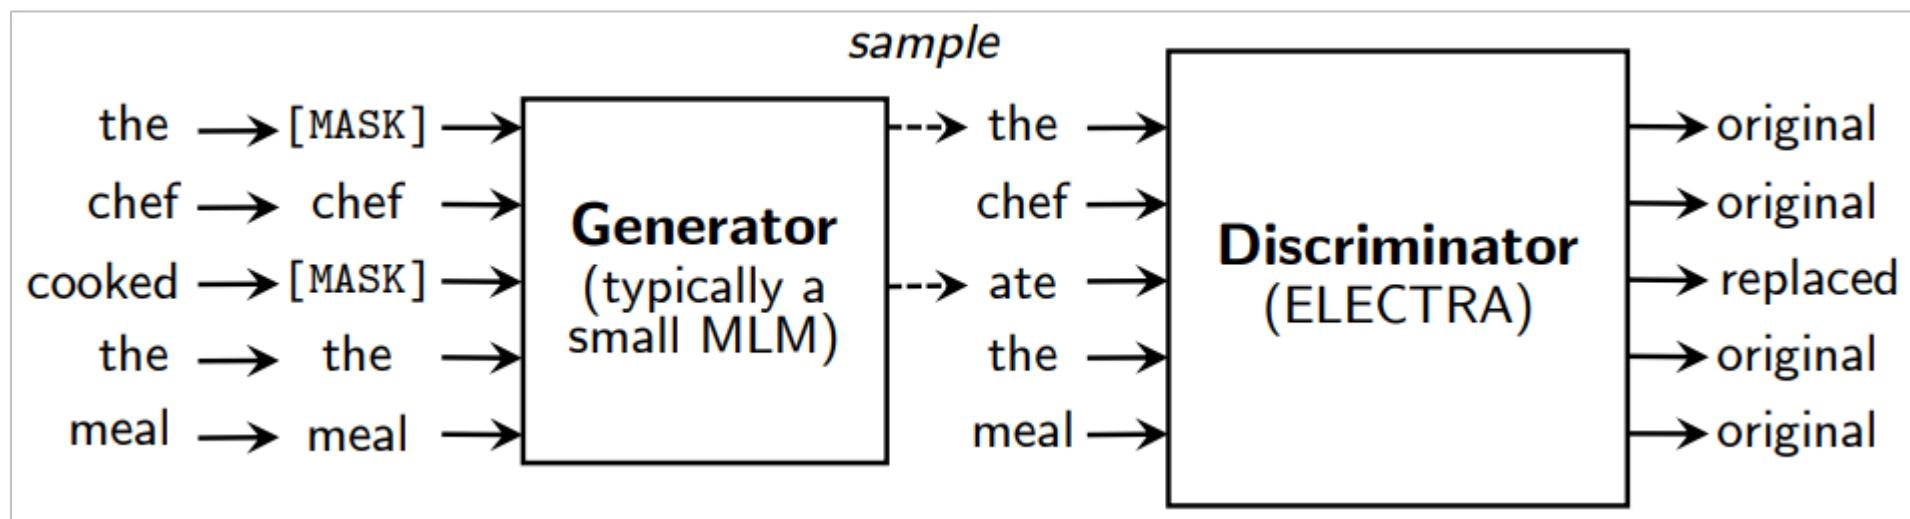
- Next Sentence Prediction (NSP) is not helpful!
- Only use Masked Language Modeling (MLM)
- Pretrain on longer sequences
- Pretrain the model for longer, with bigger batches
- Pretrain over more data
- Dynamically change the masking patterns applied to the training data in each epoch

Model	SQuAD 1.1/2.0	MNLI-m	SST-2	RACE
<i>Our reimplementation (with NSP loss):</i>				
SEGMENT-PAIR	90.4/78.7	84.0	92.9	64.2
SENTENCE-PAIR	88.7/76.2	82.9	92.1	63.0
<i>Our reimplementation (without NSP loss):</i>				
FULL-SENTENCES	90.4/79.1	84.7	92.5	64.8
DOC-SENTENCES	90.6/79.7	84.7	92.7	65.6

Model	data	bsz	steps	SQuAD (v1.1/2.0)	MNLI-m	SST-2
RoBERTa						
with BOOKS + WIKI	16GB	8K	100K	93.6/87.3	89.0	95.3
+ additional data (§3.2)	160GB	8K	100K	94.0/87.7	89.3	95.6
+ pretrain longer	160GB	8K	300K	94.4/88.7	90.0	96.1
+ pretrain even longer	160GB	8K	500K	94.6/89.4	90.2	96.4

Improving BERT: ELECTRA [Clark et al., ICLR 2020]

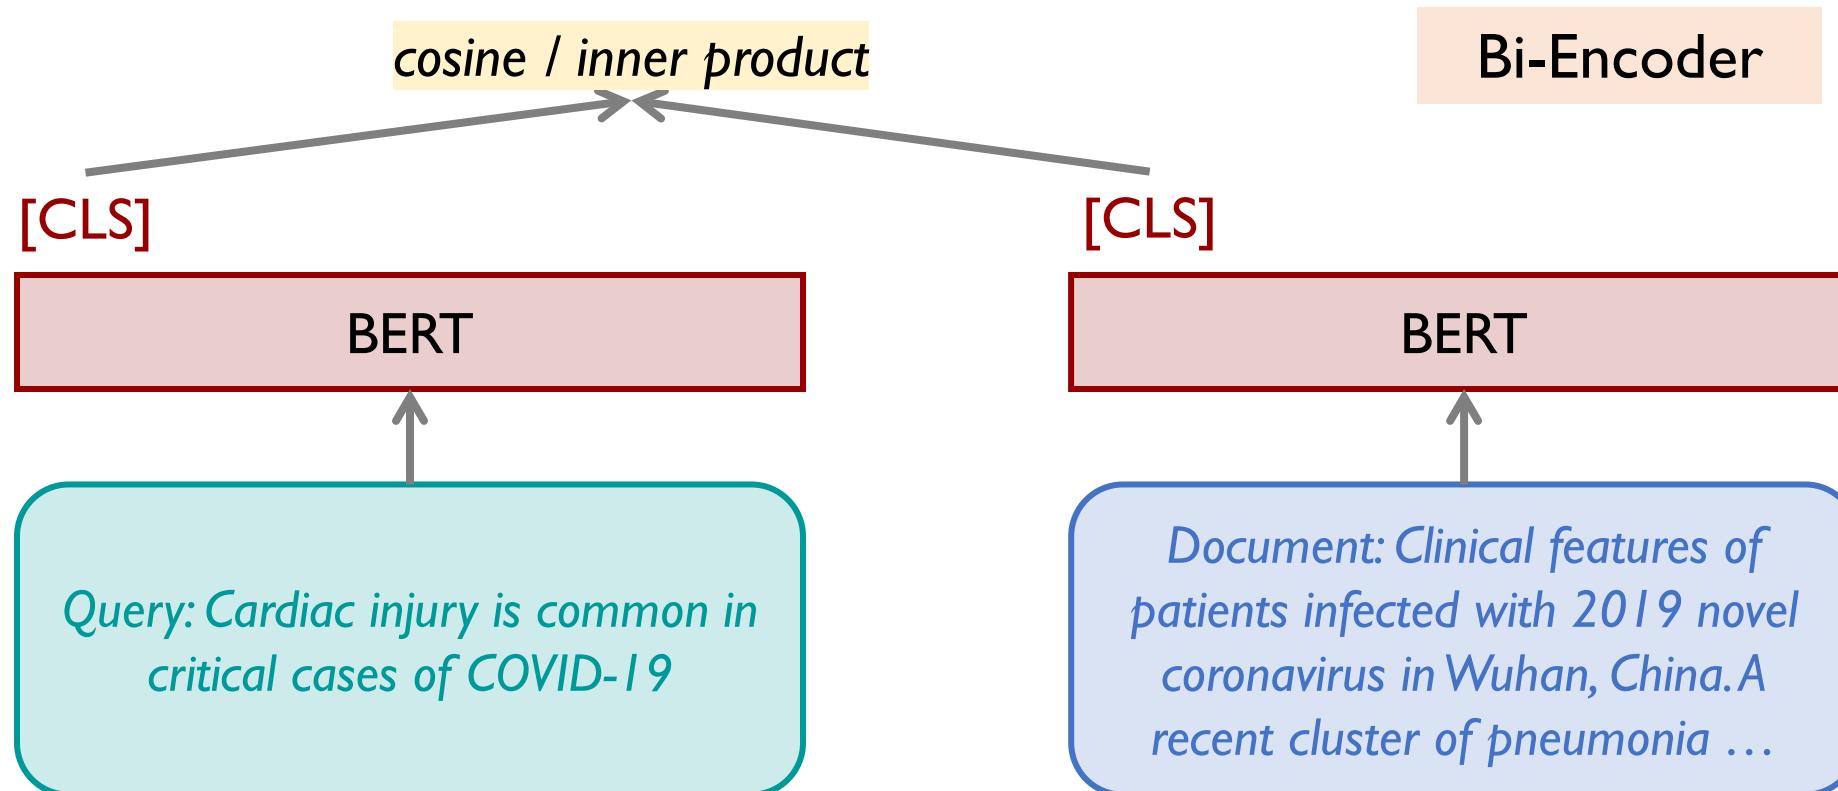
- Use a small MLM as an auxiliary generator (discarded after pretraining)
- Pretrain the main model as a discriminator
- The small auxiliary MLM and the main discriminator are jointly trained.
- The main model's pretraining task becomes **more and more challenging** in pretraining.



Content of This Lecture

How to use BERT for retrieval? – Solution 1

- Encode **query** and **document** separately
- The output vector of the [CLS] token serves as **query** / **document** embedding



Python Implementation to Encode a Query / Document

```
python

from transformers import BertTokenizer, BertModel
import torch

# Load pre-trained BERT-base model and tokenizer
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
model = BertModel.from_pretrained('bert-base-uncased')

# Input text
text = "Cardiac injury is common in critical cases of COVID-19"

# Tokenize and encode input
inputs = tokenizer(text, return_tensors='pt', truncation=True, padding=True)

# Forward pass (no gradient needed)
with torch.no_grad():
    outputs = model(**inputs)

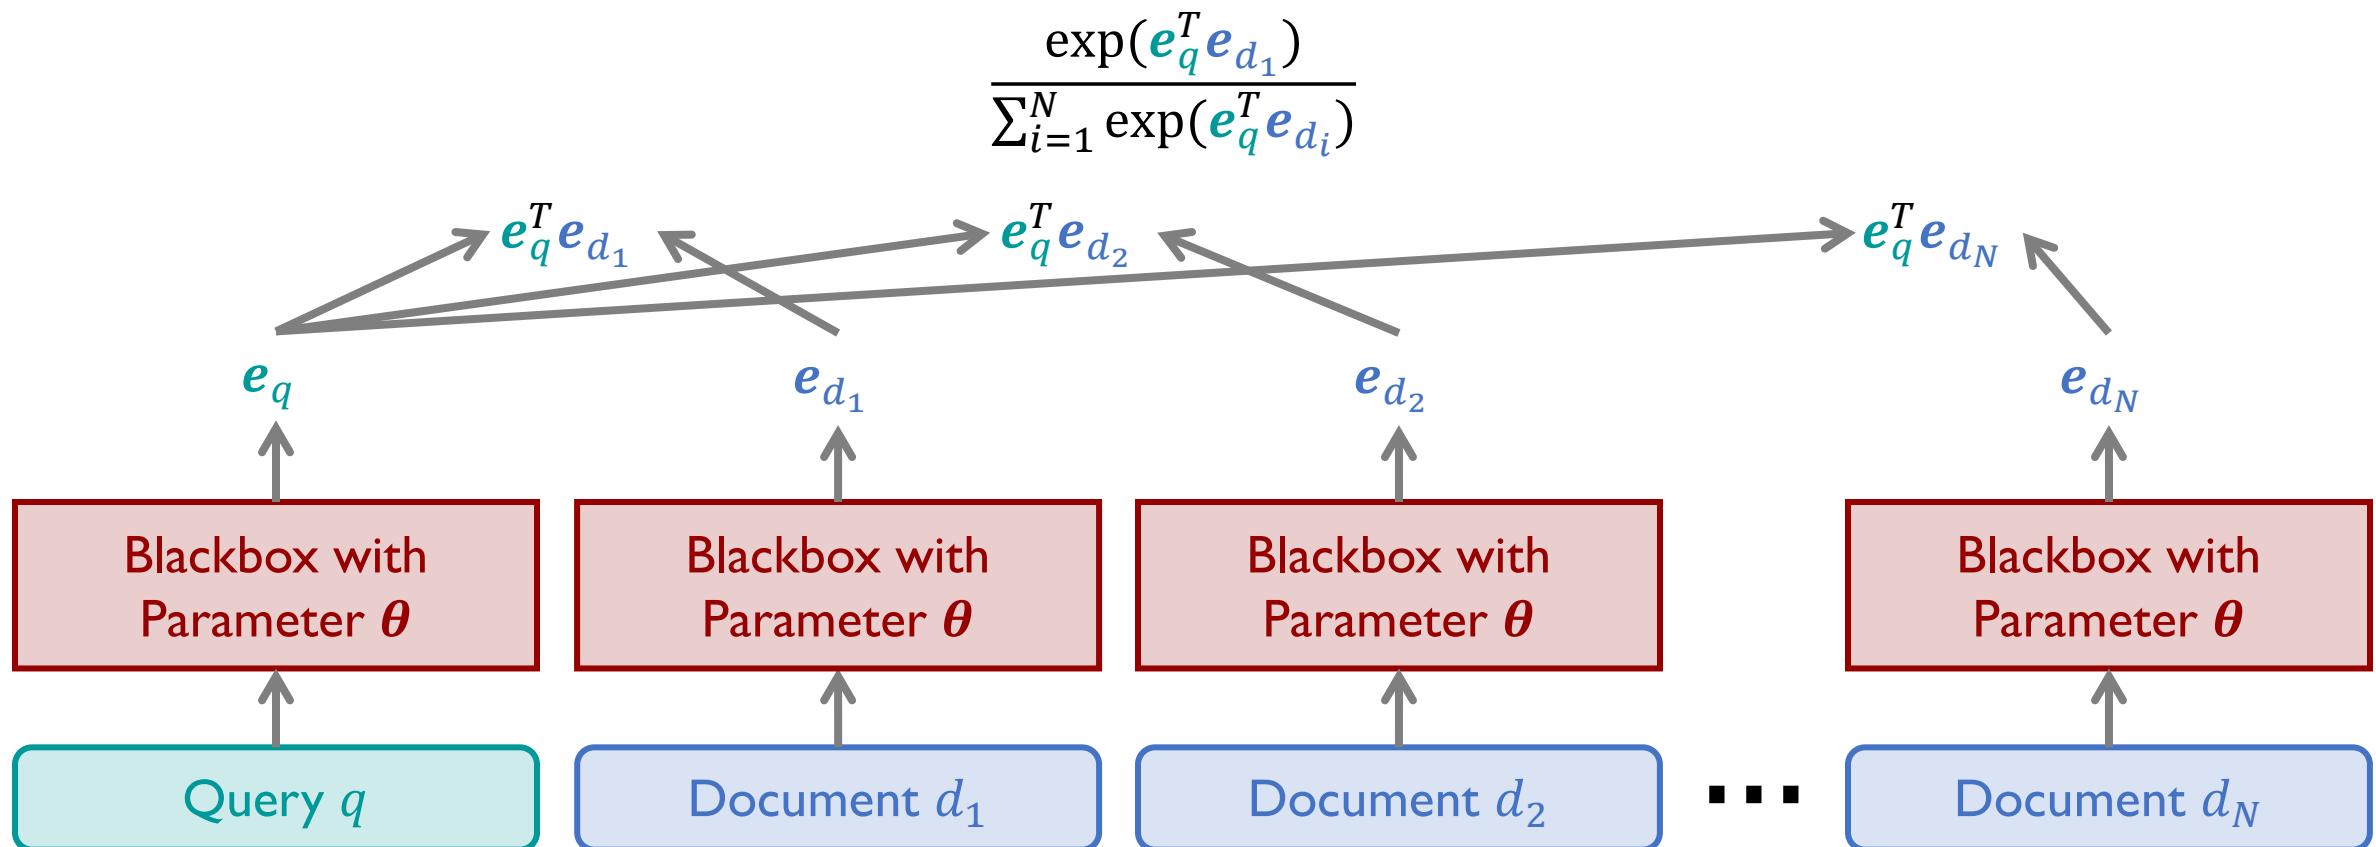
# Extract [CLS] token embedding
cls_embedding = outputs.last_hidden_state[:, 0, :] # shape: [1, 768]
```

Fine-Tune a Bi-Encoder Ranking Model

- **Fine-tuning:** continue to train a pre-trained model using supervised learning on labeled data for a specific downstream task
- What if I want to train a Bi-Encoder using learning to rank?
 - **Parameters:** All parameters in the **query** Transformer encoder and the **document** Transformer encoder
 - At the beginning of training, both encoders are initialized with **BERT**
 - In many Bi-Encoder ranking models, the **query** encoder and the **document** encoder share all parameters
 - **Learning Objective?**
 - Recall Dense Passage Retrieval!

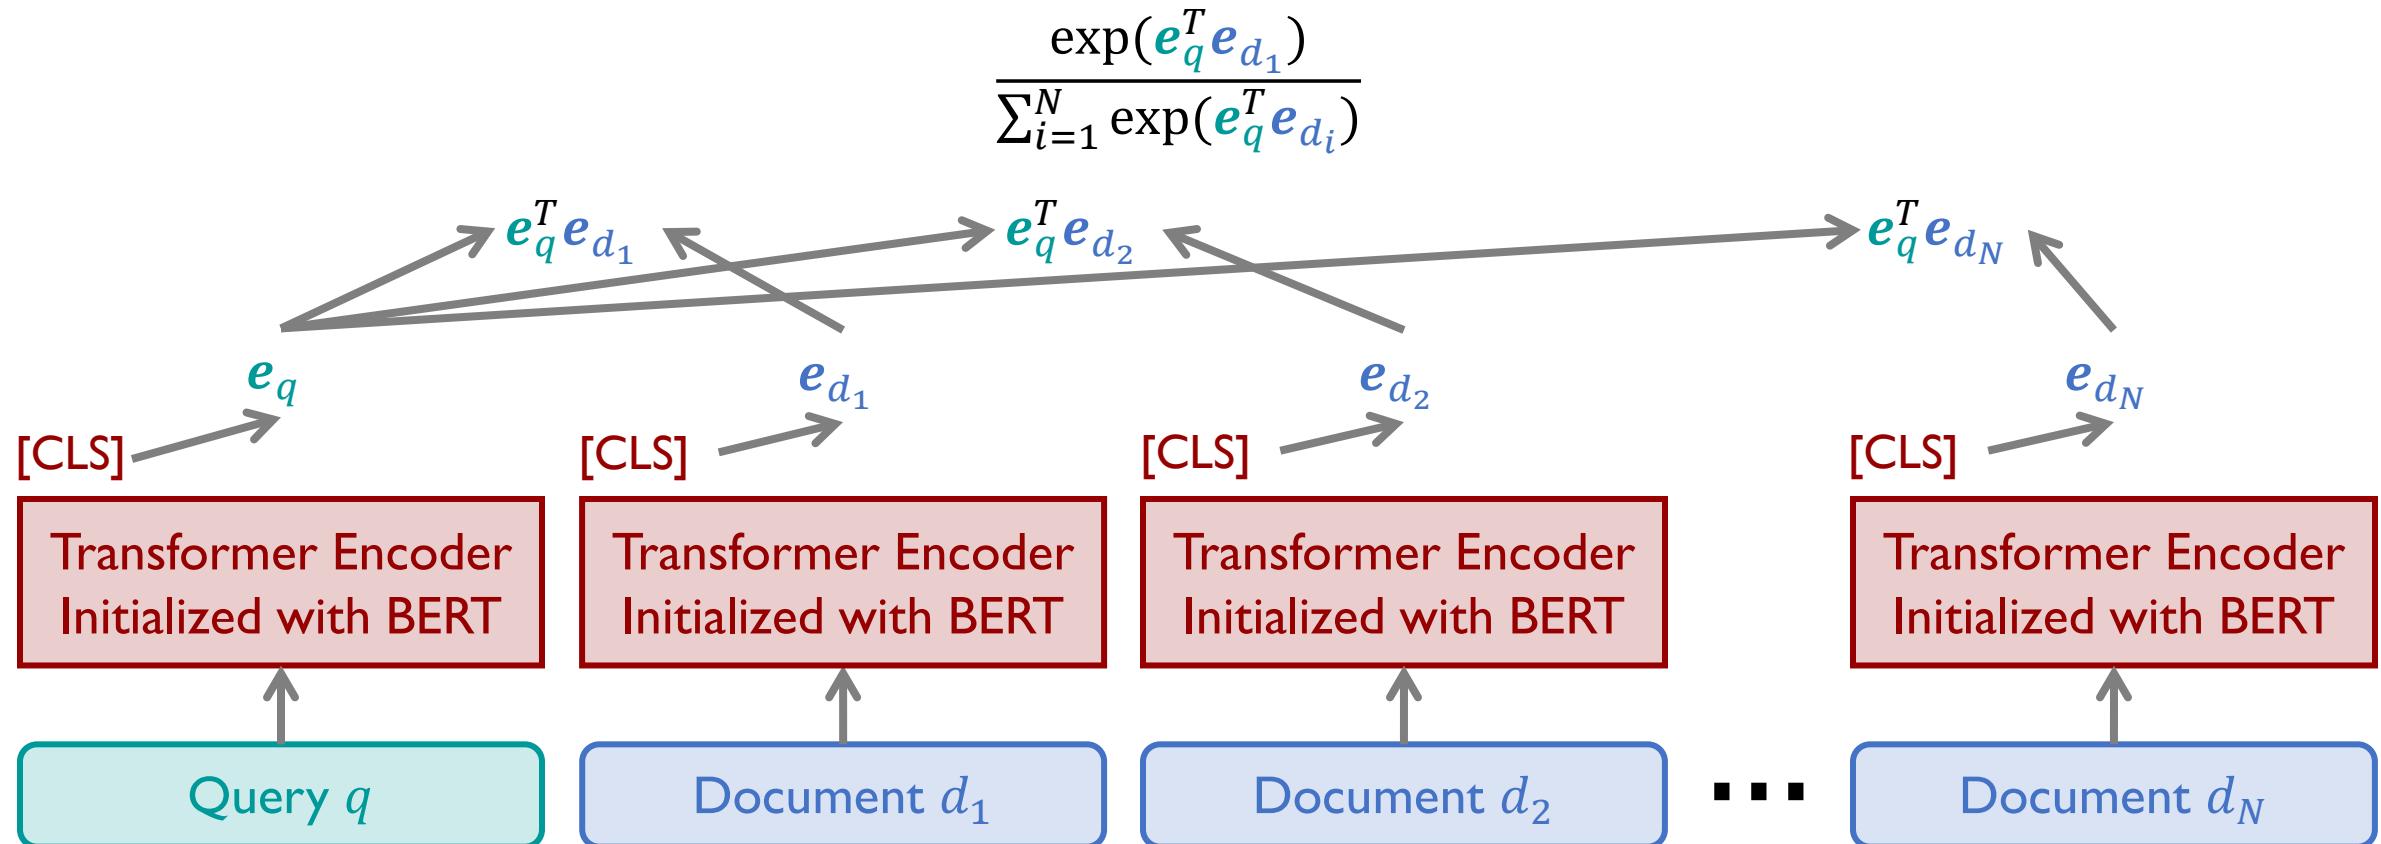
Dense Passage Retrieval [Karpukhin et al., EMNLP 2020]

- **Learning objective:** Given a query q and N documents (d_1, d_2, \dots, d_N), the ground truth tells us that d_1 is the most relevant to q among these documents



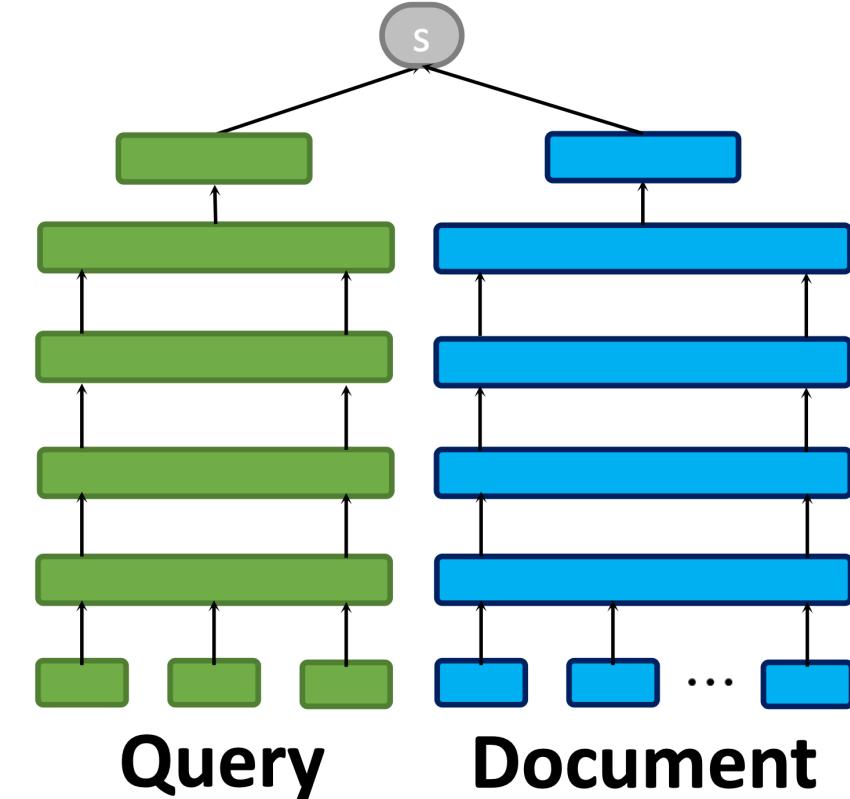
Dense Passage Retrieval (BERT Version)

- **Learning objective:** Given a query q and N documents (d_1, d_2, \dots, d_N), the ground truth tells us that d_1 is the most relevant to q among these documents



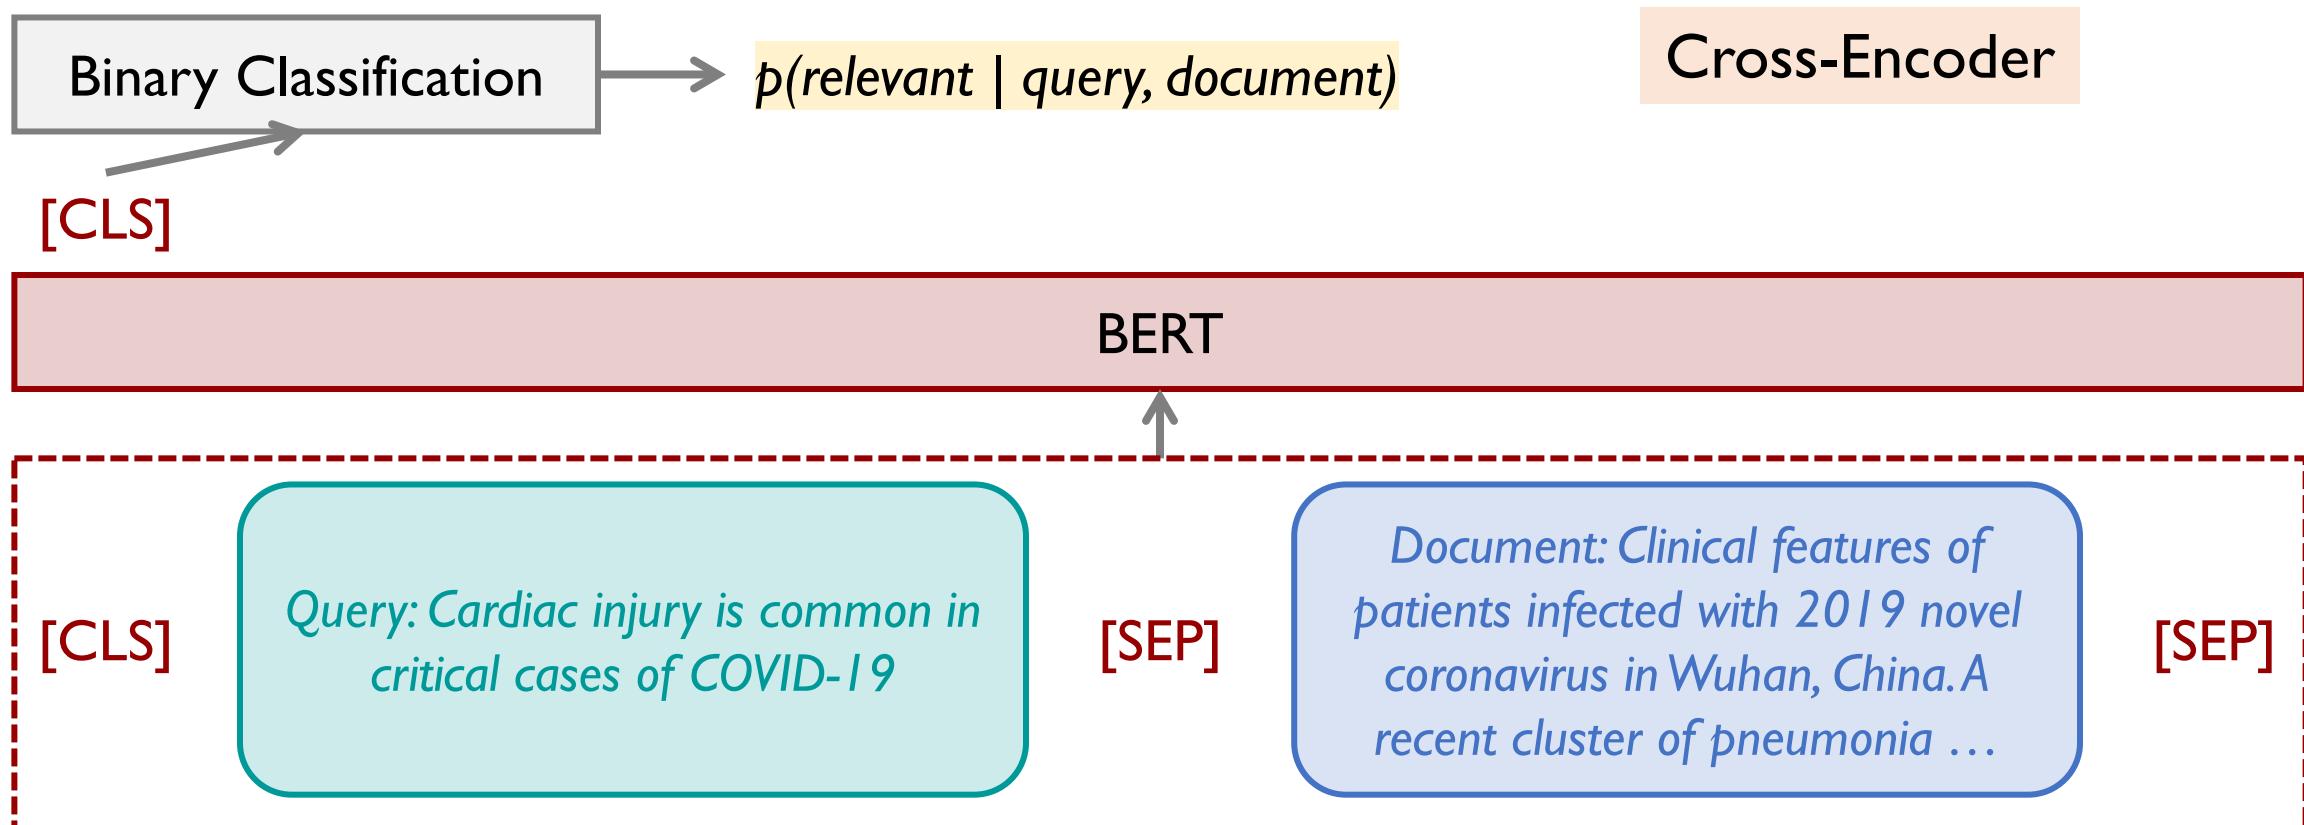
Fine-Tune a Bi-Encoder Ranking Model

- In-batch negative sampling
 - The larger the batch size, the better
- What potential issues might a Bi-Encoder have?
 - Two encoders independently encode the **query** and the **document** into vectors and calculate their similarity.
 - However, the importance of a **query word** may vary across different **documents**; the importance of a **document word** may also vary across different **queries**.



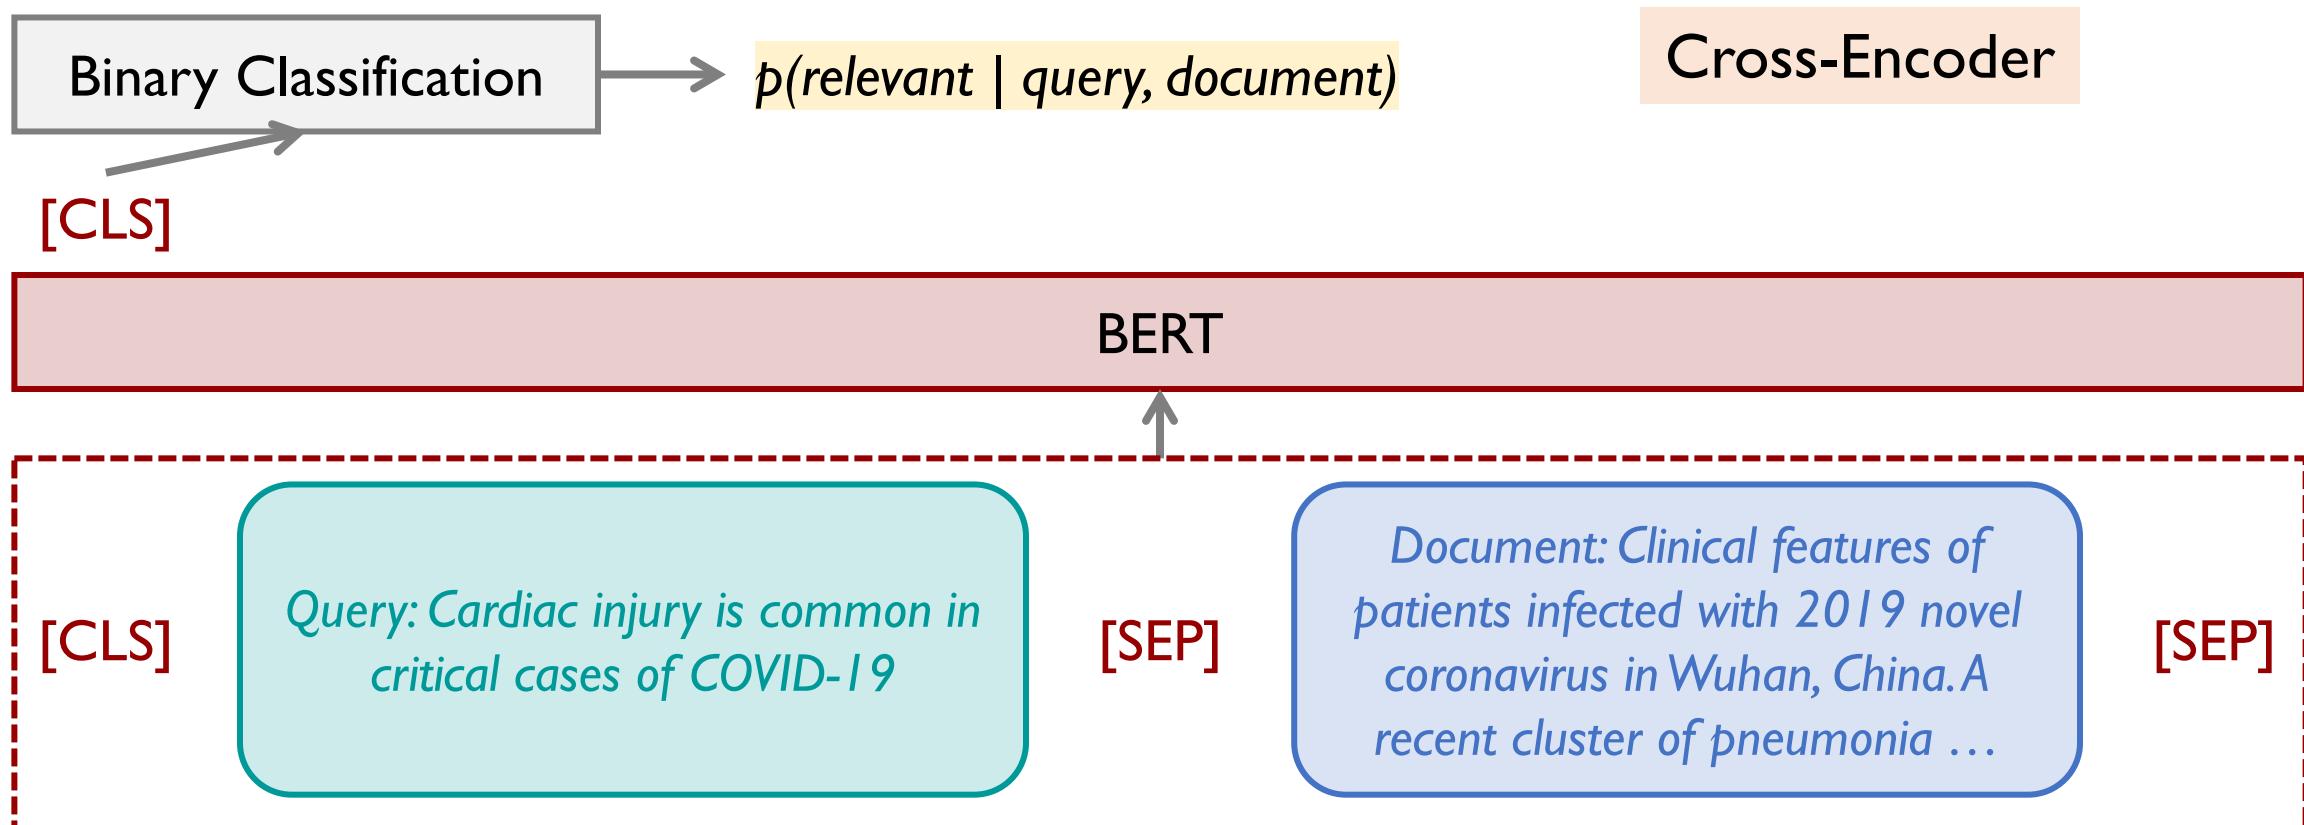
How to use BERT for retrieval? – Solution 2

- Concatenate the **query** and **document** into a single input sequence
- Get the representation of the entire sequence and perform binary classification



Fine-Tune a Cross-Encoder Ranking Model

- **Parameters:** All parameters in the Transformer encoder + the classification layer
- **Objective:** Cross-Entropy Loss, $-(y \log p + (1 - y) \log(1 - p))$



Cross-Encoder Ranking Model [Dai and Callan, SIGIR 2019]

Deeper Text Understanding for IR with Contextual Neural Language Modeling

Zhuyun Dai

Carnegie Mellon University

zhuyund@cs.cmu.edu

Jamie Callan

Carnegie Mellon University

callan@cs.cmu.edu

ABSTRACT

Neural networks provide new possibilities to automatically learn complex language patterns and query-document relations. Neural IR models have achieved promising results in learning query-document relevance patterns, but few explorations have been done on understanding the text content of a query or a document. This paper studies leveraging a recently-proposed contextual neural language model, BERT, to provide deeper text understanding for IR. Experimental results demonstrate that the contextual text representations from BERT are more effective than traditional word embeddings. Compared to bag-of-words retrieval models, the contextual language model can better leverage language structures, bringing large improvements on queries written in natural languages. Combining the text understanding ability with search knowledge leads

related words. But word co-occurrence is only a shallow bag-of-words understanding of the text. Recently, we have seen rapid progress in text understanding with the introduction of pre-trained neural language models such as ELMo [8] and BERT [3]. Different from traditional word embeddings, they are *contextual* – the representation of a word is a function of the entire input text, with word dependencies and sentence structures taken into consideration. The models are *pre-trained* on a large number of documents so that the contextual representations encode general language patterns. Contextual neural language models have outperformed traditional word embeddings on a variety of NLP tasks [3, 8].

The deeper text understanding of contextual neural language models brings new possibilities to IR. This paper explores leveraging BERT (Bidirectional Encoder Representations from Transform-

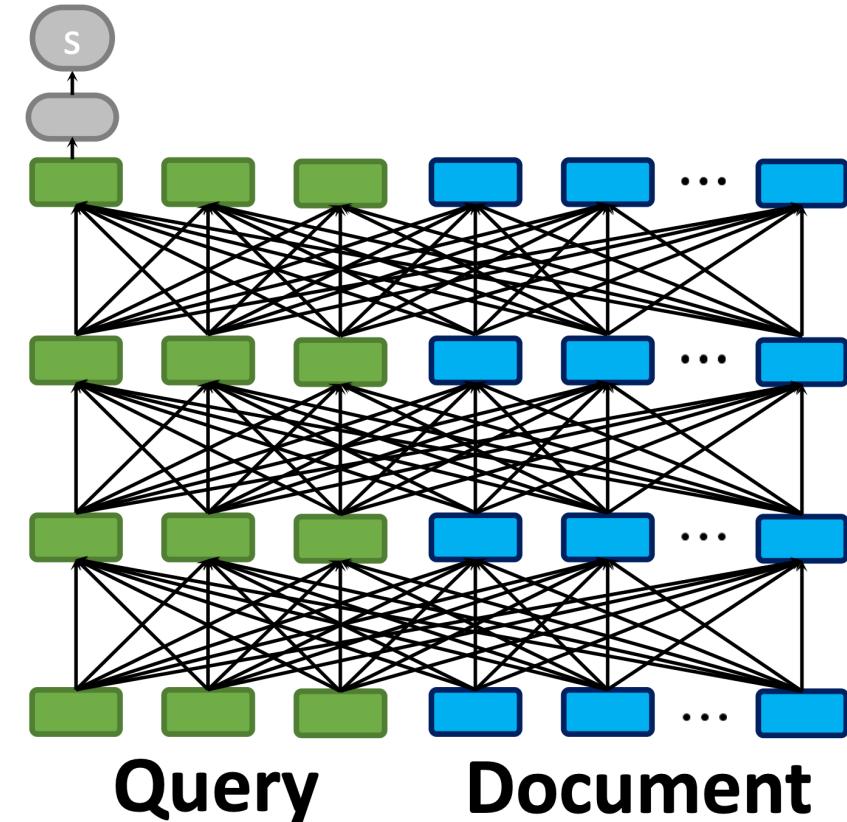
What if the document is too long?

- BERT can take at most **512** tokens
- In a Cross-Encoder architecture, the **document** needs to share the 512 tokens with the **query**
- Divide the document d into passages p_1, p_2, \dots, p_K
 - **FirstP**
$$\text{score}(q, d) = \text{score}(q, p_1)$$
 - **MaxP**
$$\text{score}(q, d) = \max_{1 \leq k \leq K} \text{score}(q, p_k)$$
 - **SumP**
$$\text{score}(q, d) = \sum_{k=1}^K \text{score}(q, p_k)$$

Model	nDCG@20			
	Robust04		ClueWeb09-B	
	Title	Description	Title	Description
BOW	0.417	0.409	0.268	0.234
SDM	0.427	0.427	0.279	0.235
RankSVM	0.420	0.435	0.289	0.245
Coor-Ascent	0.427	0.441	0.295	0.251
DRMM	0.422	0.412	0.275	0.245
Conv-KNRM	0.416	0.406	0.270	0.242
BERT-FirstP	0.444 [†]	0.491 [†]	0.286	0.272[†]
BERT-MaxP	0.469[†]	0.529[†]	0.293	0.262 [†]
BERT-SumP	0.467 [†]	0.524 [†]	0.289	0.261

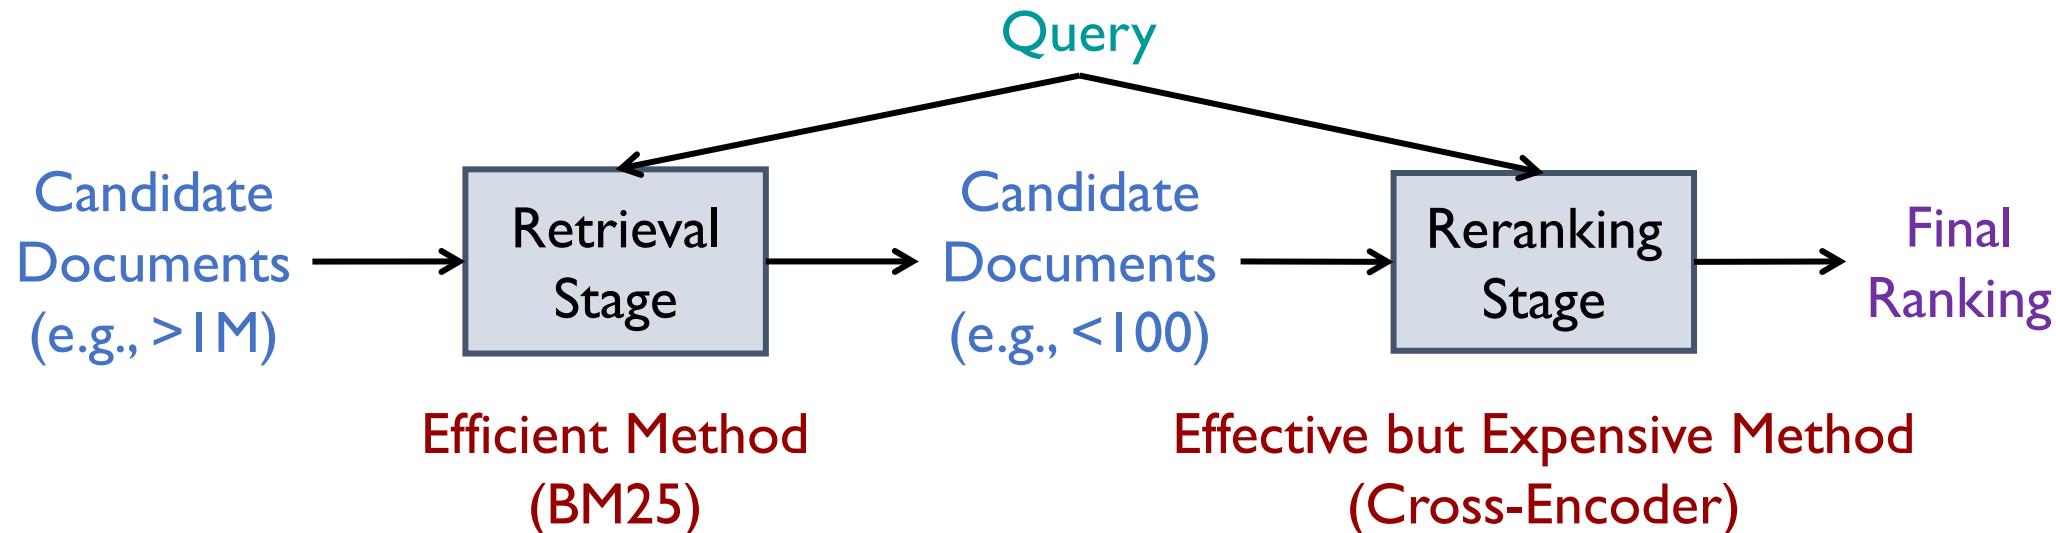
Fine-Tune a Cross-Encoder Ranking Model

- All-to-All Interaction
 - The context of each **query word** includes the entire **query** AND the entire **document**
 - The context of each **document word** includes the entire **query** AND the entire **document**
- What potential issues might a Cross-Encoder have?
 - Cross-Encoder is expensive!
 - You cannot precompute query and word representations
 - How to solve this?



Retrieval-Reranking Paradigm (BERT Version)

- We want to use effective but expensive ranking models ...
- ... only for a more fine-grained ranking of the most relevant documents.



Cross-Encoder Reranking Model [Nogueira and Cho, arXiv 2019]

PASSAGE RE-RANKING WITH BERT

Rodrigo Nogueira

New York University

rodrigonogueira@nyu.edu

Kyunghyun Cho

New York University

Facebook AI Research

CIFAR Azrieli Global Scholar

kyunghyun.cho@nyu.edu

ABSTRACT

Recently, neural models pretrained on a language modeling task, such as ELMo (Peters et al., 2017), OpenAI GPT (Radford et al., 2018), and BERT (Devlin et al., 2018), have achieved impressive results on various natural language processing tasks such as question-answering and natural language inference. In this paper, we describe a simple re-implementation of BERT for query-based passage re-ranking. Our system is the state of the art on the TREC-CAR dataset and the top entry in the leaderboard of the MS MARCO passage retrieval task, outperforming the previous state of the art by 27% (relative) in MRR@10. The code to reproduce our results is available at <https://github.com/nyu-dl/dl4marco-bert>

BERT-Based Neural Re-Ranking

Method	MS MARCO MRR@10		TREC-CAR
	Dev	Eval	MAP Test
BM25 (Lucene, no tuning)	16.7	16.5	12.3
BM25 (Anserini, tuned)	-	-	15.3
Co-PACRR* (MacAvaney et al., 2017)	-	-	14.8
KNRM (Xiong et al., 2017)	21.8	19.8	-
Conv-KNRM (Dai et al., 2018)	29.0	27.1	-
IRNet [†]	27.8	28.1	-
BERT Base	34.7	-	31.0
BERT Large	36.5	35.8	33.5

New SOTA on MS MARCO Leaderboard!

Rank	Model	Submission Date	MRR@10 On Eval
1	BERT + Small Training Rodrigo Nogueira and Kyunghyun Cho - New York University	January 7th, 2019	35.87
2	IRNet (Deep CNN/IR Hybrid Network) Dave DeBarr, Navendu Jain, Robert Sim, Justin Wang, Nirupama Chandrasekaran – Microsoft	January 2nd, 2019	28.061

MS MARCO

<https://microsoft.github.io/msmarco/>

MS MARCO

Home Document Ranking Passage Ranking Updates Submissions About



MS MARCO

X Follow @MSMarcoAI

Starting with a paper released at NIPS 2016, MS MARCO is a collection of datasets focused on deep learning in search.

The first dataset was a question answering dataset featuring 100,000 real Bing questions and a human generated answer. Since then we released a 1,000,000 question dataset, a natural language generation dataset, a passage ranking dataset, keyphrase extraction dataset, crawling dataset, and a conversational search.

MS MARCO

- Largest public IR benchmark release in 2016
 - Used by TREC from 2019 to 2024
 - Consists of more than 500K Bing search queries
 - Sparse labels: approximately one relevance document per query!
 - Passage Ranking: 9M short passages
 - Document Ranking: 3M long documents

MS MARCO Document Ranking Leaderboard

MS MARCO Document Ranking Leaderboard							Search: <input type="text"/>								
date	↑↓	↑↓	description	↑	team	↑↓	paper	↑↓	code	↑↓	type	↑↓	MRR@100 (Dev)	↑↓	MRR@100 (Eval)
2022/02/08	🏆	coCondenser(maxp) + MORES+			Luyu Gao - Carnegie Mellon University						full ranking		0.512		0.446
2021/07/14	🏆	UniRetriever			Microsoft-Research-Asia and STCA-BingAdsSelection						full ranking		0.500		0.440

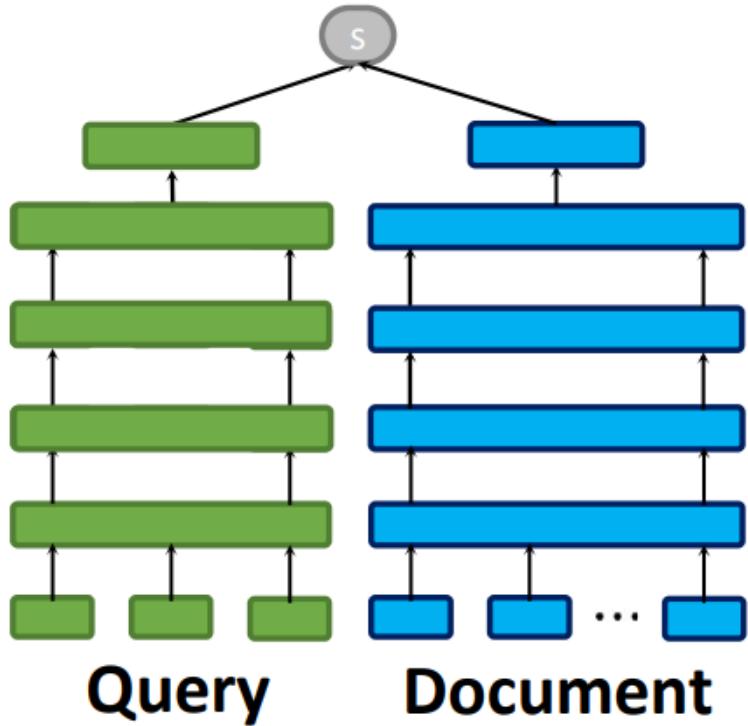
MRR (Mean Reciprocal Rank)

$$\text{MRR} = \frac{1}{|Q|} \sum_{q \in Q} \frac{1}{\text{rank}_q}$$

- Q : the set of queries
- rank_q : rank position of the first relevant result for query q
- Example
 - Suppose we have 4 candidate documents. For the following 3 queries, each query has a ranking list of the 4 documents, with the relevance ground truth as follows:
 - Query q_1 : [0, 0, 0, 1]
 - Query q_2 : [1, 1, 0, 1]
 - Query q_3 : [0, 1, 1, 1]
 - $\text{MRR} = \frac{1}{3} \times \left(\frac{1}{4} + \frac{1}{1} + \frac{1}{2} \right) = \frac{7}{12} \approx 0.583$

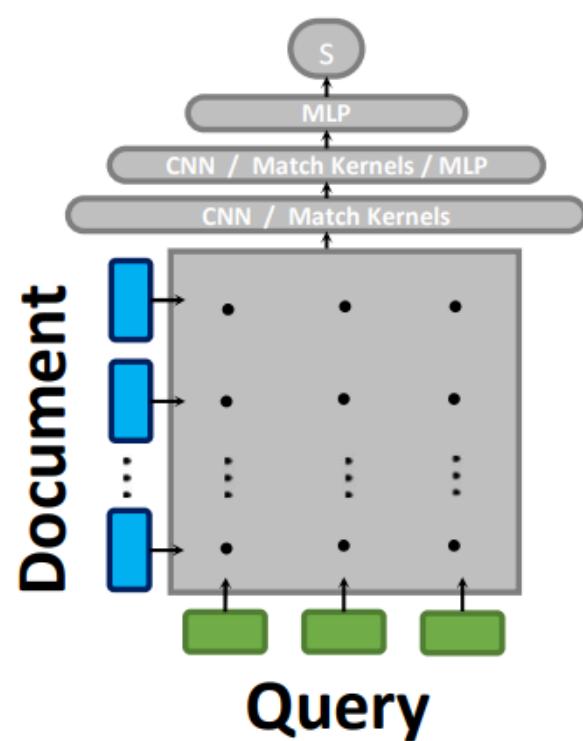
Questions?

Previously Introduced Neural Ranking Paradigms



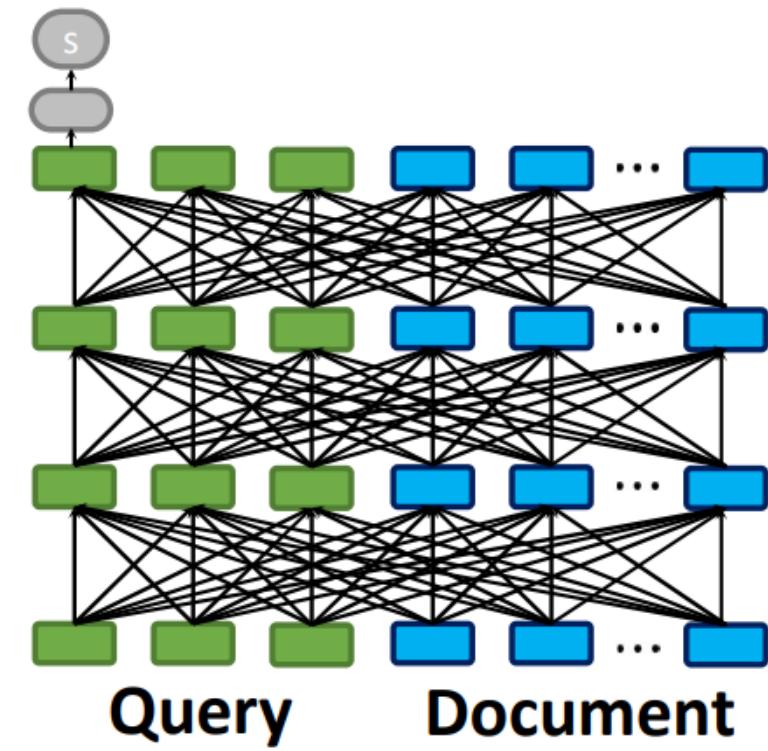
(a) Representation-based Similarity

- DESM, DPR, BERT Bi-Encoder
- Efficient
- Independent **Query/Doc** encoding



(b) Query-Document Interaction

- Duet, Conv-KNRM
- Effective but expensive
- $\text{context}(\text{Query}) = \text{Doc}$

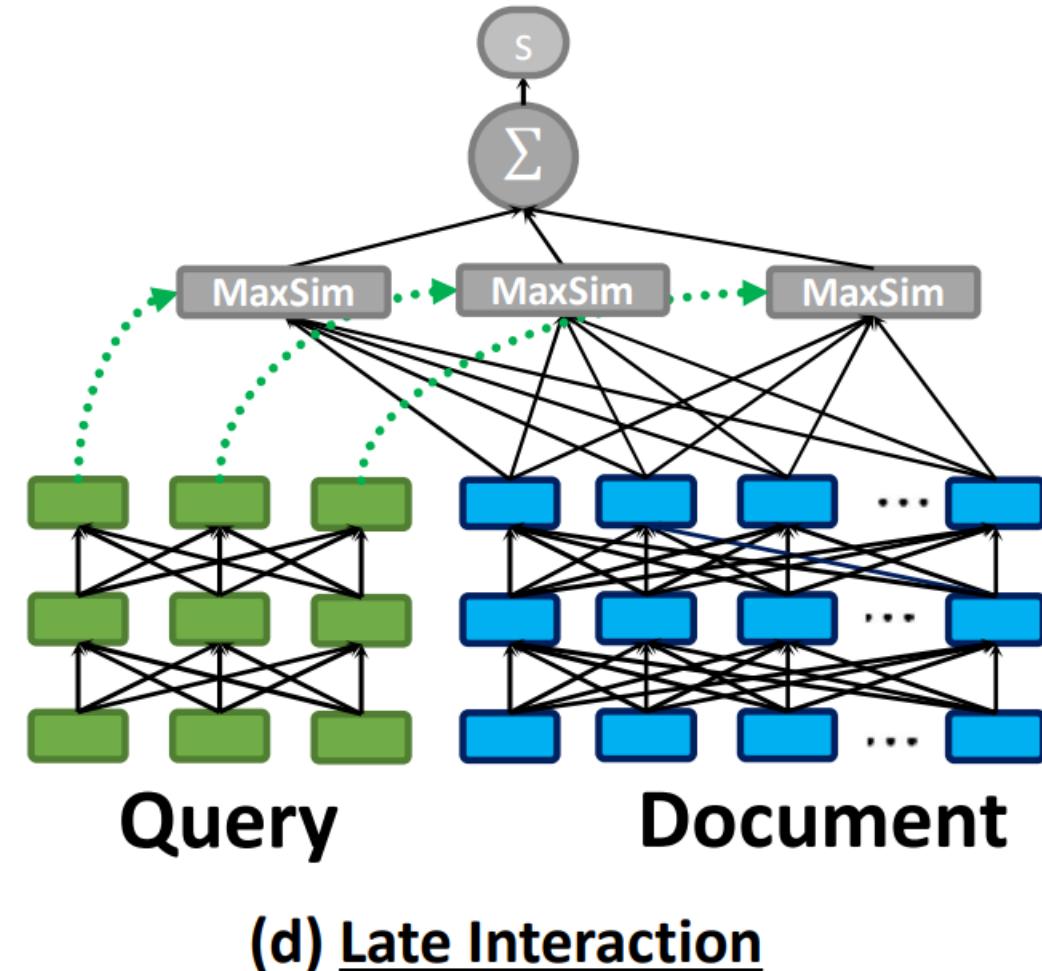


(c) All-to-all Interaction

- BERT Cross-Encoder
- Effective but expensive
- $\text{context}(\text{Query}) = \text{Query} \& \text{Doc}$

Late Interaction

- Can we keep precomputation and still have **fine-grained query-document** interactions?
- Desired Properties:
 - Independent encoding
 - Fine-grained representations
 - End-to-end retrieval (the retrieval-reranking paradigm can be used, but it is not mandatory)



ColBERT [Khattab and Zaharia, SIGIR 2020]

ColBERT: Efficient and Effective Passage Search via Contextualized Late Interaction over BERT

Omar Khattab

Stanford University

okhattab@stanford.edu

Matei Zaharia

Stanford University

matei@cs.stanford.edu

ABSTRACT

Recent progress in Natural Language Understanding (NLU) is driving fast-paced advances in Information Retrieval (IR), largely owed to fine-tuning deep language models (LMs) for document ranking. While remarkably effective, the ranking models based on these LMs increase computational cost by orders of magnitude over prior approaches, particularly as they must feed each query–document pair through a massive neural network to compute a single relevance score. To tackle this, we present ColBERT, a novel ranking model that adapts deep LMs (in particular, BERT) for efficient retrieval. ColBERT introduces a *late interaction* architecture that independently encodes the query and the document using BERT and then employs a cheap yet powerful interaction step that models their

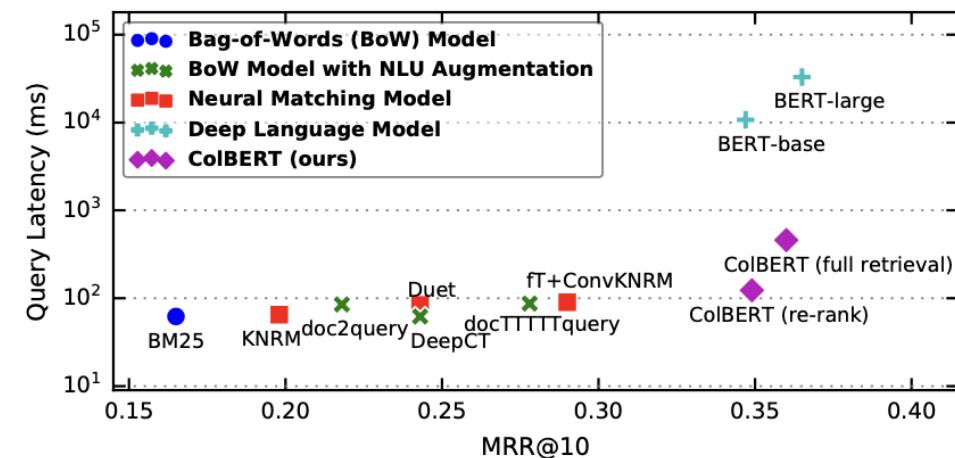


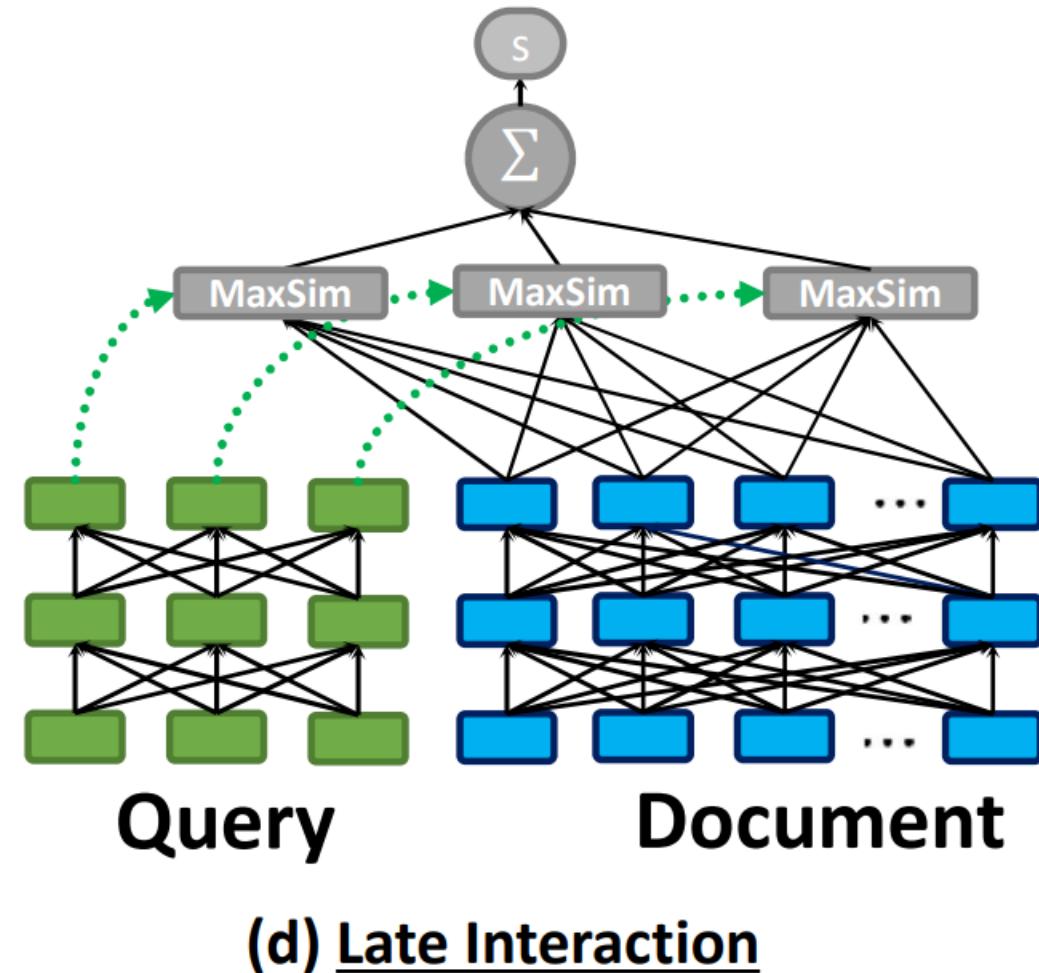
Figure 1: Effectiveness (MRR@10) versus Mean Query Latency (log-scale) for a number of representative ranking

Late Interaction

- Step 1: Encode the document into a sequence of vectors $e_{d_1}, e_{d_2}, \dots, e_{d_L}$
- Step 2: Encode the query into a sequence of vectors $e_{q_1}, e_{q_2}, \dots, e_{q_M}$
- Step 3:

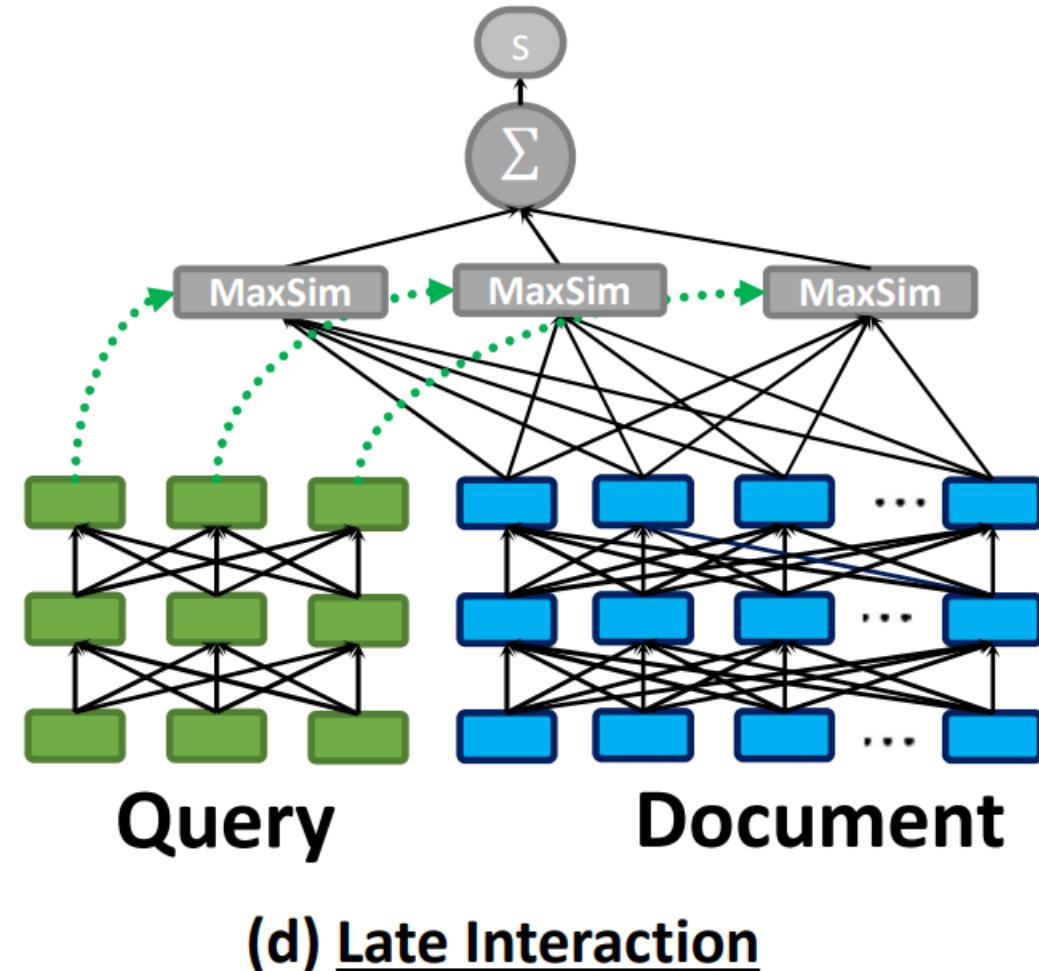
$$\text{score}(q, d) = \sum_{m=1}^M \max_{1 \leq l \leq L} e_{q_m}^T e_{d_l}$$

- Intuition: For each word in the query, find the most matching word in the current candidate document to compute the similarity



Late Interaction

- Example
 - **Query** representation
 $[1, 0], [0, 1]$
 - **Document** representation
 $[1, 1], [0.5, 0.5]$
 - For $[1, 0]$, which **document** vector is the most similar?
 - For $[0, 1]$, which **document** vector is the most similar?
 - $\text{score}(q, d)$
 $= [1, 0] \begin{bmatrix} 1 \\ 1 \end{bmatrix} + [0, 1] \begin{bmatrix} 1 \\ 1 \end{bmatrix} = 2$



Late Interaction: Real Example of Matching

when did the transformers cartoon series come out?

[...] the animated [...] The Transformers [...] [...] It was released [...] **on** August 8, 1986

when did the transformers cartoon series come out?

[...] the animated [...] The **Transformers** [...] [...] It was released [...] on August 8, 1986

when did the transformers cartoon series come out?

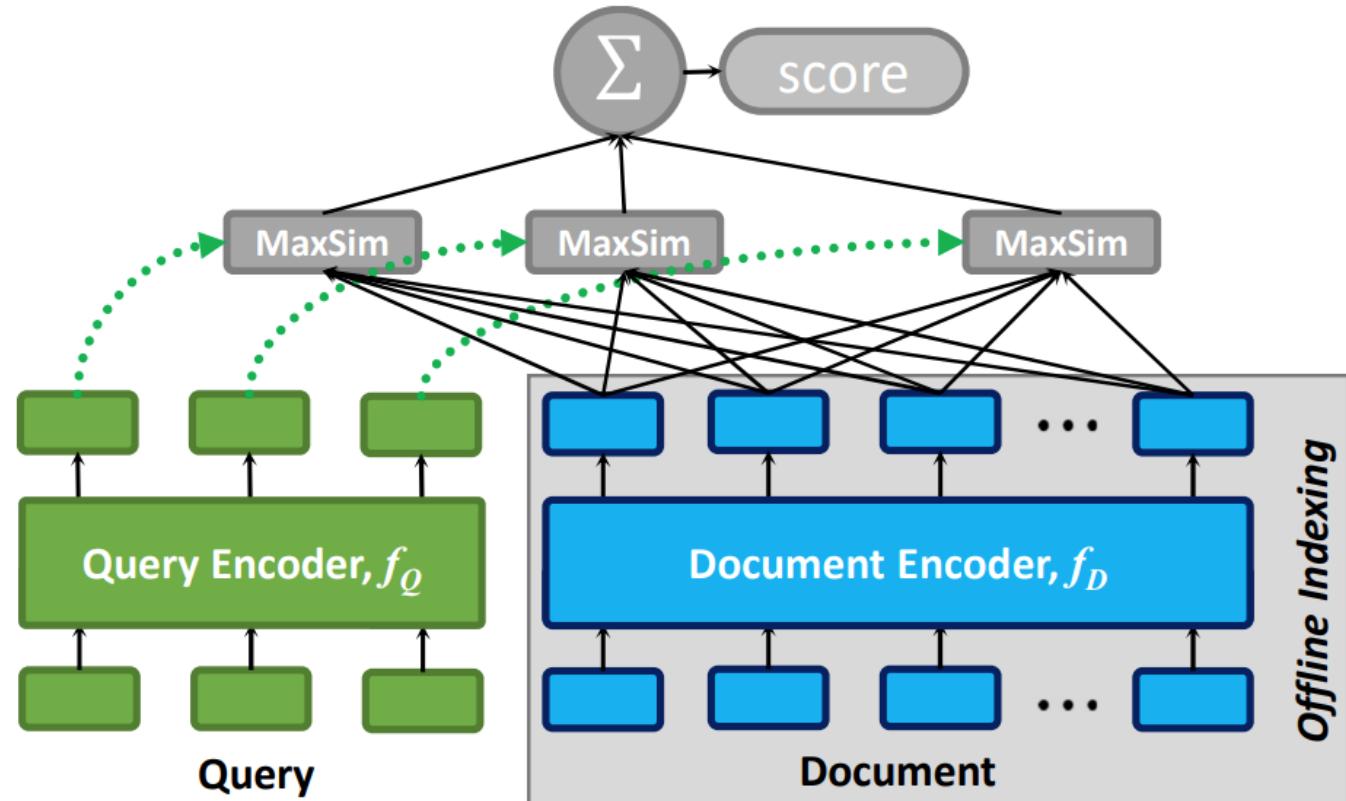
[...] the **animated** [...] The Transformers [...] [...] It was released [...] on August 8, 1986

when did the transformers cartoon series come out?

[...] the animated [...] The Transformers [...] [...] It was **released** [...] on August 8, 1986

Offline Indexing

- For Bi-Encoder, index **one vector** for each candidate document
- For ColBERT, index **a sequence of vectors** for each candidate document



Performance of ColBERT: Reranking

- On par with BERT-Base (Cross-Encoder)
- Slightly worse than BERT-Large (Cross-Encoder)
- BUT much faster!

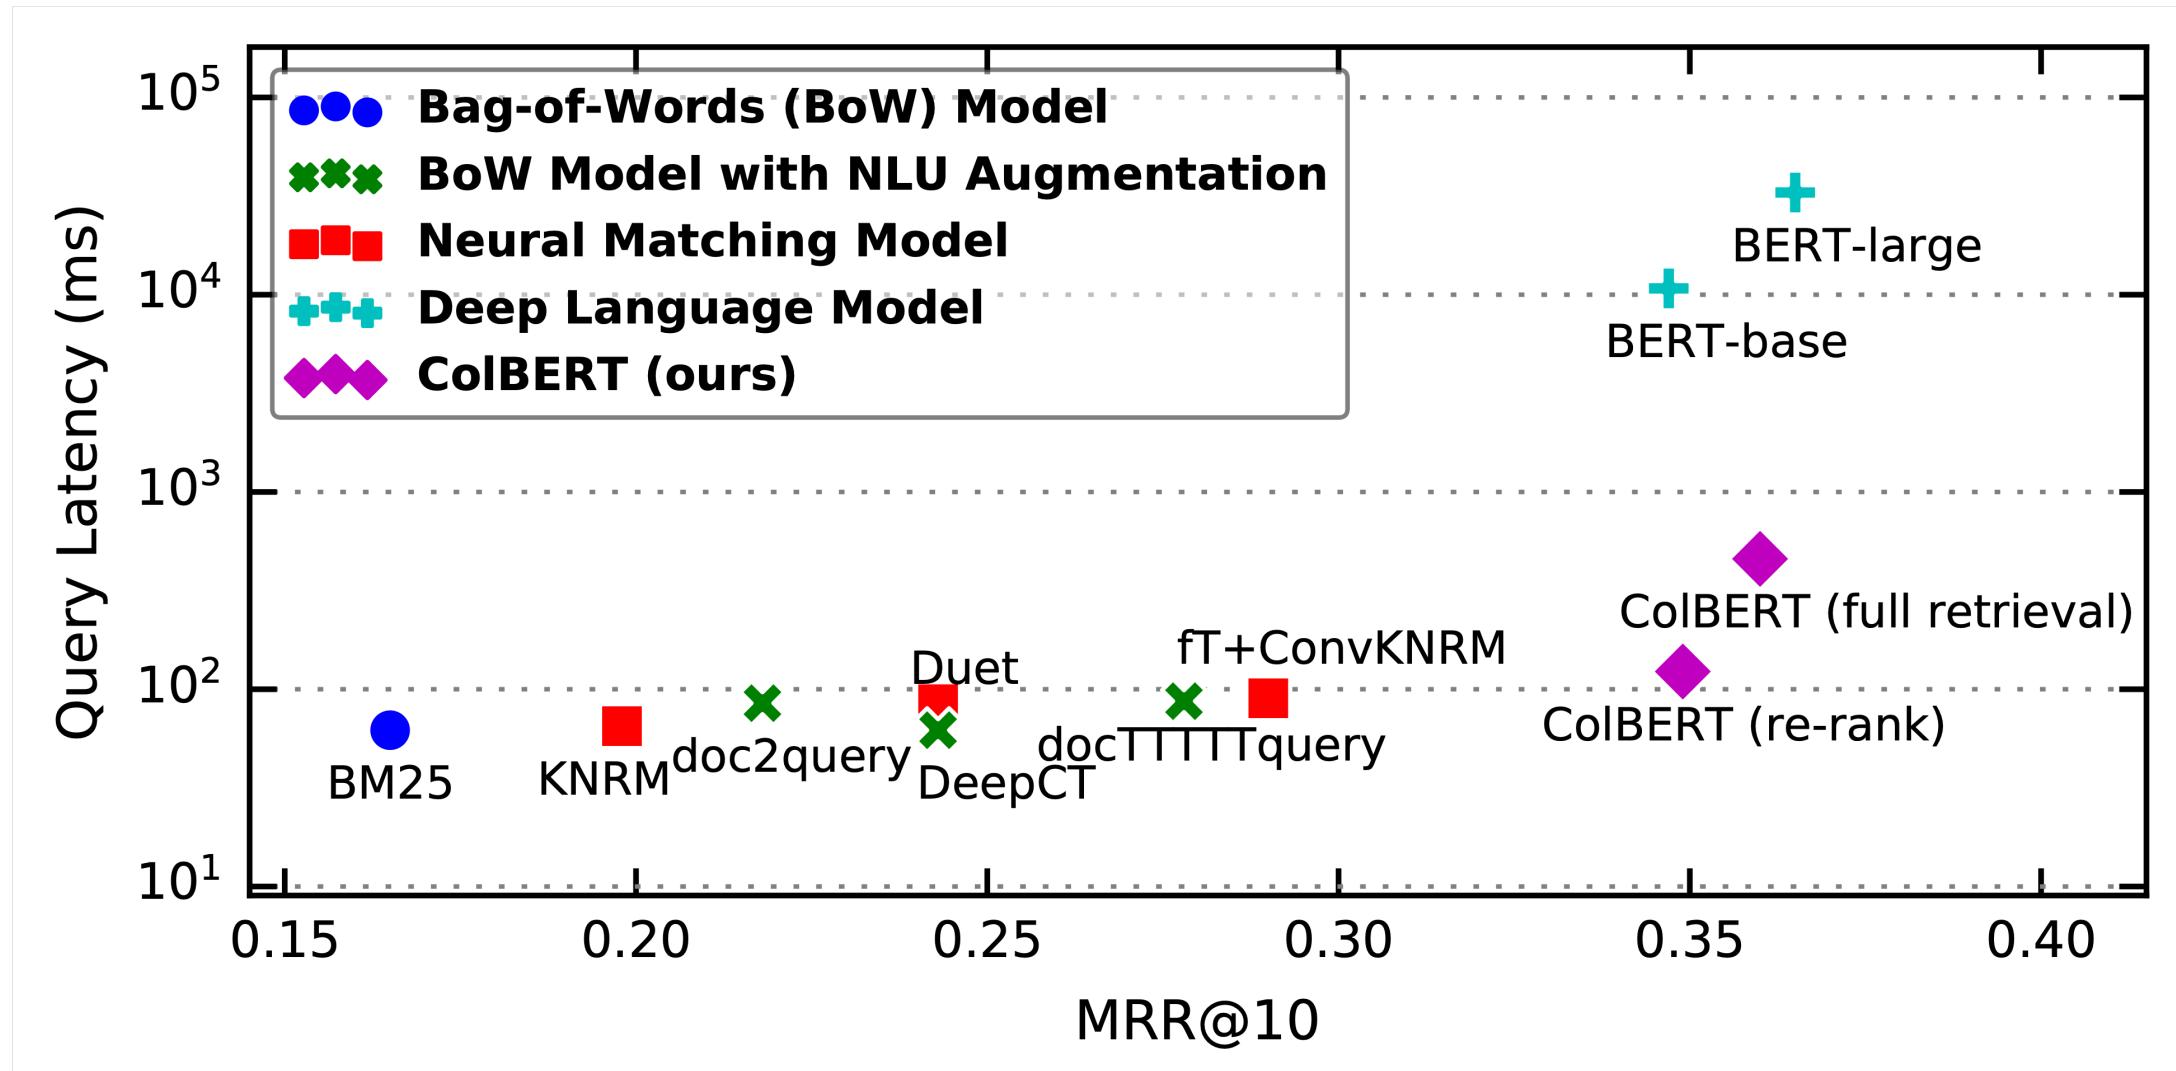
Method	MRR@10 (Dev)	MRR@10 (Eval)	Re-ranking Latency (ms)	FLOPs/query
BM25 (official)	16.7	16.5	-	-
KNRM	19.8	19.8	3	592M (0.085×)
Duet	24.3	24.5	22	159B (23×)
fastText+ConvKNRM	29.0	27.7	28	78B (11×)
BERT _{base} [25]	34.7	-	10,700	97T (13,900×)
BERT _{base} (our training)	36.0	-	10,700	97T (13,900×)
BERT _{large} [25]	36.5	35.9	32,900	340T (48,600×)
ColBERT (over BERT _{base})	34.9	34.9	61	7B (1×)

Performance of ColBERT: End-to-End Retrieval

- <10× slower than BM25 and Bi-Encoder ranking models
- BUT much more effective!

Method	MRR@10 (Dev)	MRR@10 (Local Eval)	Latency (ms)	Recall@50	Recall@200	Recall@1000
BM25 (official)	16.7	-	-	-	-	81.4
BM25 (Anserini)	18.7	19.5	62	59.2	73.8	85.7
doc2query	21.5	22.8	85	64.4	77.9	89.1
DeepCT	24.3	-	62 (est.)	69 [2]	82 [2]	91 [2]
docTTTTTquery	27.7	28.4	87	75.6	86.9	94.7
ColBERT _{L2} (re-rank)	34.8	36.4	-	75.3	80.5	81.4
ColBERT _{L2} (end-to-end)	36.0	36.7	458	82.9	92.3	96.8

Striking a Good Balance Between Bi-Encoder and Cross-Encoder



Robustness: Out-of-Domain Quality

- So far, we have looked at **in-domain** effectiveness evaluations
 - We had training and evaluation data for MS MARCO
 - We often want to use retrieval in new, **out-of-domain** settings with NO training data and NO validation data
 - This is sometimes called a “**zero-shot**” setting; it emphasizes **generalization**
 - BEIR is a popular benchmark for IR models in “zero-shot” scenarios



BEIR

Split (→)					Train	Dev	Test			Avg. Word Lengths	
Task (↓)	Domain (↓)	Dataset (↓)	Title	Relevancy	#Pairs	#Query	#Query	#Corpus	Avg. D / Q	Query	Document
Passage-Retrieval	Misc.	MS MARCO [45]	✗	Binary	532,761	—	6,980	8,841,823	1.1	5.96	55.98
Bio-Medical Information Retrieval (IR)	Bio-Medical	TREC-COVID [65]	✓	3-level	—	—	50	171,332	493.5	10.60	160.77
	Bio-Medical	NFCorpus [7]	✓	3-level	110,575	324	323	3,633	38.2	3.30	232.26
	Bio-Medical	BioASQ [61]	✓	Binary	32,916	—	500	14,914,602	4.7	8.05	202.61
Question Answering (QA)	Wikipedia	NQ [34]	✓	Binary	132,803	—	3,452	2,681,468	1.2	9.16	78.88
	Wikipedia	HotpotQA [76]	✓	Binary	170,000	5,447	7,405	5,233,329	2.0	17.61	46.30
	Finance	FiQA-2018 [44]	✗	Binary	14,166	500	648	57,638	2.6	10.77	132.32
Tweet-Retrieval	Twitter	Signal-1M (RT) [59]	✗	3-level	—	—	97	2,866,316	19.6	9.30	13.93
News Retrieval	News	TREC-NEWS [58]	✓	5-level	—	—	57	594,977	19.6	11.14	634.79
	News	Robust04 [64]	✗	3-level	—	—	249	528,155	69.9	15.27	466.40
Argument Retrieval	Misc.	ArguAna [67]	✓	Binary	—	—	1,406	8,674	1.0	192.98	166.80
	Misc.	Touché-2020 [6]	✓	3-level	—	—	49	382,545	19.0	6.55	292.37
Duplicate-Question Retrieval	StackEx.	CQA DupStack [25]	✓	Binary	—	—	13,145	457,199	1.4	8.59	129.09
	Quora	Quora	✗	Binary	—	5,000	10,000	522,931	1.6	9.53	11.44
Entity-Retrieval	Wikipedia	DBPedia [21]	✓	3-level	—	67	400	4,635,922	38.2	5.39	49.68
Citation-Prediction	Scientific	SCIDOCs [9]	✓	Binary	—	—	1,000	25,657	4.9	9.38	176.19
Fact Checking	Wikipedia	FEVER [60]	✓	Binary	140,085	6,666	6,666	5,416,568	1.2	8.13	84.76
	Wikipedia	Climate-FEVER [14]	✓	Binary	—	—	1,535	5,416,593	3.0	20.13	84.76
	Scientific	SciFact [68]	✓	Binary	920	—	300	5,183	1.1	12.37	213.63

Robustness: Out-of-Domain NDCG@10

IR Task	Classical IR BM25	Interaction Models ELECTRA re-ranker	Representation Similarity DPR	Representation Similarity SBERT	Late Interaction CoBERT
BioMed	48	49	22	34	49
QA	38	51	33	41	48
Tweet	39	31	16	26	27
News	37	43	16	37	39
Arguments	52	35	15	34	25
Duplicates	53	56	20	58	60
Entity	29	38	26	34	39
Citation	16	15	8	13	15
Fact-Check	48	52	34	47	54
Overall Avg	42	45	23	39	44

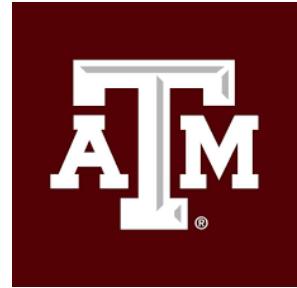
Robustness: Out-of-Domain Recall@100

IR Task	Classical IR BM25	Interaction Models ELECTRA re-ranker	Representation Similarity DPR	Representation Similarity SBERT	Late Interaction ColBERT
BioMed	45	45	23	35	45
QA	67	67	60	68	75
Tweet	38	38	16	26	28
News	40	40	22	37	37
Arguments	70	70	46	62	61
Duplicates	77	77	44	79	81
Entity	38	38	35	40	46
Citation	35	35	22	30	34
Fact-Check	71	71	65	74	75
Overall Avg	59	59	43	57	61

How to build effective “zero-shot” rankers?

Next Lecture

- How to leverage neural networks and embeddings in recommender systems (e.g., for collaborative filtering)?
- Quiz 3!
 - All policies are the same as Quiz 1 (number of questions, time limit, grading, etc.)
 - Scope:
 - Lecture 14 (Bayesian Personalized Ranking)
 - Lecture 15 (Word Embedding)
 - Lecture 16 (Neural Ranking)
 - Lecture 17 (Transformer, BERT)
 - Lecture 18 (BERT-Based Ranking)
 - Homework 2



Thank You!

Course Website: <https://yuzhang-teaching.github.io/CSCE670-F25.html>