



CSCE 670 - Information Storage and Retrieval

Week 6: Fundamentals of Recommender Systems

Yu Zhang

yuzhang@tamu.edu

Course Website: <https://yuzhang-teaching.github.io/CSCE670-F26.html>

Recap: Phase 1

- Phase 1: Search Engines

- Boolean and ranked retrieval, link analysis, evaluation, learning to rank (ML + ranking), ...

- Phase 2: Recommender Systems


- content-based recommendation, collaborative filtering, matrix factorization, recommendation with implicit feedback, ...

- Phase 3: From Foundations to Modern Methods

- word embedding learning, Transformer, “small” language models, ... (for search and recommendation)

- Phase 4: Large Language Models (!!)

Our Capabilities So Far



- Given a query, find *relevant* CDs
 - Exact matching (Boolean, phrase, proximity, wildcard)
 - Ranked retrieval (TF-IDF, BM25, learning to rank)
 - Link analysis (PageRank, HITS)

Phase 2

- **Phase 1: Search Engines**

- Boolean and ranked retrieval, link analysis, evaluation, learning to rank (ML + ranking), ...

- **Phase 2: Recommender Systems**


- content-based recommendation, collaborative filtering, matrix factorization, recommendation with implicit feedback, ...

- **Phase 3: From Foundations to Modern Methods**

- word embedding learning, Transformer, “small” language models, ... (for search and recommendation)

- **Phase 4: Large Language Models (!!)**

Example of Recommender Systems




- Customer X
 - Buys Metallica CD
 - Buys Megadeth CD



- Customer Y
 - Does search on Metallica
 - Recommender systems should suggest Megadeth from data collected about Customer X

Amazon

Related to items you've viewed




LinkedIn

You might like Pages for you



Texas A&M College of
Agriculture and Life
Sciences
Higher Education
5,952 followers

Example of Recommender Systems




Examples:



Why do we need recommendation?

- Shelf space is a **scarce** commodity for traditional retailers
 - Also: TV networks, movie theaters,...
- Web enables near-zero-cost dissemination of information about products
 - From **scarcity** to **abundance**
- More choice necessitates better filters
 - How *Into Thin Air* (published in 1997) made *Touching the Void* (published in 1988) a bestseller
 - *Touching the Void* did not become a bestseller until a similar bestseller book appears 9 years later.
 - Amazon's recommendation engine

The Long Tail




The Long Tail

THE BIT PLAYER ADVANTAGE

Beyond bricks and mortar there are two main retail models – one that gets halfway down the Long Tail and another that goes all the way. The first is the familiar hybrid model of Amazon and Netflix, companies that sell physical goods online. Digital catalogs allow them to offer unlimited selection along with search, reviews, and recommendations, while the cost savings of massive warehouses and no walk-in customers greatly expands the number of products they can sell profitably.

Pushing this even further are pure digital services, such as iTunes, which offer the additional savings of delivering their digital goods online at virtually no marginal cost. Since an extra database entry and a few megabytes of storage on a server cost effectively nothing, these retailers have no economic reason not to carry *everything* available.



Types of Recommendations

- Editorial and hand-curated (not personalized)
 - “*Store Manager’s Pick*”
 - Promoted items
- Simple aggregates (not personalized)
 - Most liked/clicked this month/week/day
 - Most recent
- Personalized approaches
 - Content-Based (this week)
 - Collaborative Filtering (this week)
 - Matrix Factorization (this week and next week)
 - Bayesian Personalized Ranking (next week)

Formal Setup

- $\mathcal{X}$: A set of users
- $\mathcal{S}$: A set of items
- **Utility function** $u: \mathcal{X} \times \mathcal{S} \rightarrow \mathcal{R}$
 - $\mathcal{R}$ = set of ratings, which is a totally ordered set
 - E.g., 1-5 stars
 - E.g., real number in $[0,1]$

Utility Matrix U

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

Key Problems

- **Problem 1:** Gathering “known” ratings for matrix
 - How to collect the data in the utility matrix?
 - Evaluating the quality of an item solely based on its average rating?
- **Problem 2:** Extrapolate unknown ratings from the known ones
 - Mainly interested in high unknown ratings
 - We are not interested in knowing what you don't like but what you like
- **Problem 3:** Evaluating extrapolation methods
 - How to measure success/performance of recommendation methods?

Problem 1: Gathering Ratings

- Explicit
 - Ask people to rate items
 - Doesn't work well in practice – people can't be bothered
- Implicit
 - Learn ratings from user actions
 - E.g., purchase implies high rating
 - What about low ratings?


Problem 2: Extrapolating Utilities

- **Key problem:** Utility matrix U is sparse
 - Most people have not rated most items
 - **Cold start:**
 - New items have no ratings
 - New users have no history
- Solutions to be introduced today:
 - **Content-Based Approach**
 - **Collaborative Filtering**

Content-Based Approach (Calculating User-Item Similarity)

Content-Based Recommender Systems

- **Idea:** Recommend items to user x similar to previous items rated highly by x
- **Example:**
 - **Movie recommendations:** Recommend movies with same actor(s), director, genre, ...
 - **Websites, blogs, news:** Recommend other sites with “similar” content



Profiles

- Item Profile i
 - A set (vector) of features
 - Movies: author, title, actor, director, ...
 - Text: Set of “important” words in document (e.g., based on TF-IDF)
- User Profile x
 - Weighted average of rated item profiles
- Prediction
 - Given x and i , estimate $u(x, i) = \cos(x, i) = \frac{x^T i}{\|x\| \cdot \|i\|}$

Example

- Rating scale: $\mathcal{R} = \{1, 0, -1\}$
- User x has rated 3 songs
 - Song 1: “sunny day”, rating: 1
 - Song 2: “cloudy day”, rating: 0
 - Song 3: “rainy day”, rating: -1
- Let’s simplify the model and use Boolean representation (rather than TF-IDF)
- Item profiles

	<i>sunny</i>	<i>cloudy</i>	<i>rainy</i>	<i>day</i>
Song 1	1	0	0	1
Song 2	0	1	0	1
Song 3	0	0	1	1

Example

- User x has rated 3 songs
 - Song 1: “sunny day”, rating: 1
 - Song 2: “cloudy day”, rating: 0
 - Song 3: “rainy day”, rating: -1
- Item profiles

	<i>sunny</i>	<i>cloudy</i>	<i>rainy</i>	<i>day</i>
Song 1	1	0	0	1
Song 2	0	1	0	1
Song 3	0	0	1	1

- User profile: $\text{User } x = 1 \times \text{Song 1} + 0 \times \text{Song 2} + (-1) \times \text{Song 3}$

	<i>sunny</i>	<i>cloudy</i>	<i>rainy</i>	<i>day</i>
User x	1	0	-1	0

Example

- User profile: User $x = 1 \times \text{Song 1} + 0 \times \text{Song 2} + (-1) \times \text{Song 3}$

	<i>sunny</i>	<i>cloudy</i>	<i>rainy</i>	<i>day</i>
User x	1	0	-1	0

- New Song: “sunny cloudy”
- Will User x like it?
- Profile of New Song

	<i>sunny</i>	<i>cloudy</i>	<i>rainy</i>	<i>day</i>
New Song	1	1	0	0

- $u(\text{User } x, \text{New Song}) = \cos \left(\left[\begin{array}{c} 1 \\ 0 \\ -1 \\ 0 \end{array} \right], \left[\begin{array}{c} 1 \\ 1 \\ 0 \\ 0 \end{array} \right] \right) = \frac{1}{\sqrt{2} \times \sqrt{2}} = \frac{1}{2}$

Content-Based Recommender Systems: Pros

- + No need for data on other users
 - No cold-start or sparsity problems
- + Able to recommend to users with unique tastes
- + Able to recommend new & unpopular items
 - No first-rater problem
- + Able to provide explanations
 - Can provide explanations of recommended items by listing content-features that caused an item to be recommended

Content-Based Recommender Systems: Cons


- – Finding the appropriate features is hard
 - E.g., images, movies, music
- – Recommendations for new users
 - How to build a user profile?
- – Overspecialization
 - Never recommends items outside user's content profile
 - People might have multiple interests
 - Unable to exploit quality judgments of other users

Collaborative Filtering

(Harnessing Quality Judgments of Other Users)

Collaborative Filtering

- Consider user x
- **Step 1:** Find a set $\mathcal{N}$ of other users whose ratings are “similar” to x ’s ratings
- **Step 2:** Estimate x ’s ratings based on ratings of users in $\mathcal{N}$



Step 1: Finding Similar Users

- How to define User-User similarity?
- Example:

	Item 1	Item 2	Item 3	Item 4	Item 5
User x	*			*	***
User y	*		**	**	

- (Bad) Solution 1: Jaccard Similarity (between two sets)

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

$$x: \{1, 4, 5\}, \quad y: \{1, 3, 4\}, \quad J(x, y) = \frac{1}{2}$$

Problem: Ignores the **value** of the rating

Step 1: Finding Similar Users

- How to define User-User similarity?
- Example:

	Item 1	Item 2	Item 3	Item 4	Item 5
User x	*			*	***
User y	*		**	**	

- (Bad) Solution 2: Cosine Similarity (between two vectors)
 - $x: [1, 0, 0, 1, 3]^T$
 - $y: [1, 0, 2, 2, 0]^T$
 - Problem: Treats missing ratings as “negative”

Step 1: Finding Similar Users

- How to define User-User similarity?
- Example:

	Item 1	Item 2	Item 3	Item 4	Item 5
User x	*			*	***
User y	*		**	**	

- **Solution 3:** Pearson Correlation Coefficient
 - Consider $\mathcal{S}_{xy}$ = items rated by both users x and y

$$sim(x, y) = \frac{\sum_{i \in \mathcal{S}_{xy}} (U_{xi} - \bar{U}_x)(U_{yi} - \bar{U}_y)}{\sqrt{\sum_{i \in \mathcal{S}_{xy}} (U_{xi} - \bar{U}_x)^2} \sqrt{\sum_{i \in \mathcal{S}_{xy}} (U_{yi} - \bar{U}_y)^2}}$$

$\bar{U}_x, \bar{U}_y$: average rating given by x and y

How to understand the Pearson Correlation Coefficient?

$$\text{sim}(x, y) = \frac{\sum_{i \in \mathcal{S}_{xy}} (U_{xi} - \bar{U}_x)(U_{yi} - \bar{U}_y)}{\sqrt{\sum_{i \in \mathcal{S}_{xy}} (U_{xi} - \bar{U}_x)^2} \sqrt{\sum_{i \in \mathcal{S}_{xy}} (U_{yi} - \bar{U}_y)^2}}$$

- Original Table

	Item 1	Item 2	Item 3	Item 4	Item 5
User x	*			*	***
User y	*		**	**	

- **Step 1:** Subtract the (row) mean

	Item 1	Item 2	Item 3	Item 4	Item 5
User x	$1 - 5/3 = -2/3$			$1 - 5/3 = -2/3$	$3 - 5/3 = 4/3$
User y	$1 - 5/3 = -2/3$		$2 - 5/3 = 1/3$	$2 - 5/3 = 1/3$	

How to understand the Pearson Correlation Coefficient?

$$sim(x, y) = \frac{\sum_{i \in \mathcal{S}_{xy}} (U_{xi} - \bar{U}_x)(U_{yi} - \bar{U}_y)}{\sqrt{\sum_{i \in \mathcal{S}_{xy}} (U_{xi} - \bar{U}_x)^2} \sqrt{\sum_{i \in \mathcal{S}_{xy}} (U_{yi} - \bar{U}_y)^2}}$$

- Original Table

	Item 1	Item 2	Item 3	Item 4	Item 5
User x	*			*	***
User y	*		**	**	

- **Step 2:** Only keep the column rated by both x and y

	Item 1	Item 4
User x	$1 - 5/3 = -2/3$	$1 - 5/3 = -2/3$
User y	$1 - 5/3 = -2/3$	$2 - 5/3 = 1/3$

- **Step 3:** Calculate the **Cosine Similarity**
 - Pearson is equivalent to Cosine after some data normalization steps!

Step 2: Rating Prediction

- Let $\mathcal{N}$ be the set of k users that are most similar to x who have rated item i
- Prediction for item i of user x :

- Simple average:

$$U_{xi} = \frac{1}{k} \sum_{y \in \mathcal{N}} U_{yi}$$

- Weighted by User-User similarity:

$$U_{xi} = \frac{\sum_{y \in \mathcal{N}} \text{sim}(x, y) \cdot U_{yi}}{\sum_{y \in \mathcal{N}} \text{sim}(x, y)}$$

- Many other tricks possible...
- So far, User-User Collaborative Filtering
 - Using User-User similarity to predict User-Item similarity

How about Item-Item Collaborative Filtering?

- Using **Item-Item** similarity to predict **User-Item** similarity
- **Step 1**: For item i , find other similar items
- **Step 2**: Estimate rating for item i based on ratings for similar items
- Can use same similarity metrics and prediction functions as in user-user model
 - E.g., Weighted by Item-Item similarity:

$$U_{xi} = \frac{\sum_{j \in \mathcal{N}} \text{sim}(i, j) \cdot U_{xj}}{\sum_{j \in \mathcal{N}} \text{sim}(i, j)}$$

$\text{sim}(i, j)$: Pearson Correlation Coefficient between item i and item j

$\mathcal{N}$: the set of **items** rated by x that are similar to i

Example ($|\mathcal{N}| = 2$)

users

CDs

	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3			5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	

- unknown rating


- rating between 1 to 5

Example ($|\mathcal{N}| = 2$)

users

	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3		?	5			5		4	
2			5	4			4			2	1	3
3	2	4		1	2		3		4	3	5	
4		2	4		5			4			2	
5			4	3	4	2					2	5
6	1		3		3			2			4	

CDs

 - estimate rating of CD 1 by user 5

Example ($|\mathcal{N}| = 2$)

		users												
		1	2	3	4	5	6	7	8	9	10	11	12	$sim(1,\cdot)$
CDs	1	1		3		?	5			5		4		1.00
	2			5	4			4			2	1	3	-0.18
	<u>3</u>	2	4		1	2		3		4	3	5		0.41
	4		2	4		5			4			2		-0.10
	5			4	3	4	2					2	5	-0.31
	<u>6</u>	1		3		3			2			4		0.59

Neighbor selection: Identify movies that are similar to CD 1, rated by user 5

Example ($|\mathcal{N}| = 2$)

		users												
		1	2	3	4	5	6	7	8	9	10	11	12	$sim(1,\cdot)$
CDs	1	1		3		2.6	5			5		4		1.00
	2			5	4			4			2	1	3	-0.18
	<u>3</u>	2	4		1	2		3		4	3	5		0.41
	4		2	4		5			4			2		-0.10
	5			4	3	4	2					2	5	-0.31
	<u>6</u>	1		3		3			2			4		0.59

Predict by taking weighted average:

$$U_{51} = (0.41 \times 2 + 0.59 \times 3) / (0.41 + 0.59) = 2.6$$

Collaborative Filtering: Common Practice

- So far,

$$U_{xi} = \frac{\sum_{j \in \mathcal{N}} \text{sim}(i, j) \cdot U_{xj}}{\sum_{j \in \mathcal{N}} \text{sim}(i, j)}$$

- In practice,

$$U_{xi} = b_{xi} + \frac{\sum_{j \in \mathcal{N}} \text{sim}(i, j) \cdot (U_{xj} - b_{xj})}{\sum_{j \in \mathcal{N}} \text{sim}(i, j)}$$

- b_{xi} : baseline estimate for U_{xi} ($b_{xi} = \mu + b_x + b_i$)
- μ : overall mean CD rating
- b_x : rating deviation of user x , which is the (avg. rating given by user x) $- \mu$
- b_i : rating deviation of item i , which is the (avg. rating given to item i) $- \mu$

Item-Item vs. User-User

- In practice, it has been observed that item-item often works better than user-user!
- Why? Items are simpler, users have multiple tastes
 - It is impossible for a piece of music to be both 60's rock and 1700's baroque.
 - There are individuals who like both 60's rock and 1700's baroque, and who buy examples of both types of music.

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.8	
Bob		0.5		0.3
Carol	0.9		1	0.8
David			1	0.4

Collaborative Filtering: Pros and Cons

- + Works for any kind of item
 - No feature selection (e.g., text information) needed
- - Cold start
 - Need enough users in the system to find a match
- - Sparsity
 - The user/ratings matrix is sparse
 - Hard to find users that have rated the same items
- - First rater
 - Cannot recommend an item that has not been previously rated
 - New items & esoteric items
- - Popularity bias
 - Cannot recommend items to someone with unique taste
 - Tends to recommend popular items

Hybrid Methods

- Implement two or more different recommenders (e.g., **content-based** and **collaborative filtering**) and combine predictions
 - Perhaps using a linear model
 - “*Learning to Recommend*”
- Add **content-based approaches** to **collaborative filtering**
 - Building item profiles to deal with the new item problem
 - Building demographics to deal with the new user problem

Evaluation of Recommender Systems

Utility Matrix

users

CDs

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			2		2
				5	
	2	1			1
	3			3	
1					

Training vs. Testing

users

CDs

Training Set

Test Set

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			?		?
				?	
	2	1			?
	3			?	
1					

Evaluation Metrics

- Compare predictions with *known* ratings. (There are many predictions whose ground truth is *unknown*.)
 - **Root-mean-square error (RMSE)**
 - $\sqrt{\frac{1}{M} \sum_{x,i} (U_{xi} - U_{xi}^*)^2}$ where U_{xi} is predicted, and U_{xi}^* is the true rating of x on i ;
 M is the number of testing samples
 - **Precision@10**
 - $\frac{\text{Among the top 10 items with known ratings, how many are relevant (e.g., } \geq 4 \text{ stars)}}{10}$
 - **nDCG@10**
 - **Recall@10**
 - $\frac{\text{Among the top 10 items with known ratings, how many are relevant (e.g., } \geq 4 \text{ stars)}}{\text{Among ALL items with known ratings, how many are relevant (e.g., } \geq 4 \text{ stars)}}$

Additional Comments


- **Problem with RMSE:** In practice, our main interest is in accurately predicting **high** ratings.
 - It is far more important to predict whether you would give 5 stars or 4 stars to a CD you might like than to predict whether you would give 2 stars or 1 star to one you dislike.
 - However, RMSE may penalize methods that perform well on high ratings but poorly on others.
- **Tip:** Leverage **ALL** the data
 - Don't try to reduce data size in an effort to make fancy algorithms work
 - Simple methods on large data do best
 - More data beats better algorithms:
<http://anand.typepad.com/datawocky/2008/03/more-data-usual.html>

The Netflix Prize



- Training data
 - 100 million ratings, 480,000 users, 17,770 movies
 - 6 years of data: 2000-2005
- Test data
 - Last few ratings of each user (2.8 million)
 - Evaluation criterion: **RMSE**
 - $\sqrt{\frac{1}{M} \sum_{x,i} (U_{xi} - U_{xi}^*)^2}$ where U_{xi} is predicted, and U_{xi}^* is the true rating of x on i ; M is the number of testing samples
 - Netflix's system RMSE: 0.9514
- Competition
 - 2,700+ teams
 - \$1 million prize for 10% improvement on Netflix


The Netflix Prize



RMSE:

$$\sqrt{\frac{1}{M} \sum_{x,i} (U_{xi} - U_{xi}^*)^2}$$

Performance of Various Models



Recap: Modeling Deviations

- Basic Collaborative Filtering:

$$U_{xi} = \frac{\sum_{j \in \mathcal{N}} \text{sim}(i, j) \cdot U_{xj}}{\sum_{j \in \mathcal{N}} \text{sim}(i, j)}$$

- In practice,

$$U_{xi} = b_{xi} + \frac{\sum_{j \in \mathcal{N}} \text{sim}(i, j) \cdot (U_{xj} - b_{xj})}{\sum_{j \in \mathcal{N}} \text{sim}(i, j)}$$

- b_{xi} : baseline estimate for U_{xi} ($b_{xi} = \mu + b_x + b_i$)
- μ : overall mean movie rating
- b_x : rating deviation of user x , which is the (avg. rating given by user x) $- \mu$
- b_i : rating deviation of item i , which is the (avg. rating given to item i) $- \mu$

One Step Further: Learning the Weight

$$U_{xi} = b_{xi} + \sum_{j \in \mathcal{N}} w_{ij} (U_{xj} - b_{xj})$$

- w_{ij} is learned from training data
 - We allow $\sum_{j \in \mathcal{N}} w_{ij} \neq 1$.
- w_{ij} models the interaction between pairs of movies.
 - It does not depend on user x .
- What is the objective?
 - **RMSE!** $\sqrt{\frac{1}{M} \sum_{x,i} (U_{xi} - U_{xi}^*)^2}$
 - Or equivalently: $\sum_{x,i} (U_{xi} - U_{xi}^*)^2$

Recommendations via Optimization

$$J(\mathbf{w}) = \sum_{x,i} (U_{xi} - U_{xi}^*)^2 = \sum_{x,i} \left(\left[b_{xi} + \sum_{j \in \mathcal{N}} w_{ij} (U_{xj} - b_{xj}) \right] - U_{xi}^* \right)^2$$

- How to find the values of w_{ij} ?

- Gradient descent!


$$\frac{\partial J}{\partial w_{ij}} = 2 \sum_{x,i} \left(\left[b_{xi} + \sum_{j \in \mathcal{N}} w_{ij} (U_{xj} - b_{xj}) \right] - U_{xi}^* \right) (U_{xj} - b_{xj})$$

(for $j \in \mathcal{N}$)

$$\frac{\partial J}{\partial w_{ij}} = 0$$


(for $j \notin \mathcal{N}$)

Performance of Various Models



Matrix Factorization (Latent-Factor Models)


There are certain latent factors that influence users' ratings.



Latent-Factor Models

- $U \approx QP^T$

The number of factors is small.
In other words, Q and P are “thin”.



- For now, let's assume this is mathematically doable.
 - U has missing entries but let's first ignore that!
 - Basically, we will want the reconstruction error to be small on known ratings and we don't care about the values on the missing ones.

How to interpret Q and P ?

Diagram illustrating the matrix multiplication $Q \times P^T$.

Matrix Q (items × factors):

.1	-.4	.2
-.5	.6	.5
-.2	.3	.5
1.1	2.1	.3
-.7	2.1	-.2
-1	.7	.3

Matrix P^T (factors × users):

1.1	-.2	.3	.5	-.2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

The multiplication is represented as $Q \times P^T$.

- Let's assume that the first factor is *the level of seriousness*.

How to interpret Q and P ?

Diagram illustrating the matrix multiplication $Q \times P^T$.

Matrix Q (items × factors):

	factors		
items	.1	-.4	.2
	-.5	.6	.5
	-.2	.3	.5
	1.1	2.1	.3
	-.7	2.1	-.2
	-1	.7	.3

Matrix P^T (factors × users):

factors		users										
	1.1	-.2	.3	.5	-.2	-.5	.8	-.4	.3	1.4	2.4	-.9
	-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
	2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

$Q \times P^T$

- Let's assume that the first factor is *the level of seriousness*.
 - The seriousness of **User 1** is 1.1
 - The seriousness of **Movie 4** is 1.1
 - So, **User 1** may like **Movie 4**

How to interpret Q and P ?

Diagram illustrating the matrix multiplication $Q \times P^T$.

Matrix Q (items × factors):

	factors		
items	.1	-.4	.2
	-.5	.6	.5
	-.2	.3	.5
	1.1	2.1	.3
	-.7	2.1	-2
	-1	.7	.3

Matrix P^T (factors × users):


	users											
factors	1.1	-.2	.3	.5	-2	-.5	.8	-.4	.3	1.4	2.4	-.9
	-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
	2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

$Q \times P^T$

- Let's assume that the first factor is *the level of seriousness*.
 - The seriousness of **User 1** is 1.1
 - The seriousness of **Movie 5** is -0.7
 - So, **User 1** may NOT like **Movie 5**

Of course, we need to consider all the factors.


- Ratings as “sum of products”



$$U_{xi} = \sum_{\phi: \text{all factors}} Q_{i\phi} \cdot P_{x\phi}$$

Estimating the Missing Rating

- Ratings as “sum of products”




$$U_{xi} = \sum_{\phi: \text{all factors}} Q_{i\phi} \cdot P_{x\phi} = (-0.5) \times (-2) + 0.6 \times 0.3 + 0.5 \times 2.4 = 2.38$$

How to find Q and P ?

Singular Value Decomposition (SVD)

$$A \approx U \Sigma V^T$$


- Input matrix: A
- **Step 1:** Compute $A^T A$
- **Step 2:** Find the eigenvalues of eigenvectors of $A^T A$
 - Eigenvalues $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_n \geq 0$
 - Eigenvectors $v_1, v_2, \dots, v_n$
- **Step 3:** Consider the largest k eigenvalues and their corresponding eigenvectors only. (The choice of k depends on how closely you wish to approximate)
 - $V = [v_1, v_2, \dots, v_k]$



Singular Value Decomposition (SVD)

$$A \approx U \Sigma V^T$$

- **Step 2:** Find the eigenvalues of eigenvectors of $A^T A$
 - Eigenvalues $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_n \geq 0$
 - Eigenvectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$
- **Step 3:** Consider the largest k eigenvalues and their corresponding eigenvectors only.
 - $V = [\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_k]$
 - $\Sigma = \text{diag}\{\sqrt{\lambda_1}, \sqrt{\lambda_2}, \dots, \sqrt{\lambda_k}\}$
- **Step 4:** $U = AV\Sigma^{-1}$
 - Or you can do Steps 1-3 again for AA^T (rather than $A^T A$) to get U



SVD is good, but ...

- SVD gives the **minimum reconstruction error** if we know all entries in A .

$$\min_{U, \Sigma, V} \sum_{i,j} (A_{ij} - [U\Sigma V^T]_{ij})^2$$

- Exactly our objective!
- Using SVD for our matrix factorization task?


Latent Factor Model	User-Item Matrix: U	User-Factor Matrix: Q	Factor-Item Matrix: P^T
SVD	Input Matrix: A	U	ΣV^T

- **BUT, our user-item matrix U has missing values!**
 - How to interpret missing values? (as 0? a bad idea)
 - Does the property of **minimum reconstruction error** still hold if there are missing values? (we don't know)

Factorizing a Matrix with Missing Values

$$\min_{\mathbf{Q}, \mathbf{P}} \sum_{(x,i) \text{ known}} (U_{xi} - [\mathbf{Q}\mathbf{P}^T]_{xi})^2 = \sum_{(x,i) \text{ known}} (U_{xi} - \mathbf{q}_i \mathbf{p}_x^T)^2$$

- $\mathbf{q}_i$ (item vector): the row corresponding to item i in $\mathbf{Q}$
- $\mathbf{p}_x^T$ (user vector): the column corresponding to user x in $\mathbf{P}^T$



Overfitting

$$\min_{\mathbf{Q}, \mathbf{P}} \sum_{(x,i) \text{ known}} (U_{xi} - \mathbf{q}_i \mathbf{p}_x^T)^2$$

- $\mathbf{q}_i$ (item vector): the row corresponding to item i in $\mathbf{Q}$
- $\mathbf{p}_x^T$ (user vector): the column corresponding to user x in $\mathbf{P}^T$
- No closed form solution.
- All item vectors and user vectors are parameters to be learned!
- **Overfitting**: With too much freedom (too many free parameters) the model starts fitting noise in the training data, thus not generalizing well to unseen test data.

1	3	4			
	3	5			5
		4	5		5
		3			
		3			
2			?	?	?
				?	
	2	1			?
	3			?	
1					

Regularization

- Model parameters can be “complicated” where there are sufficient training data
- Model parameters should be “simple” where training data are scarce

$$\min_{\mathbf{Q}, \mathbf{P}} \sum_{(x,i) \text{ known}} (U_{xi} - \mathbf{q}_i \mathbf{p}_x^T)^2 + \left[c_1 \sum_x \|\mathbf{p}_x\|^2 + c_2 \sum_i \|\mathbf{q}_i\|^2 \right]$$


Original Objective

Regularization Term

(c_1 and c_2 are hyperparameters)


- How to understand the **Regularization Term**?

The Effect of Regularization




The Effect of Regularization

- If **the user** has rated hundreds of movies, it stays about where the data places it.



The Effect of Regularization

- If **the user** has rated only a handful, it is pulled back towards the origin.



Gradient Descent


$$\min_{\mathbf{Q}, \mathbf{P}} J = \sum_{(x,i) \text{ known}} (U_{xi} - \mathbf{q}_i \mathbf{p}_x^T)^2 + \left[c_1 \sum_x \|\mathbf{p}_x\|^2 + c_2 \sum_i \|\mathbf{q}_i\|^2 \right]$$

- **Step 1:** Initialize $\mathbf{Q}$ and $\mathbf{P}$ using SVD (pretend missing ratings are 0)


- **Step 2:** Gradient descent

- $$P_{x\phi} = P_{x\phi} - \eta \frac{\partial J}{\partial P_{x\phi}}$$
 - $$\frac{\partial J}{\partial P_{x\phi}} = \sum_{(x,i) \text{ known}} (-2(U_{xi} - \mathbf{q}_i \mathbf{p}_x^T) Q_{i\phi} + 2c_1 P_{x\phi})$$
- $$Q_{i\phi} = Q_{i\phi} - \eta \frac{\partial J}{\partial Q_{i\phi}}$$
 - $$\frac{\partial J}{\partial Q_{i\phi}} = \sum_{(x,i) \text{ known}} (-2(U_{xi} - \mathbf{q}_i \mathbf{p}_x^T) P_{x\phi} + 2c_2 Q_{i\phi})$$

Learned Item Vectors in the Latent Factor Space



Performance of Various Models



Extending Latent Factor Models to Include Bias

Bias, Again

- Basic Latent Factor Model:

$$U_{xi} = \mathbf{q}_i \mathbf{p}_x^T$$

- Latent Factor Model with Bias:

$$U_{xi} = \mu + b_x + b_i + \mathbf{q}_i \mathbf{p}_x^T$$

- μ : overall mean movie rating
- b_x : rating deviation of user x
- b_i : rating deviation of item i


user bias



movie bias



user-movie interaction



Bias, Again

- Latent Factor Model with Bias:

$$U_{xi} = \mu + b_x + b_i + \mathbf{q}_i \mathbf{p}_x^T$$

- μ : overall mean movie rating
 - E.g., $\mu = 2.7$
- b_x : rating deviation of user x (to be learned)
 - E.g., Bob is a critical reviewer. Based on the training data, his rating will be 0.7 star lower than the mean $\Rightarrow b_x = -0.7$.
- b_i : rating deviation of item i (to be learned)
 - E.g., Star Wars will get a mean rating of 0.5 higher than the average $\Rightarrow b_i = 0.5$
- $\mathbf{q}_i$ and $\mathbf{p}_x$: vector of user x and item i in the latent factor space (to be learned)
 - E.g., based on the genre, Bob likes Star Wars $\Rightarrow \mathbf{q}_i \mathbf{p}_x^T = 0.3$
- $U_{xi} = 2.7 - 0.7 + 0.5 + 0.3 = 2.8$

Fitting the New Model

$$\min_{Q, P, b_x, b_i} J = \sum_{(x,i) \text{ known}} (U_{xi} - (\mu + b_x + b_i + \mathbf{q}_i \mathbf{p}_x^T))^2 \\ + \left[c_1 \sum_x \|\mathbf{p}_x\|^2 + c_2 \sum_i \|\mathbf{q}_i\|^2 + c_3 \sum_x \|b_x\|^2 + c_4 \sum_i \|b_i\|^2 \right]$$


- Both biases b_x, b_i as well as interactions $\mathbf{q}_i, \mathbf{p}_x$ are treated as parameters to be learned via gradient descent

- $P_{x\phi} = P_{x\phi} - \eta \frac{\partial J}{\partial P_{x\phi}}, \quad Q_{i\phi} = Q_{i\phi} - \eta \frac{\partial J}{\partial Q_{i\phi}}$


- $b_x = b_x - \eta \frac{\partial J}{\partial b_x}, \quad b_i = b_i - \eta \frac{\partial J}{\partial b_i}$

Performance of Various Models

- Which hyperparameter determines the number of parameters?
 - Number of factors



Performance of Various Models



What were the next steps?
Next week!



Thank You!

Course Website: <https://yuzhang-teaching.github.io/CSCE670-F26.html>