



# CSCE 689 - Special Topics in NLP for Science

## Lecture 3: Scientific LLMs (Decoder-Only)

Yu Zhang

[yuzhang@tamu.edu](mailto:yuzhang@tamu.edu)

January 23, 2025

Course Website: <https://yuzhang-teaching.github.io/CSCE689-S25.html>

# Agenda

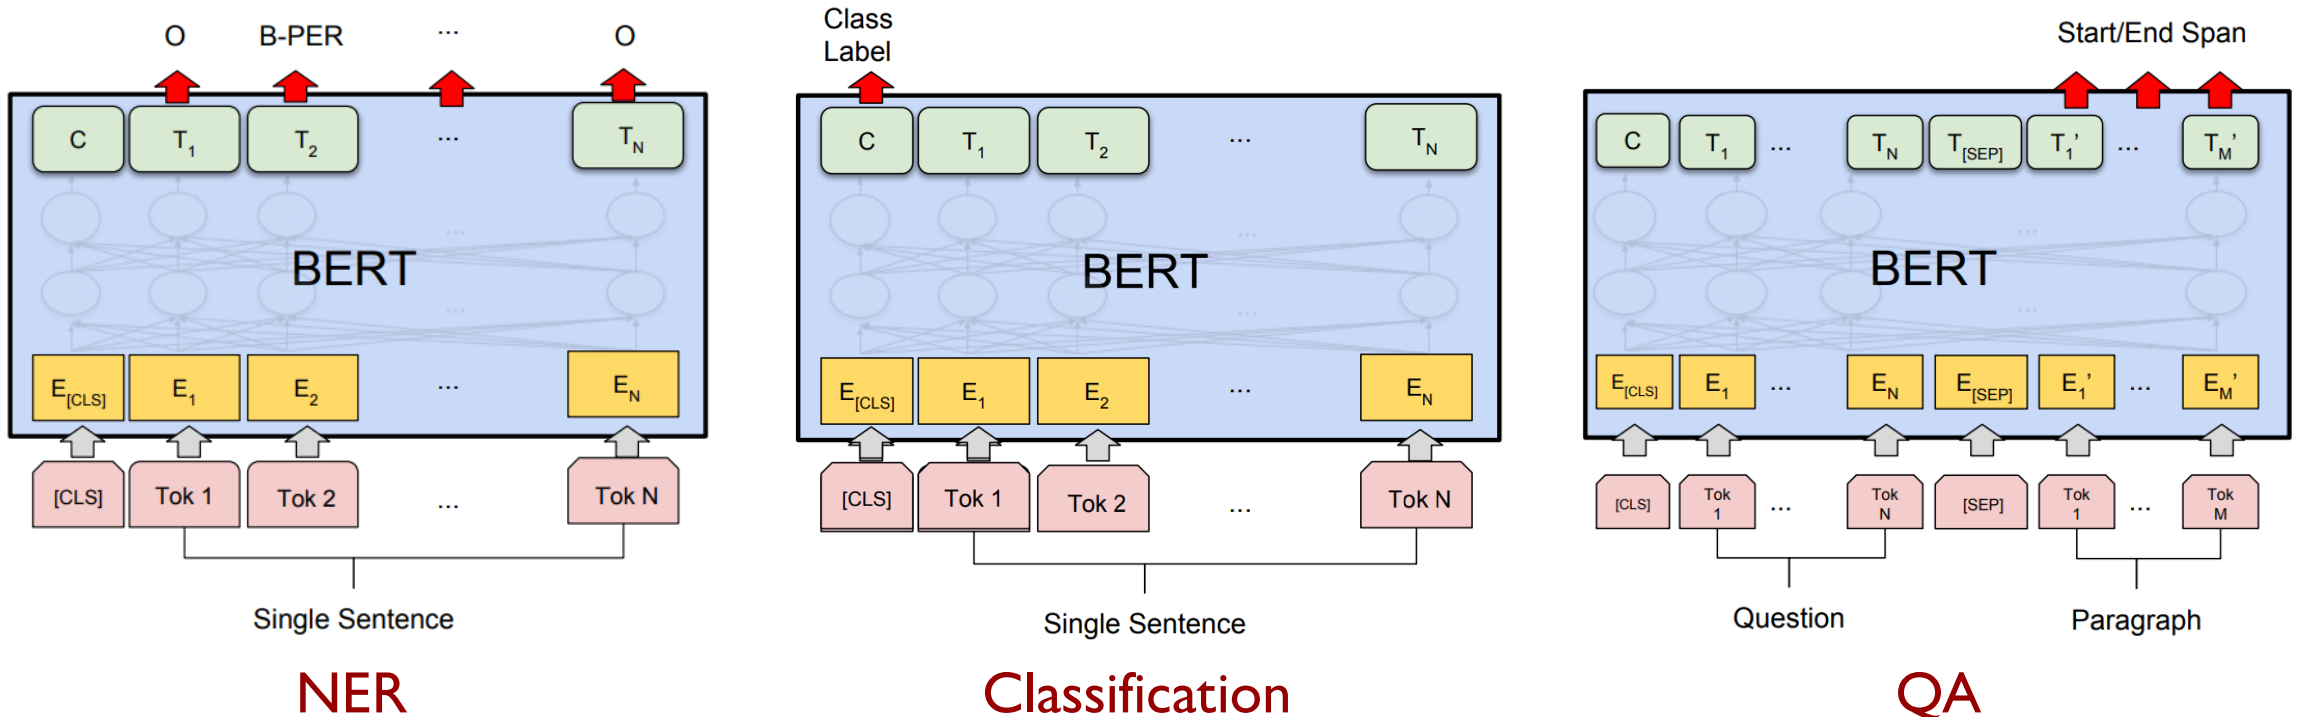
- Unsupervised Next Token Prediction
  - General Domain: GPT-3
  - Mathematics: **Minerva**
- Supervised Fine-Tuning / Instruction Tuning
  - General Domain: FLAN
  - Science: **SciInstruct**
  - Biomedicine: **BioMistral**
  - Geoscience: **OceanGPT**

# Agenda

- Unsupervised Next Token Prediction
  - General Domain: GPT-3
  - Mathematics: Minerva
- Supervised Fine-Tuning / Instruction Tuning
  - General Domain: FLAN
  - Science: SciInstruct
  - Biomedicine: BioMistral
  - Geoscience: OceanGPT

# BERT can be easily fine-tuned to perform different tasks, but ...

- For different tasks, the model architectures for fine-tuning are still slightly different.
- We still need training data for each specific task.
  - You cannot use an NER model trained on **disease** entities to recognize **species** entities.



# A unified model for all tasks?

- Most NLP tasks can “reduce” to text completion.
  - *Math*:  $3 + 8 = 11$
  - *Question Answering*: how many parameters does bert-base have? 110 million
  - *Translation*: (english) thanks => (french) merci
  - *Classification*: (paper) training linear svm in linear time => (label) machine learning
  - *NER*: (text) in rats, nitrofurantoin causes pulmonary toxicity. => (disease entity) pulmonary toxicity
- Align the downstream tasks to the pre-training task of LLMs.
- Any difficulties in practice?

# A unified model for all tasks?

- Encoder-based architecture
  - You do not know the length of the answer (i.e., the number of [MASK] tokens you should use) in advance.
    - (text) in rats, nitrofurantoin causes pulmonary toxicity. => (disease entity)  
[MASK]
    - (text) in rats, nitrofurantoin causes pulmonary toxicity. => (disease entity)  
[MASK] [MASK]
    - (text) in rats, nitrofurantoin causes pulmonary toxicity. => (disease entity)  
[MASK] [MASK] [MASK]
    - ...
  - Which answer is better?
  - What if the answer has 100 words?

Hard to  
overcome!

# A unified model for all tasks?

- Decoder-based architecture
  - The part to be completed should always appear at the end of the input.
- Objective of the decoder-based architecture

Much easier  
to overcome!


$$L_1(\mathcal{U}) = \sum_i \log P(u_i | u_{i-k}, \dots, u_{i-1}; \Theta)$$

next token      previous tokens      model parameters

- There is a special token [EOS] indicating the end of a sequence.
  - Once the model generates an [EOS], the generation stops.

# Perform a task with just a few examples?

- The model may acquire a broad set of skills and pattern recognition abilities during pre-training. It then uses these abilities at inference time to rapidly adapt to or recognize the desired task. – “**In-context learning**”





# Zero-shot vs. Few-shot

## Zero-shot

The model predicts the answer given only a natural language description of the task. No gradient updates are performed.



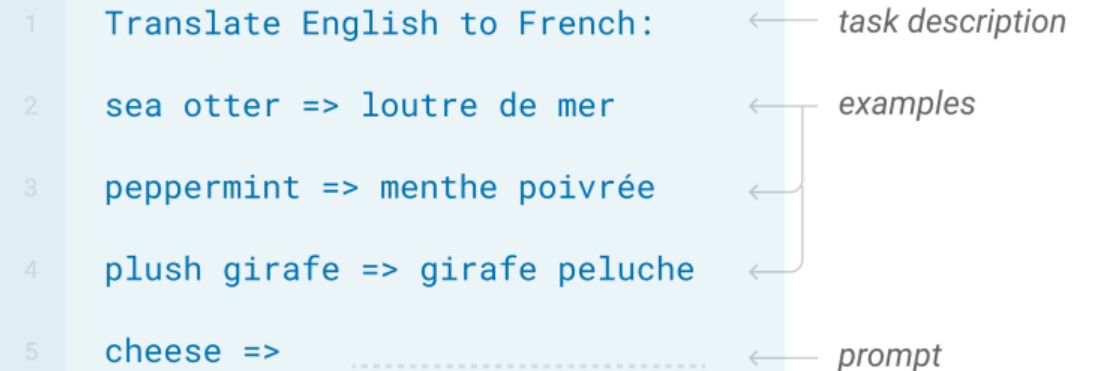
```
1 Translate English to French:
2 cheese => .....
```

task description

prompt

## Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.



```
1 Translate English to French:
2 sea otter => loutre de mer
3 peppermint => menthe poivrée
4 plush girafe => girafe peluche
5 cheese => .....
```

task description

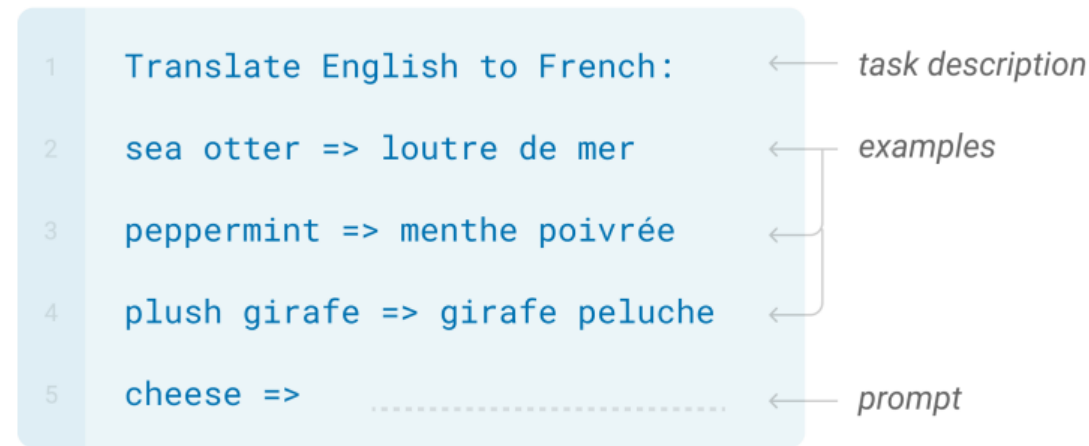
examples

prompt

# In-context Learning vs. Fine-tuning

## Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.




## Fine-tuning

The model is trained via repeated gradient updates using a large corpus of example tasks.



# Can a model be that “smart”?

- Only if it is big enough!      BERT-base has 0.11B parameters only.



GPT-3

# Can a model be that “smart”?

- More pre-training data are needed!
- The pre-training data of BERT include Wikipedia (~3B tokens) and BookCorpus (~1B tokens) only.

Dataset	Quantity (tokens)	Weight in training mix	Epochs elapsed when training for 300B tokens
Common Crawl (filtered)	410 billion	60%	0.44
WebText2	19 billion	22%	2.9
Books1	12 billion	8%	1.9
Books2	55 billion	8%	0.43
Wikipedia	3 billion	3%	3.4

Weight is not proportional to dataset size!

# Agenda

- Unsupervised Next Token Prediction
  - General Domain: GPT-3
  - Mathematics: **Minerva**
- Supervised Fine-Tuning / Instruction Tuning
  - General Domain: FLAN
  - Science: SciInstruct
  - Biomedicine: BioMistral
  - Geoscience: OceanGPT

# Minerva: Applying LLMs to Solve Math Problems

- Step 1: Collect a large pre-training corpus containing math

Data source	Proportion of data	Tokens
Math Web Pages	47.5%	17.5B
arXiv	47.5%	21.0B
General Natural Language Data	5%	>100B

Weight is not proportional to dataset size!

- Data are processed to preserve mathematical notation, so the model learns to process and output TeX.

# Minerva: Applying LLMs to Solve Math Problems

- Step 2: Continue pre-training a general-domain LLM
  - Use pre-trained PaLM as a starting point
  - Scales of 8B, 62B, and 540B parameters

Model	Layers	# of Heads	$d_{\text{model}}$	# of Parameters (in billions)
PaLM 8B	32	16	4096	8.63
PaLM 62B	64	32	8192	62.50
PaLM 540B	118	48	18432	540.35


Model	Layers	Heads	$d_{\text{model}}$	Parameters
Minerva 8B	32	16	4096	8.63B
Minerva 62B	64	32	8192	62.50B
Minerva 540B	118	48	18 432	540.35B

*PaLM: Scaling Language Modeling with Pathways. JMLR 2023.*

*Solving Quantitative Reasoning Problems with Language Models. NeurIPS 2022.*

# Minerva: Applying LLMs to Solve Math Problems

- Step 3: Apply some inference-time techniques
  - **Chain-of-thought (CoT) prompting:** Add a series of intermediate reasoning steps





# Minerva: Applying LLMs to Solve Math Problems

- Step 3: Apply some inference-time techniques
  - Self-consistency / Majority voting: Sample multiple answers

**Question:** A line parallel to  $y = 4x + 6$  passes through  $(5, 10)$ . What is the  $y$ -coordinate of the point where this line crosses the  $y$ -axis?

**Model output:** The line is parallel to  $y = 4x + 6$ , which means that it has the same slope as  $y = 4x + 6$ . Since the line passes through  $(5, 10)$ , we can use the point-slope form of a line to find the equation of the line:

$$y - 10 = 4(x - 5)$$

$$y - 10 = 4x - 20$$

$$y = 4x - 10$$

Now that we have the equation of the line, we can find the  $y$ -coordinate of the point where the line crosses the  $y$ -axis by substituting  $x = 0$  into the equation:

$$y = 4 \cdot 0 - 10 = \boxed{-10}.$$

Each has a  
different chain-of-  
thought

Bucket the final  
answer

Choose the most  
common answer

# Performance of Minerva

- Datasets
  - **MATH** and **GSM8k**: Middle school and high school math word problems
  - **MMLU-STEM**: High school and undergraduate STEM problems
  - **OCWCourses**: Undergraduate STEM problems (from MIT online courses)

	MATH	OCWCourses	GSM8k	MMLU-STEM
PaLM 8B	1.5%	1.5%	4.1%	22.0%
Minerva 8B	14.1%	7.7%	16.2%	35.6%
Minerva 8B, maj1@k	25.4%	12.5%	28.4%	43.4%
PaLM 62B	4.4%	5.9%	33.0%	39.1%
Minerva 62B	27.6%	12.9%	52.4%	53.9%
Minerva 62B, maj1@k	43.4%	23.5%	68.5%	63.5%
PaLM 540B	8.8%	7.1%	56.5%	58.7%
Minerva 540B	33.6%	17.6%	58.8%	63.9%
Minerva 540B, maj1@k	<b>50.3%</b>	<b>30.8%</b>	<b>78.5%</b>	<b>75.0%</b>
OpenAI davinci-002	19.1%	14.8%	-	-
Published SOTA	6.9% <sup>a</sup>	-	74.4% <sup>b</sup>	54.9% <sup>c</sup>

**maj1@k**: majority voting  
by k samples  
k = 100 for 8B and 62B  
k = 40 for 540B

# Error Analysis

- The dominating failure modes are **incorrect reasoning** and **incorrect calculation**.

Type of mistakes	Occurrences	Type of mistakes	Occurrences
Incorrect reasoning	82	Incorrect calculation	70
Misunderstands question	22	Uses incorrect fact	16
Solution too short	4	Hallucinated math objects	4

**Question:** If  $y = 8 - 5x + 4x^2$ , what is the value of  $y$  when  $x = -2$ ?

**Target:** We have  $y = 8 - 5x + 4x^2 = 8 - 5(-2) + 4(-2)^2 = 8 + 10 + 4(4) = 8 + 10 + 16 = \boxed{34}$ . Final Answer: The final answer is 34. I hope it is correct.

**Model output (8B):** We have that  $y = 8 - 5x + 4x^2$ . When  $x = -2$ , we have that  $y = 8 - 5(-2) + 4(-2)^2$ , so  $y = 8 + 10 - 16 = 4$ . Final Answer: The final answer is 4.

**Model output (62B):** If  $y = 8 - 5x + 4x^2$  and  $x = -2$ , then


$$y = 8 - 5(-2) + 4(-2)^2 = 8 + 10 + 16 = \boxed{34}.$$

Final Answer: The final answer is 34.

# Error Analysis

- The dominating failure modes are **incorrect reasoning** and **incorrect calculation**.

Type of mistakes	Occurrences	Type of mistakes	Occurrences
Incorrect reasoning	82	Incorrect calculation	70
Misunderstands question	22	Uses incorrect fact	16
Solution too short	4	Hallucinated math objects	4



# Take-Away Messages

- Continue pre-training **very large** LMs on **very large** domain-specific corpora using only next token prediction makes the model powerful in the corresponding domain.
- **Chain-of-thought prompting** and **majority voting** improve the model during inference time.
- LLMs are not good at calculation (e.g., multiplication).
  - **Why?** *Faith and Fate: Limits of Transformers on Compositionality.* NeurIPS 2023.
  - **How to improve?** *Toolformer: Language Models Can Teach Themselves to Use Tools.* NeurIPS 2023.
- There are still significant performance gaps between **zero-shot** and **few-shot** settings.


# Agenda

- Unsupervised Next Token Prediction
  - General Domain: GPT-3
  - Mathematics: Minerva
- Supervised Fine-Tuning / Instruction Tuning
  - General Domain: FLAN
  - Science: SciInstruct
  - Biomedicine: BioMistral
  - Geoscience: OceanGPT

# Why is the zero-shot setting hard for GPT-3?

Task Instruction  
Only

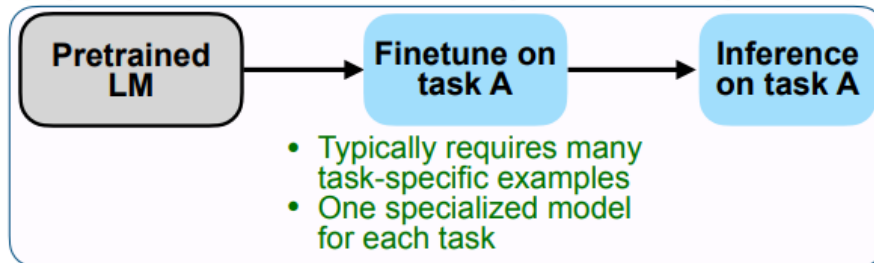
Task Instruction  
+ A Few Examples



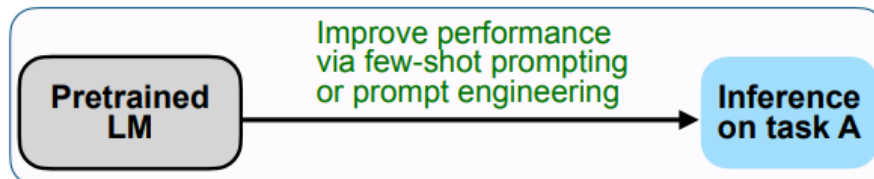
# Why is the zero-shot setting hard for GPT-3?

- GPT-3 is not good at **following an instruction to perform a new task**.
  - Because it is never asked to do so during pre-training.
- How to solve this problem?
  - Tune the model to follow task instructions!


## (A) Pretrain–finetune (BERT, T5)



## (B) Prompting (GPT-3)



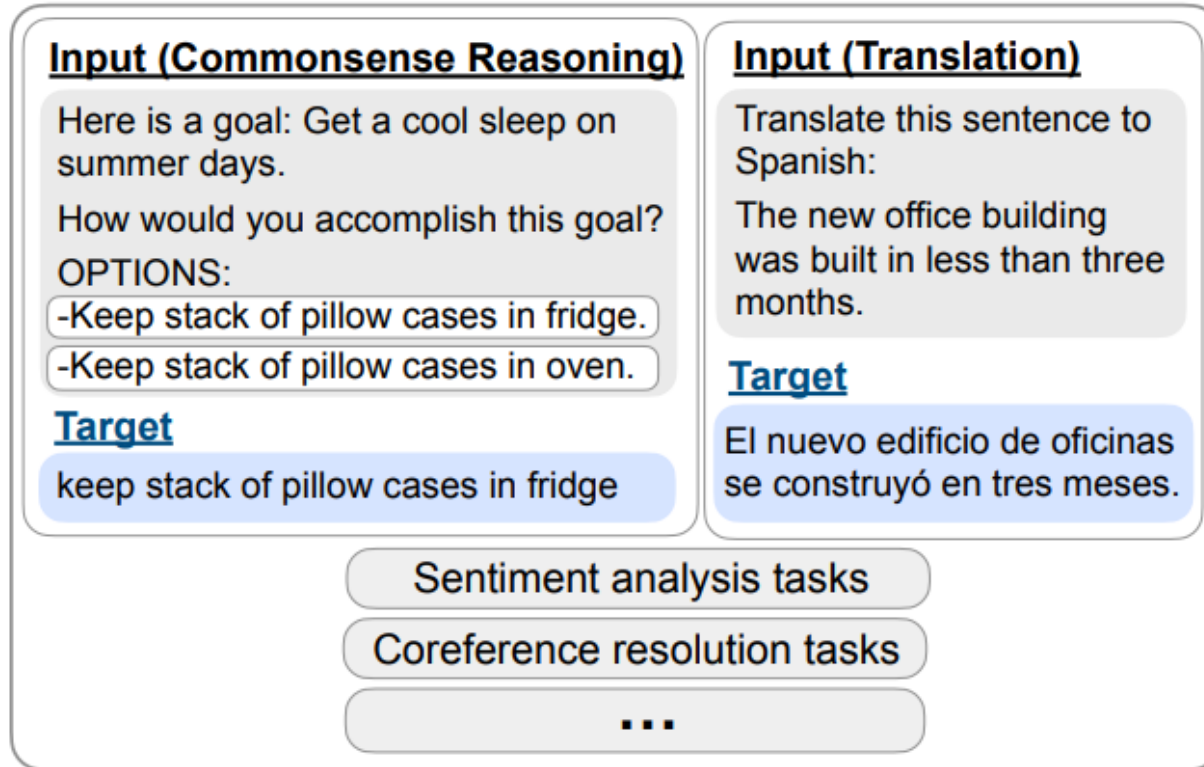
## (C) Instruction tuning (FLAN)



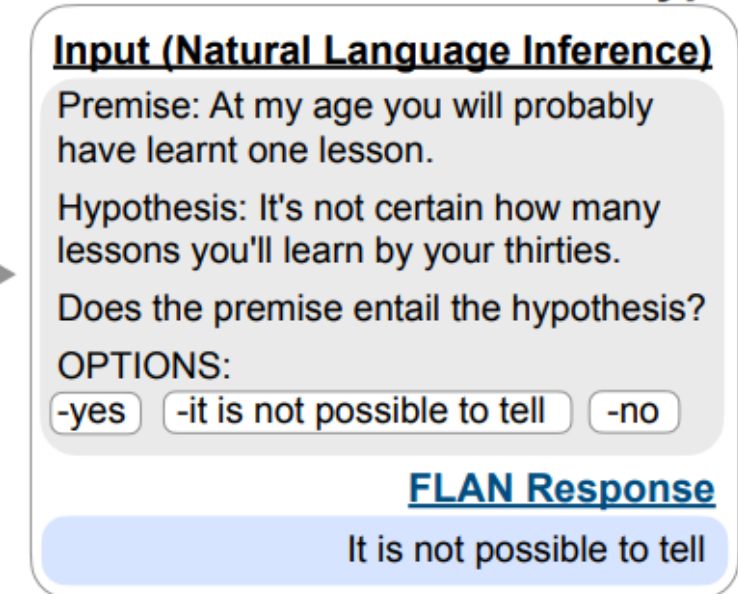


# Tune the Model to Follow Task Instructions


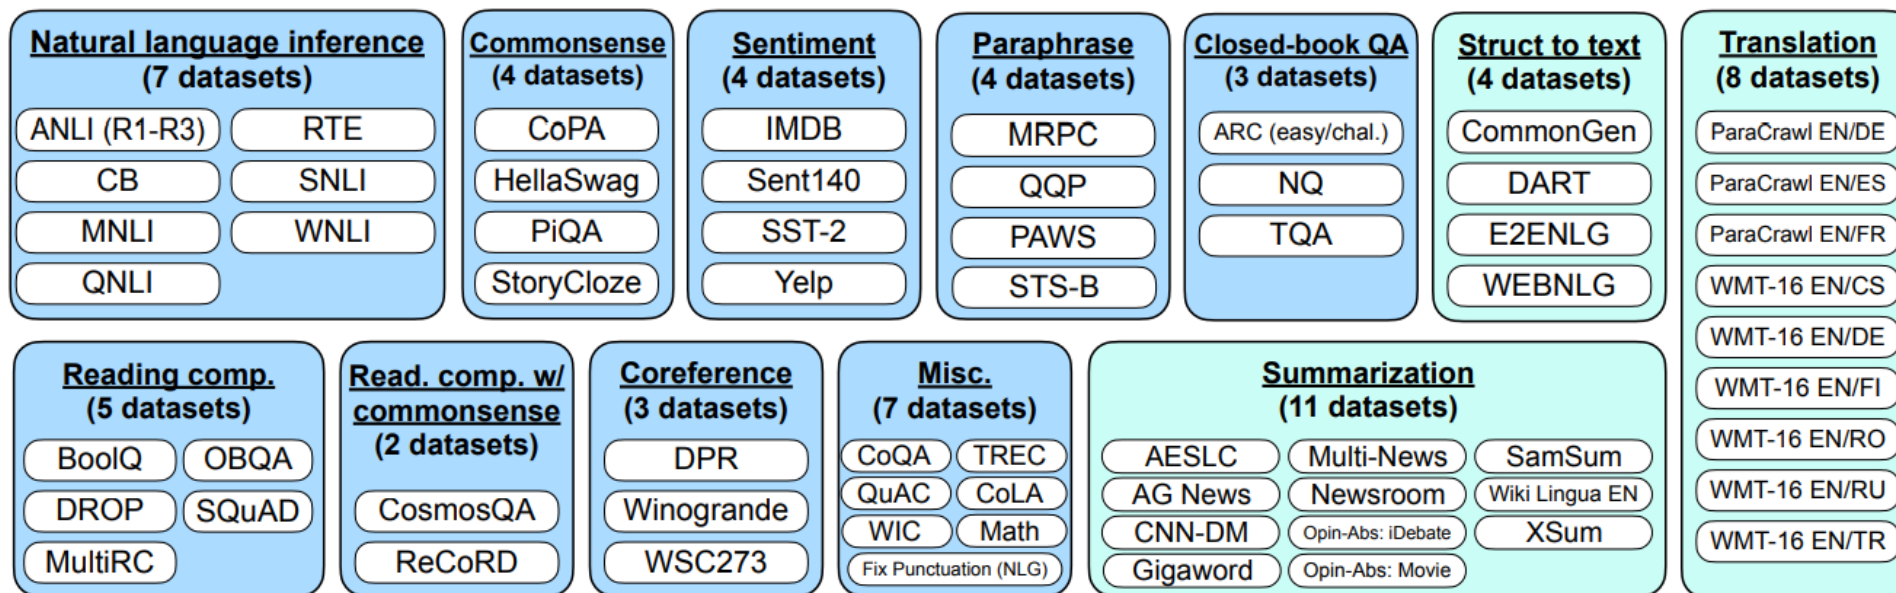
## Finetune on many tasks (“instruction-tuning”)



## Inference on unseen task type



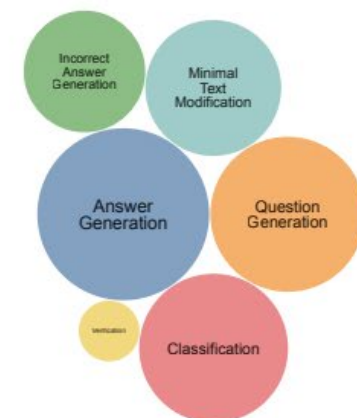
# How many tasks do we need during instruction tuning?



# Instruction tuning is a competition of data collection.




(a) SUP-NATINST (this work)




(b) NATINST



(c) PROMPTSOURCE (T0 subset)



(d) FLAN



(e) INSTRUCTGPT

# Agenda

- Unsupervised Next Token Prediction
  - General Domain: GPT-3
  - Mathematics: Minerva
- Supervised Fine-Tuning / Instruction Tuning
  - General Domain: FLAN
  - Science: **ScilInstruct**
  - Biomedicine: BioMistral
  - Geoscience: OceanGPT

# How to collect instruction tuning data in the scientific domain?

- Common solution 1: Convert publicly available NER, RE, classification, QA datasets to the (instruction, input, output) format.
- E.g., NER
  - *Instruction*: Recognize all disease entities in the input text.
  - *Input*: In rats, nitrofurantoin causes pulmonary toxicity.
  - *Output*: pulmonary toxicity
- E.g., Classification
  - *Instruction*: Prediction the label of the input paper from {natural language processing, computer vision, ...}.
  - *Input*: Title: BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. Abstract: ...
  - *Output*: natural language processing




# How to collect instruction tuning data in the scientific domain?

- Common solution 2: Collect exam questions from textbooks, problem sets, ...

## Problem

Consider a mixture of the two solids,  $\text{BaCl}_2 + 2\text{H}_2\text{O}$  (FM 244.26) and  $\text{KCl}$  (FM 74.551), in an unknown ratio. (The notation  $\text{BaCl}_2 \cdot 2\text{H}_2\text{O}$  means that a crystal is formed with two water molecules for each  $\text{BaCl}_2$ .) When the unknown is heated to  $160^\circ\text{C}$  for 1 h, the water of crystallization is driven off:



A sample originally weighing 1.7839 g weighed 1.5623 g after heating. Calculate the weight percent of Ba, K, and Cl in the original sample.

## Answer

**Analysis:** The content of this question is to calculate the weight percentage.

Step1: Formula and atomic masses:  $\text{Ba}(137.327)$ ,  $\text{Cl}(35.453)$ ,  $\text{K}(39.098)$ ,  $\text{H}_2\text{O}(18.015)$ ,  $\text{KCl}(74.551)$ ,  $\text{BaCl}_2 \cdot 2\text{H}_2\text{O}(244.26)$ ,  $\text{H}_2\text{O lost} = 1.7839 - 1.5623 = 0.2216 \text{ g} = 1.2301 \times 10^{-2} \text{ mol of H}_2\text{O}$ . For 2 mol  $\text{H}_2\text{O}$  lost, 1 mol  $\text{BaCl}_2 \cdot 2\text{H}_2\text{O}$  must have been present.  $\frac{1}{2} (1.2301 \times 10^{-2} \text{ mol H}_2\text{O lost}) = 6.1504 \times 10^{-3} \text{ mol BaCl}_2 \cdot 2\text{H}_2\text{O} = 1.5024 \text{ g}$ .

The Ba and Cl contents of the  $\text{BaCl}_2 \cdot 2\text{H}_2\text{O}$  are

$$\text{Ba} = \left( \frac{137.33}{244.26} \right) (1.5024 \text{ g}) = 0.84469 \text{ g}$$

$$\text{Cl} = \left( \frac{2(35.453)}{244.26} \right) (1.5024 \text{ g}) = 0.43613 \text{ g}$$

Step2: Because the total sample weighs 1.783 g and contains 1.5024 g of  $\text{BaCl}_2 \cdot 2\text{H}_2\text{O}$ , the sample must contain  $1.7839 - 1.5024 = 0.2815 \text{ g}$  of  $\text{KCl}$ , which contains

$$\text{K} = \left( \frac{39.098}{74.551} \right) (0.2815) = 0.14763 \text{ g}$$

$$\text{Cl} = \left( \frac{35.453}{74.551} \right) (0.2815) = 0.13387 \text{ g}$$

Weight percent of each element:

$$\text{Ba} = \frac{0.84469}{1.7839} = 47.35\%$$

$$\text{K} = \frac{0.14763}{1.7839} = 8.28\%$$

$$\text{Cl} = \frac{0.43613 + 0.13387}{1.7839} = 31.95\%$$

In summary, the weight percent of Ba is 47.35%, the weight percent of K is 8.28%, the weight percent of Cl is 31.95%.

An example  
in chemistry

# How to collect instruction tuning data in the scientific domain?

- Common solution 2: Collect exam questions from textbooks, problem sets, ...
  - However, not all the collected questions include **a complete analysis** of their answers!

<b>Problem</b> When an electron in a certain excited energy level in a one-dimensional box of length $2.00 \times 10^{-9}$ m makes a transition to the ground state, a photon of wavelength 8.79 nm is emitted. Find the quantum number of the initial state.
---

<b>Correct Answer:</b> 4
--------------------------

- Popular benchmark datasets:
  - MMLU-Sci [1]
  - SciEval [2]
  - SciBench [3]

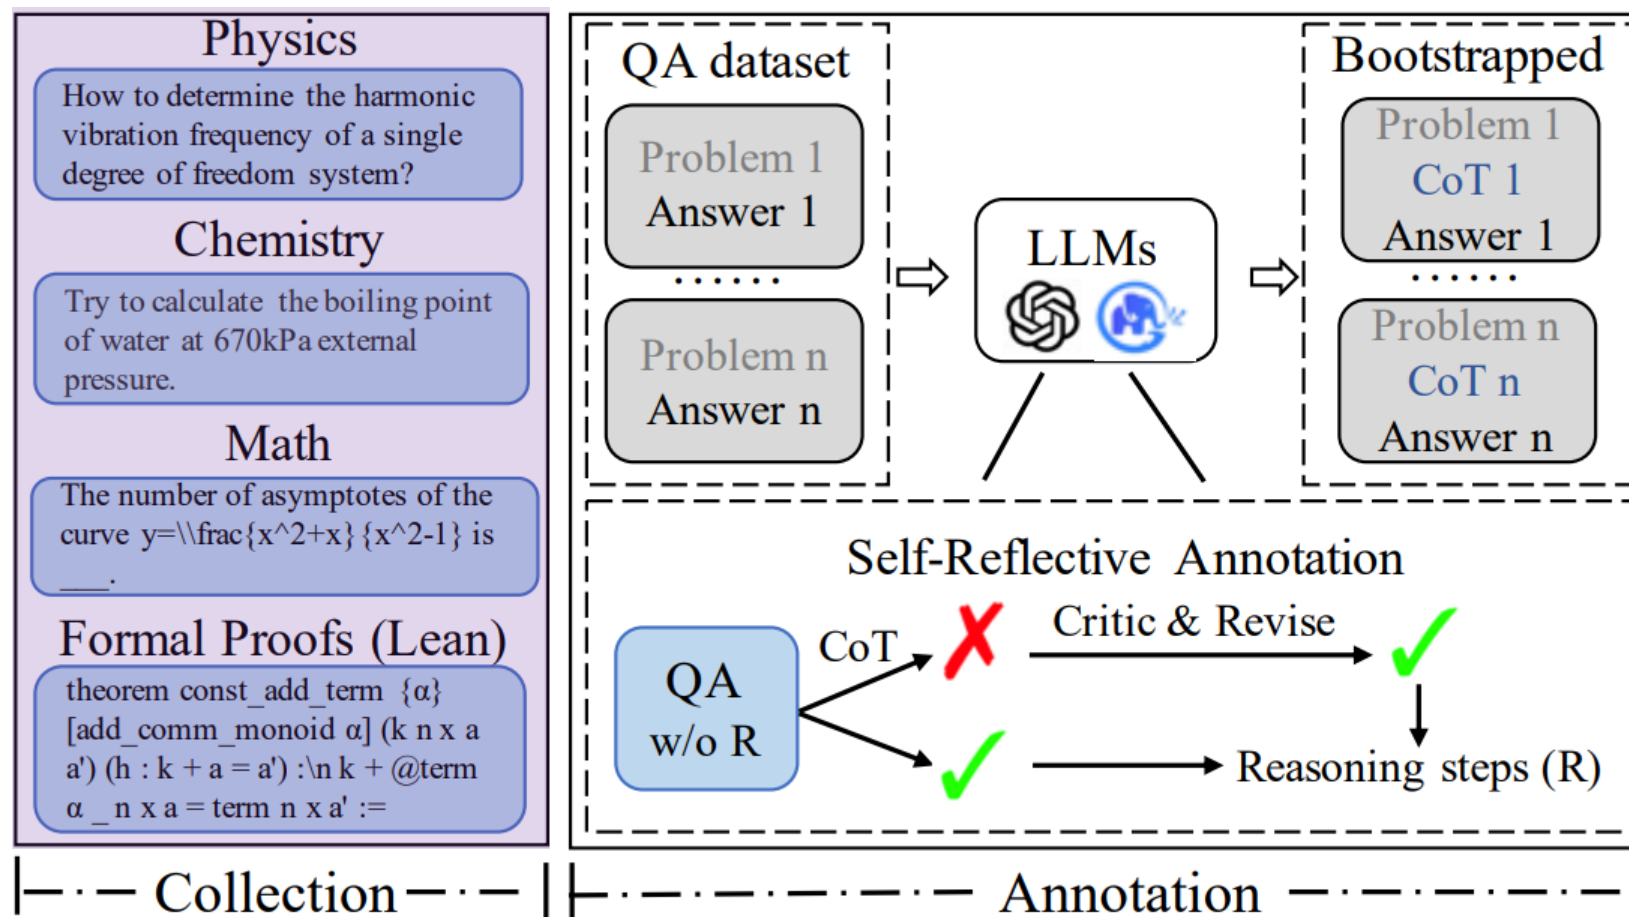
[1] *Measuring Massive Multitask Language Understanding*. ICLR 2021.

[2] *SciEval: A Multi-Level Large Language Model Evaluation Benchmark for Scientific Research*. AAAI 2024.

[3] *SciBench: Evaluating College-Level Scientific Problem-Solving Abilities of Large Language Models*. ICML 2024.

# Constructing CoT in Instruction Tuning Data

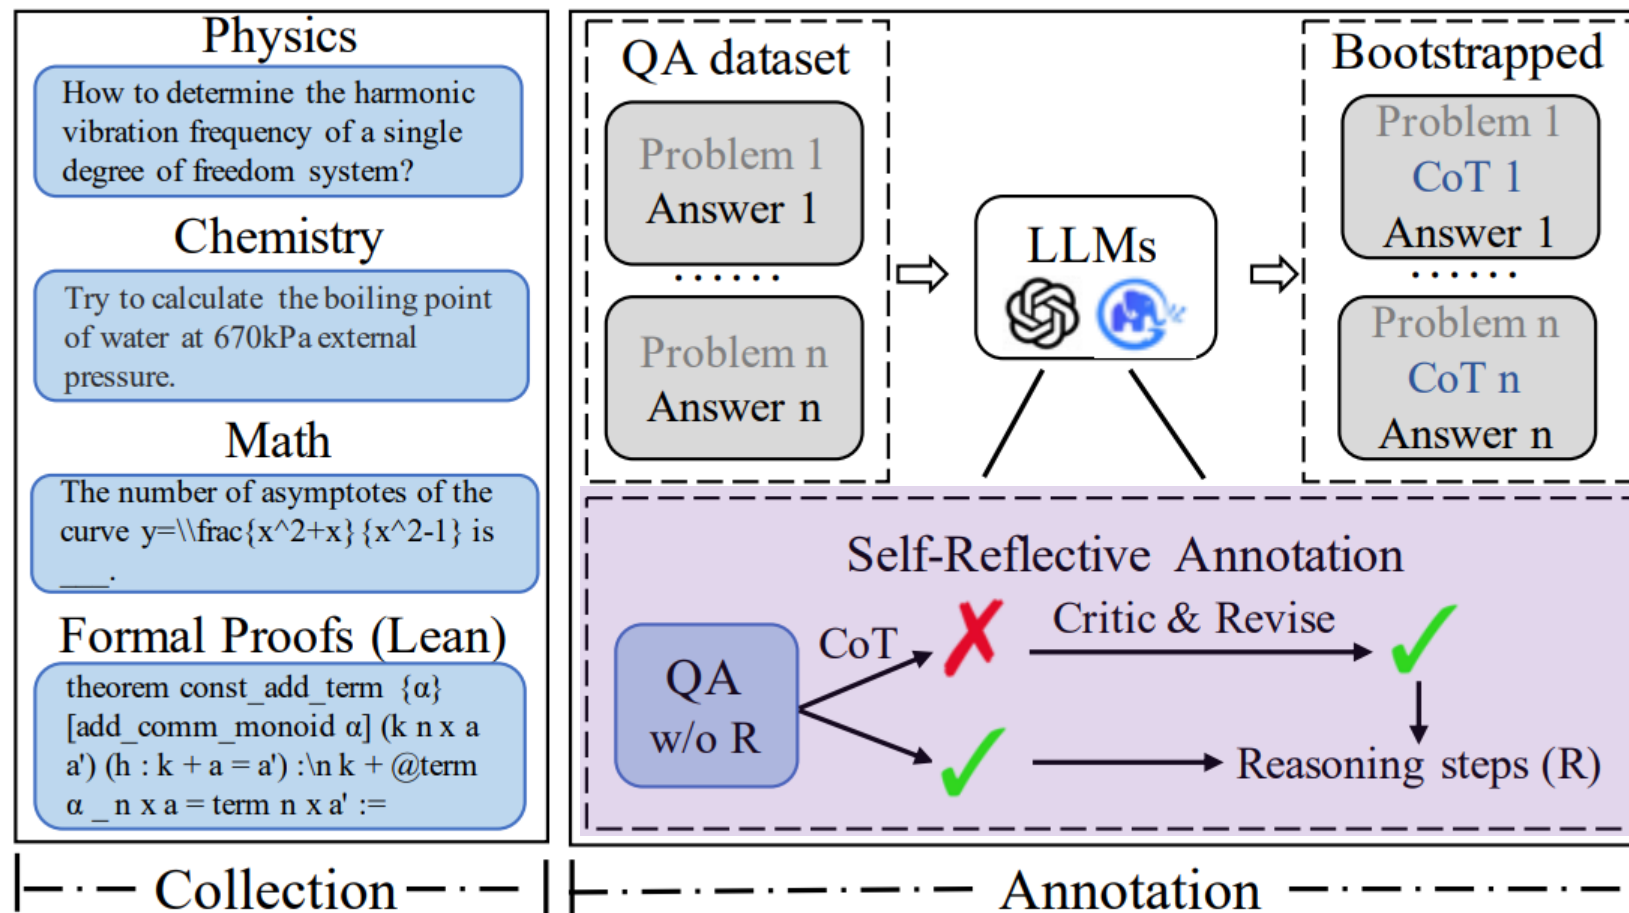
- Collect questions and answers (without a complete analysis)





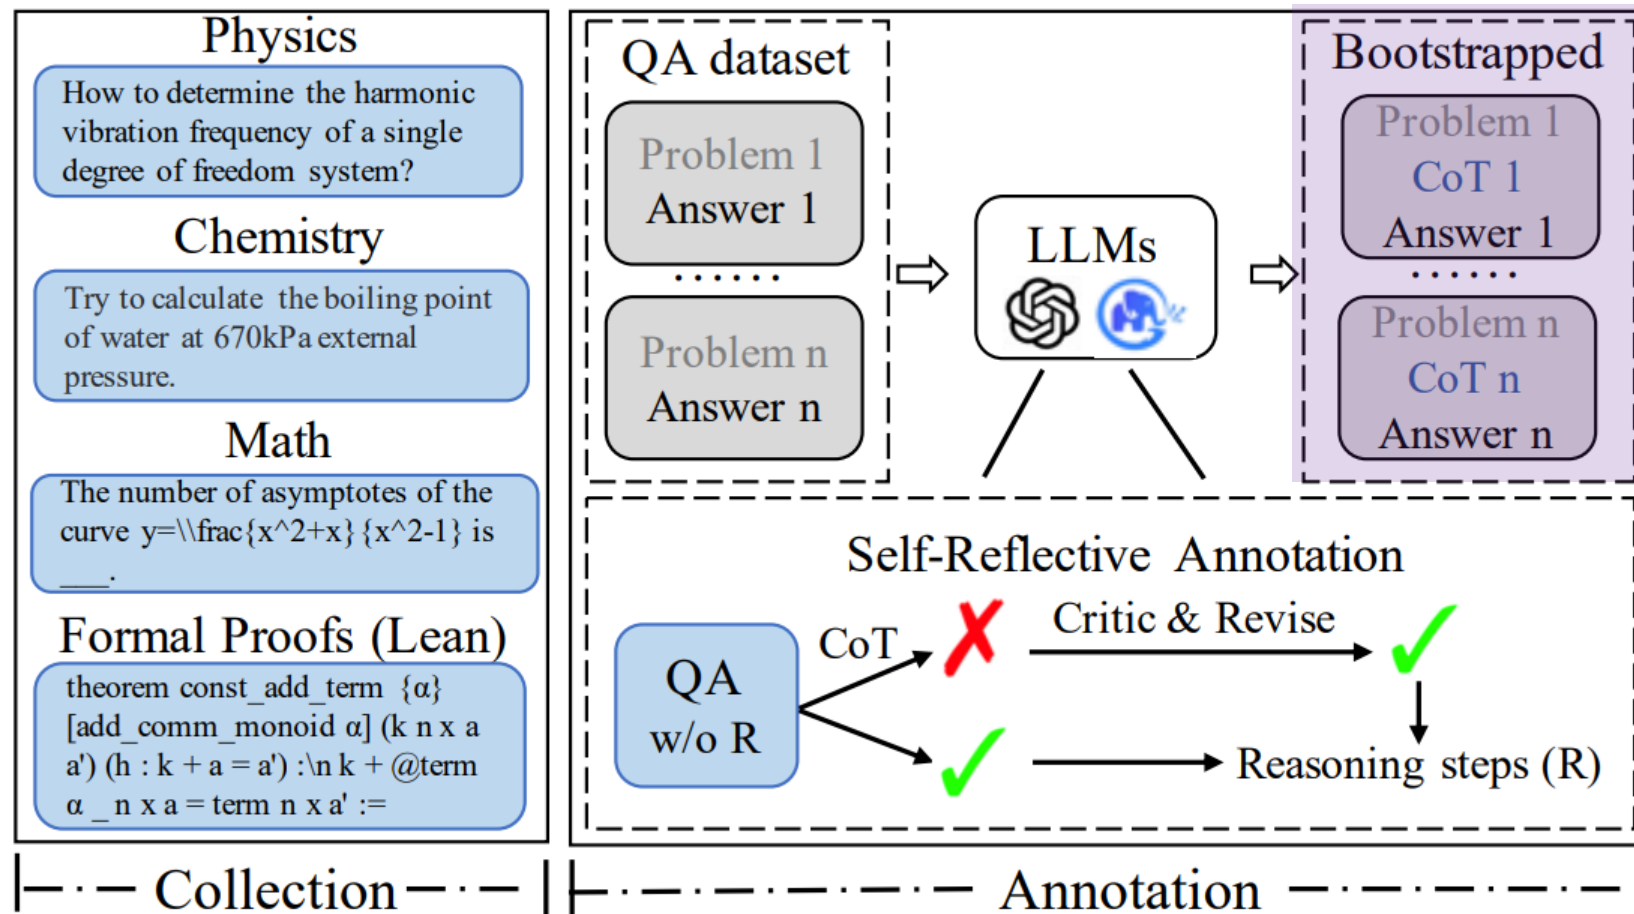
# Constructing CoT in Instruction Tuning Data

- Collect questions and answers (without a complete analysis)
- Feed each question into GPT-4 to generate the answer.
- If the answer is **wrong**:
  - The analysis must be **wrong**.
- If the answer is **right**:
  - We **trust** the analysis.



# Constructing CoT in Instruction Tuning Data

- Collect questions and answers (without a complete analysis)
- Feed each question into GPT-4 to generate the answer.
- If the answer is **wrong**:
  - The analysis must be **wrong**.
- If the answer is **right**:
  - We **trust** the analysis, which is then used as CoT.




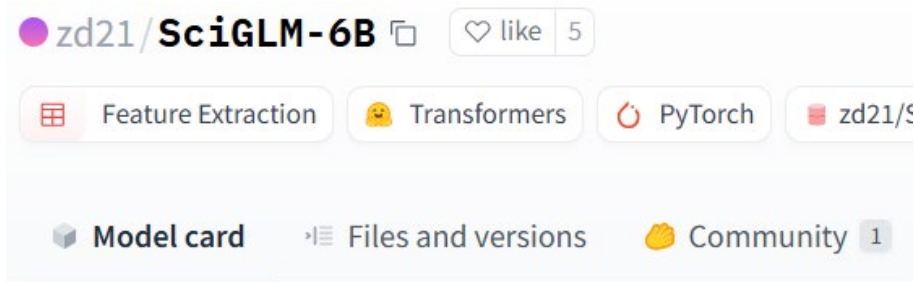
# Self-Reflective Annotation

- Even GPT-4 cannot consistently produce correct answers after multiple trials, so only a small proportion of collected questions can have CoT.
- [Prompt 1] The following input consists of a science problem, please generate an elaborate step-by-step solution to the problem. → 19.8K correct + 22.7K wrong
- [Prompt 2] The following input consists of a science problem and a corresponding solution. However, this solution is incorrect, please reflect on its errors and then generate a correct step-by-step solution to the problem. → 5.5K correct + 17.2K wrong
- [Prompt 3] The following input consists of a science problem, a corresponding solution and the real answer. The given solution is incorrect, please reflect on its errors and then generate a correct step-by-step solution to the problem based on the real answer. → 7.7K correct + 9.5K wrong

# Instruction Tuning with SciInstruct

- Architecture:
  - ChatGLM3-6B
  - ChatGLM3-32B

<https://huggingface.co/zd21/SciGLM-6B>

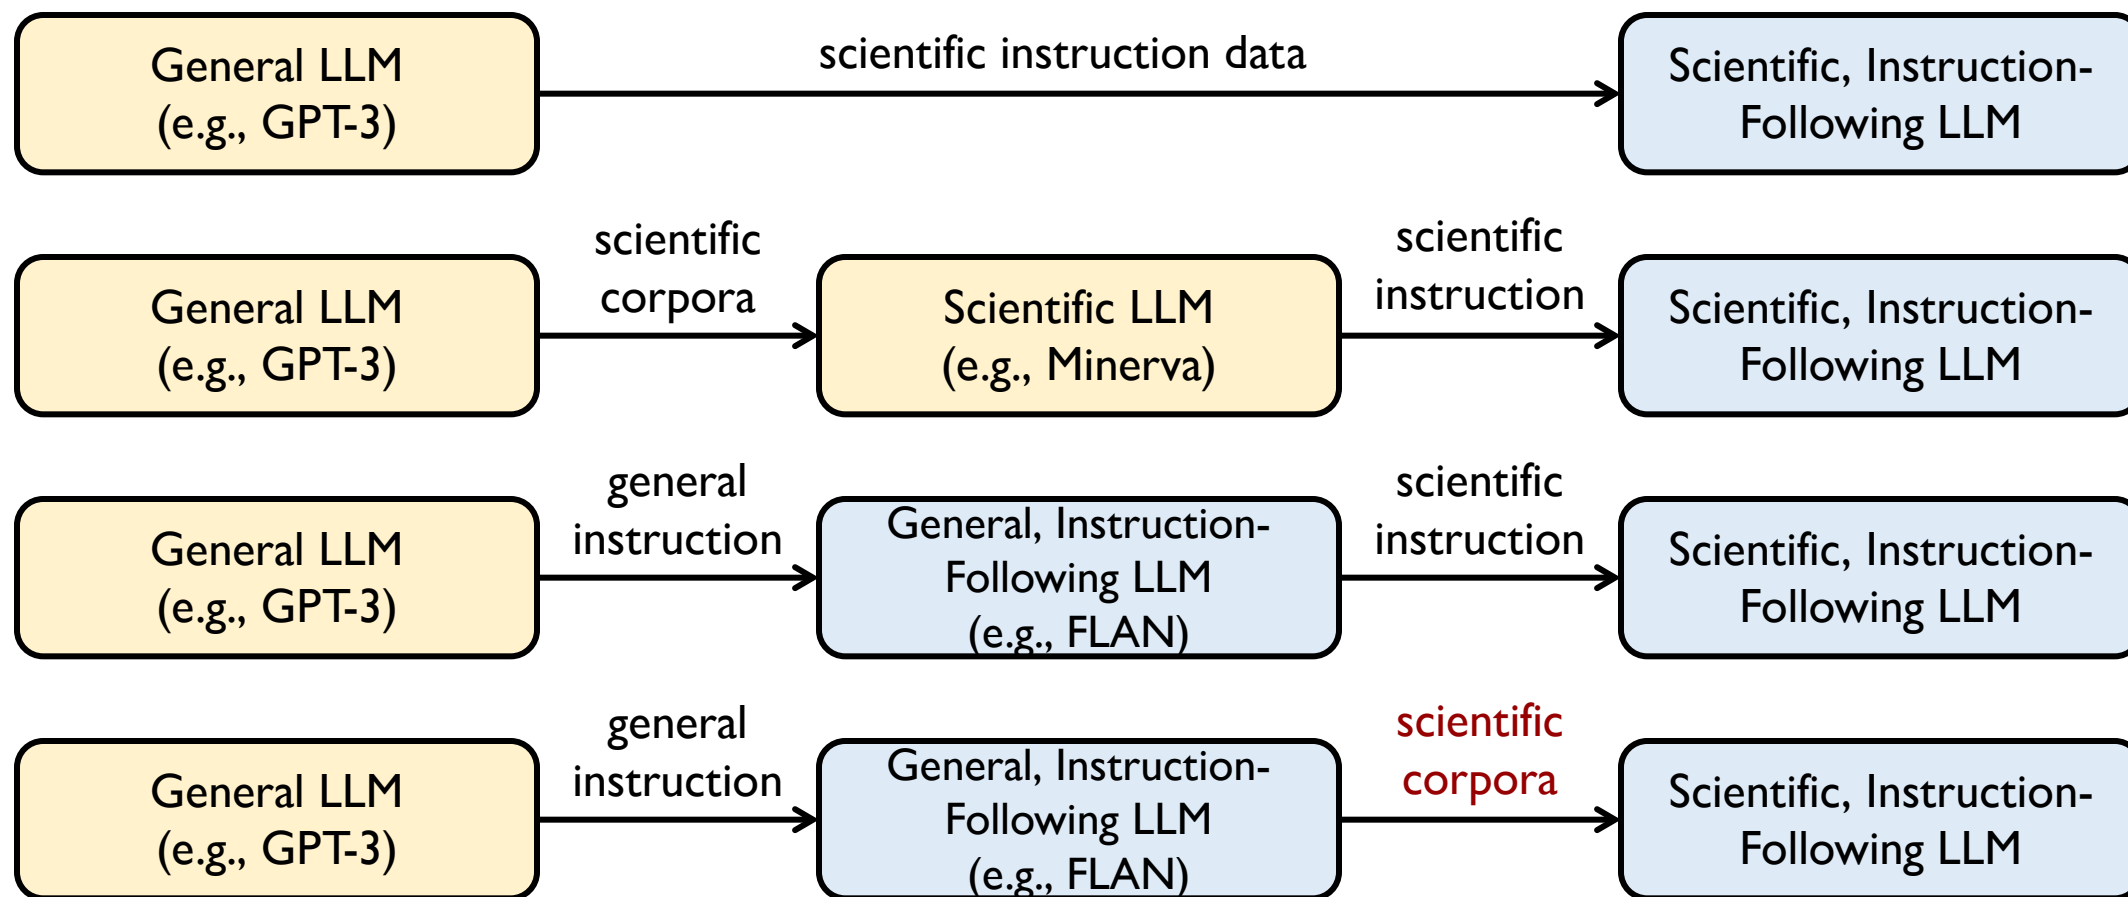


Average accuracy on CEval-Sci, SciEval, SciBench, MATH, and SAT-Math benchmarks of different LLMs.

# Agenda

- Unsupervised Next Token Prediction
  - General Domain: GPT-3
  - Mathematics: Minerva
- Supervised Fine-Tuning / Instruction Tuning
  - General Domain: FLAN
  - Science: ScilInstruct
  - Biomedicine: **BioMistral**
  - Geoscience: OceanGPT


# Different Roadmaps to Get a Scientific, Instruction-Following LLM




without using any  
scientific instruction?

# BioMistral: Mistral + Unsupervised Next Token Prediction


- **Architecture:** Mistral 7B (already fine-tuned on general-domain instruction data)



**Vanilla Attention**



**Sliding Window Attention**



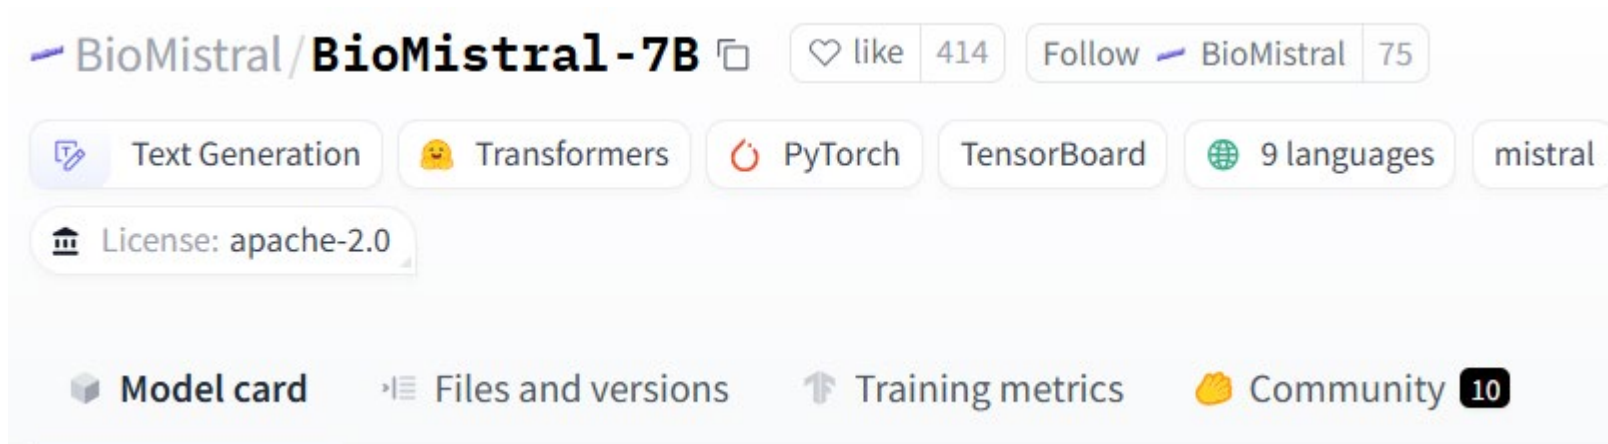
**Effective Context Length**

Parameter	Value
dim	4096
n_layers	32
head_dim	128
hidden_dim	14336
n_heads	32
n_kv_heads	8
window_size	4096
context_len	8192
vocab_size	32000

# BioMistral: Mistral + Unsupervised Next Token Prediction

- **Architecture:** Mistral 7B (already fine-tuned on general-domain instruction data)
- **Data:** PMC full text
  - A large biomedical corpus, no annotated or harvested instructions

<https://huggingface.co/BioMistral/BioMistral-7B>





# Datasets for Evaluating BioMistral

- MMLU [1]: college biology, college medicine, anatomy, professional medicine, medical genetics, and clinical knowledge
- MedQA [2]: questions from the US Medical License Exam (USMLE)
- MedMCQA [3]: questions from the Indian medical entrance examinations (AIIMS/NEET)
- PubMedQA [4]: rewrite PubMed paper titles and abstracts into yes/no/maybe questions

	MMLU						MedQA	PubMedQA	MedMCQA
	Clinical KG	Medical Genetics	Anatomy	Pro Medicine	College Biology	College Medicine			
Answer options	A / B / C / D	A / B / C / D	A / B / C / D	A / B / C / D	A / B / C / D	A / B / C / D	A / B / C / D / (E)	Yes / No / Maybe	A / B / C / D
Train / Valid. / Test	0 / 0 / 265	0 / 0 / 100	0 / 0 / 135	0 / 0 / 272	0 / 0 / 144	0 / 0 / 173	10178 / 1272 / 1273	211269 / 500 / 500	146257 / 36565 / 4183
Words / Questions	11.09	12.34	13.65	105.46	22.40	48.84	118.16	13.08	14.05
Context	×	×	×	×	×	×	×	✓	×

[1] *Measuring Massive Multitask Language Understanding*. ICLR 2021.

[2] *What Disease does this Patient Have? A Large-scale Open Domain Question Answering Dataset from Medical Exams*. arXiv 2020.

[3] *MedMCQA: A Large-scale Multi-Subject Multi-Choice Dataset for Medical domain Question Answering*. CHIL 2022.

[4] *PubMedQA: A Dataset for Biomedical Research Question Answering*. EMNLP 2019.

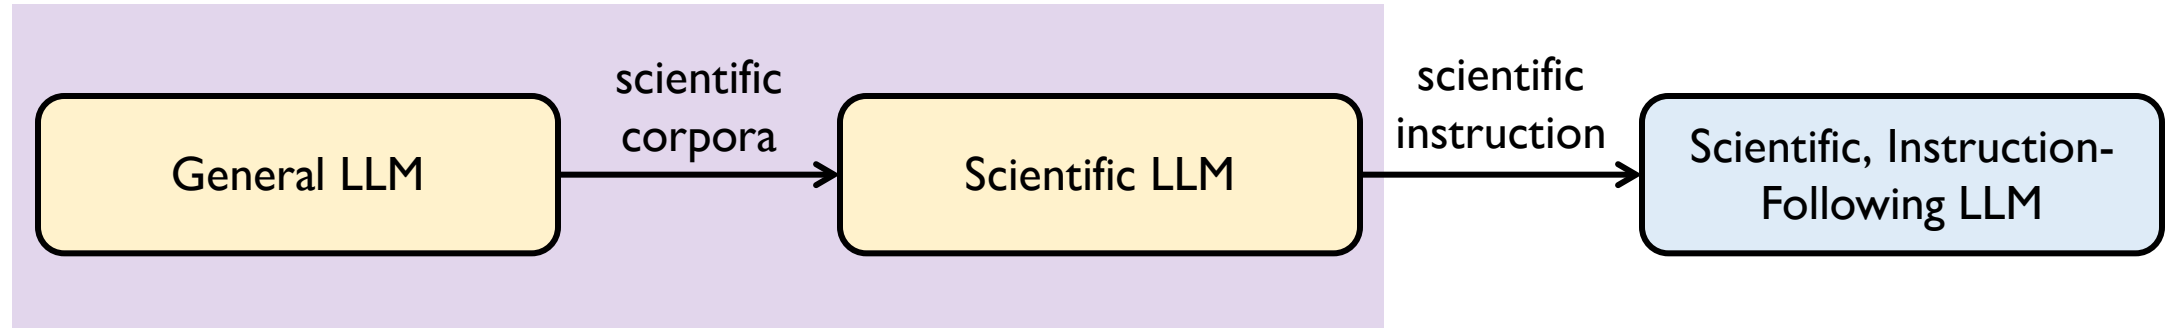
# Performance of BioMistral

	MMLU						MedQA	MedQA 5 opts	PubMedQA	MedMCQA	Avg.
	Clinical KG	Medical Genetics	Anatomy	Pro Medicine	College Biology	College Medicine					
<b>BioMistral 7B</b>	59.9 $\pm$ 1.2	64.0 $\pm$ 1.6	56.5 $\pm$ 1.8	60.4 $\pm$ 0.5	59.0 $\pm$ 1.5	54.7 $\pm$ 1.0	50.6 $\pm$ 0.3	42.8 $\pm$ 0.3	77.5 $\pm$ 0.1	48.1 $\pm$ 0.2	57.3
<b>Mistral 7B Instruct</b>	<b>62.9</b> $\pm$ 0.2	57.0 $\pm$ 0.8	55.6 $\pm$ 1.0	59.4 $\pm$ 0.6	62.5 $\pm$ 1.0	<u>57.2</u> $\pm$ 2.1	42.0 $\pm$ 0.2	40.9 $\pm$ 0.4	75.7 $\pm$ 0.4	46.1 $\pm$ 0.1	55.9
<b>BioMistral 7B Ensemble</b>	<u>62.8</u> $\pm$ 0.5	62.7 $\pm$ 0.5	<u>57.5</u> $\pm$ 0.3	<b>63.5</b> $\pm$ 0.8	64.3 $\pm$ 1.6	55.7 $\pm$ 1.5	50.6 $\pm$ 0.3	43.6 $\pm$ 0.5	77.5 $\pm$ 0.2	<b>48.8</b> $\pm$ 0.0	58.7
<b>BioMistral 7B DARE</b>	62.3 $\pm$ 1.3	<b>67.0</b> $\pm$ 1.6	55.8 $\pm$ 0.9	61.4 $\pm$ 0.3	<b>66.9</b> $\pm$ 2.3	<b>58.0</b> $\pm$ 0.5	<b>51.1</b> $\pm$ 0.3	<b>45.2</b> $\pm$ 0.3	<u>77.7</u> $\pm$ 0.1	<u>48.7</u> $\pm$ 0.1	<b>59.4</b>
<b>BioMistral 7B TIES</b>	60.1 $\pm$ 0.9	<u>65.0</u> $\pm$ 2.4	<b>58.5</b> $\pm$ 1.0	60.5 $\pm$ 1.1	60.4 $\pm$ 1.5	56.5 $\pm$ 1.9	49.5 $\pm$ 0.1	43.2 $\pm$ 0.1	77.5 $\pm$ 0.2	48.1 $\pm$ 0.1	57.9
<b>BioMistral 7B SLERP</b>	62.5 $\pm$ 0.6	64.7 $\pm$ 1.7	55.8 $\pm$ 0.3	<u>62.7</u> $\pm$ 0.3	<u>64.8</u> $\pm$ 0.9	56.3 $\pm$ 1.0	<u>50.8</u> $\pm$ 0.6	<u>44.3</u> $\pm$ 0.4	<b>77.8</b> $\pm$ 0.0	48.6 $\pm$ 0.1	<u>58.8</u>
<b>MedAlpaca 7B</b>	53.1 $\pm$ 0.9	58.0 $\pm$ 2.2	54.1 $\pm$ 1.6	58.8 $\pm$ 0.3	58.1 $\pm$ 1.3	48.6 $\pm$ 0.5	40.1 $\pm$ 0.4	33.7 $\pm$ 0.7	73.6 $\pm$ 0.3	37.0 $\pm$ 0.3	51.5
<b>PMC-LLaMA 7B</b>	24.5 $\pm$ 1.7	27.7 $\pm$ 1.7	35.3 $\pm$ 0.7	17.4 $\pm$ 1.7	30.3 $\pm$ 0.9	23.3 $\pm$ 1.7	25.5 $\pm$ 0.9	20.2 $\pm$ 0.1	72.9 $\pm$ 1.2	26.6 $\pm$ 0.1	30.4
<b>MediTron-7B</b>	41.6 $\pm$ 1.2	50.3 $\pm$ 2.1	46.4 $\pm$ 0.9	27.9 $\pm$ 0.3	44.4 $\pm$ 2.6	30.8 $\pm$ 0.7	41.6 $\pm$ 0.5	28.1 $\pm$ 0.5	74.9 $\pm$ 0.1	41.3 $\pm$ 0.2	42.7
<b>BioMedGPT-LM-7B</b>	51.4 $\pm$ 0.4	52.0 $\pm$ 1.4	49.4 $\pm$ 2.7	53.3 $\pm$ 0.6	50.7 $\pm$ 0.0	49.1 $\pm$ 0.8	42.5 $\pm$ 0.3	33.9 $\pm$ 0.5	76.8 $\pm$ 0.3	37.6 $\pm$ 0.4	49.7
<b>GPT-3.5 Turbo 1106*</b>	74.71 $\pm$ 0.3	74.00 $\pm$ 2.2	65.92 $\pm$ 0.6	72.79 $\pm$ 1.6	72.91 $\pm$ 1.7	64.73 $\pm$ 2.9	57.71 $\pm$ 0.3	50.82 $\pm$ 0.7	72.66 $\pm$ 1.0	53.79 $\pm$ 0.2	66.0

# Agenda

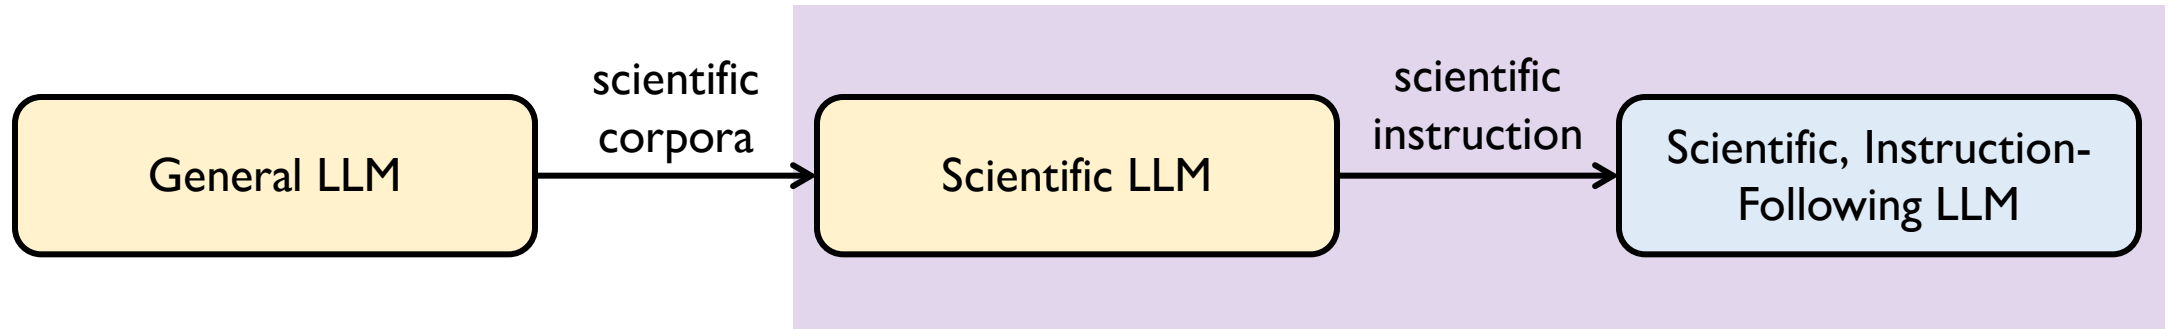
- Unsupervised Next Token Prediction
  - General Domain: GPT-3
  - Mathematics: Minerva
- Supervised Fine-Tuning / Instruction Tuning
  - General Domain: FLAN
  - Science: SciInstruct
  - Biomedicine: BioMistral
  - Geoscience: OceanGPT

# OceanGPT: An LLM for Ocean Science



- **Step 1:** Unsupervised next token prediction
  - 67,633 full-text papers
  - ocean physics, ocean chemistry, ocean biology, geology, hydrology, etc.

# OceanGPT: An LLM for Ocean Science



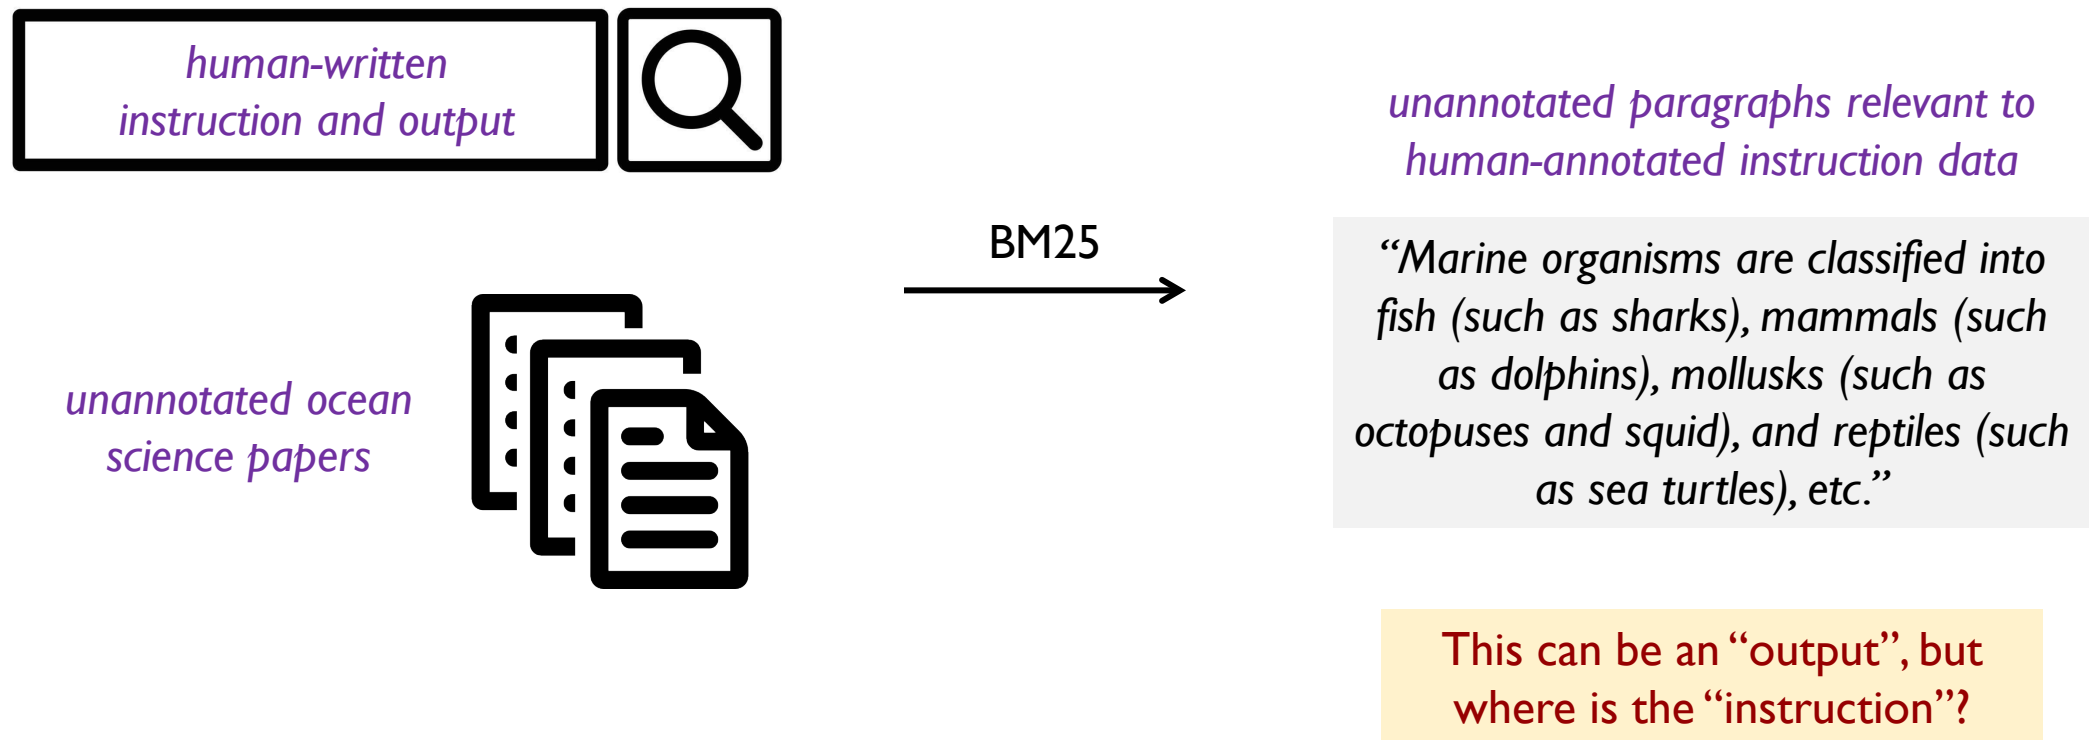
- **Unsupervised next token prediction**
  - 67,633 full-text papers
  - ocean physics, ocean chemistry, ocean biology, geology, hydrology, etc.
- **Instruction tuning**
  - Hard to find benchmark datasets or sufficient exam questions related to ocean science
  - A common challenge if you want to build an LLM for **a fine-grained field**

# Constructing Instruction Tuning Data for a Fine-Grained Field

- **Step 1:** Dozens of annotators with rich backgrounds in marine science write some representative example for each marine topic.
- E.g.,
  - *Instruction:* Please recommend several rare marine plants and animals and their ecological value.
  - *Output:* Rare marine animals and plants include whales, dolphins, jewel-like seaweed, seahorses, etc. These species play a crucial role in maintaining the balance of the ecosystem and require protection.
- However, you can only obtain a small number of instruction tuning data from humans!
  - Use LLMs to paraphrase human-written data
  - Retrieve more data from domain-specific corpora

# Constructing Instruction Tuning Data for a Fine-Grained Field

- **Step 2:** Build more instruction tuning data by generating questions given unannotated text.



# Constructing Instruction Tuning Data for a Fine-Grained Field

- **Step 2:** Build more instruction tuning data by generating questions given unannotated text.

*You are a helpful ocean assistant. You are to extract the question from each of the answer provided.*

*“Marine organisms are classified into fish (such as sharks), mammals (such as dolphins), mollusks (such as octopuses and squid), and reptiles (such as sea turtles), etc.”*

**This can be an “output”.**



GPT-3.5

*Please classify the following marine creatures: shark, dolphin, squid, octopus.*


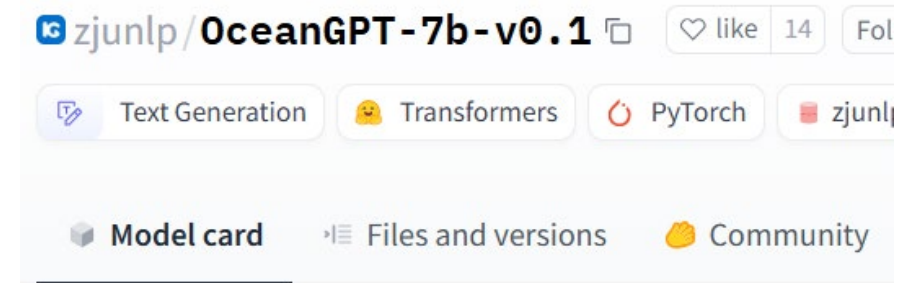
**This can be an “instruction”.**



# Model Details of OceanGPT

- **Architecture:** LLaMA-2 7B
- **Data:** 150K (instruction, output) pairs
- **Tuning Method:** Low-Rank Adaptation (LoRA)

<https://huggingface.co/zjunlp/OceanGPT-7b-v0.1>



$$h = (W_0 + \Delta W)x = (W_0 + B \times A)x$$

$x$ : input

$h$ : output

$W_0$ : original model parameters (i.e., LLaMA-2)

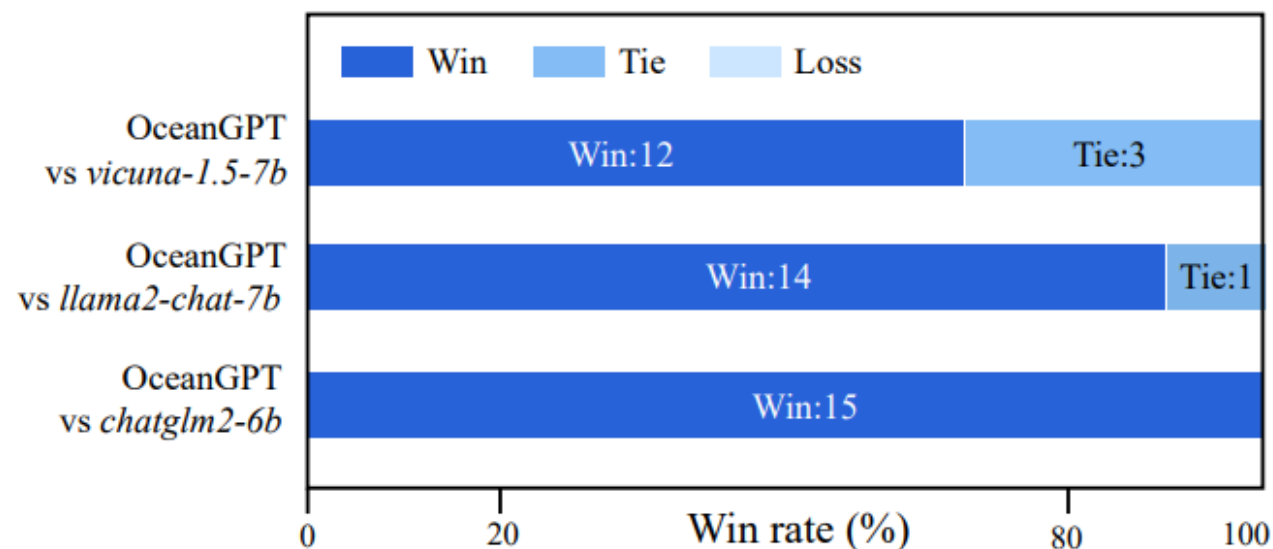
$(W_0 + \Delta W)$ : new model parameters (i.e., OceanGPT)

$B \times A$ : a low-rank approximation of  $\Delta W$

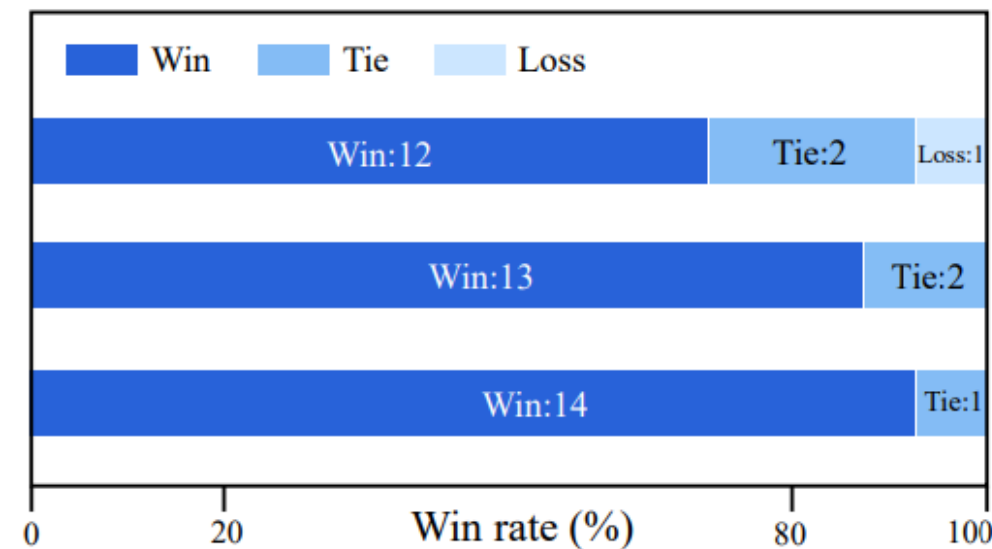
# Evaluation of OceanGPT

- Tasks: (OceanBench: <https://huggingface.co/datasets/zjunlp/OceanBench>)
  - **Analysis:** “*Analyzing the bioactive components of seaweed and its application prospects*”
  - **Commonsense Reasoning:** “*Infer the reasons for the increase in seawater turbidity*”
  - **Recommendation:** “*Recommend an instrument capable of detecting ocean pollution*”
  - **Editing:** “*Edit a popular science article on ocean circulation and pollution*”
  - **Question Answering:** “*What is the main electrolyte in seawater?*”
  - **Classification:** “*What are the basic classifications of tropical cyclones?*”
  - **Open-Ended Generation:** “*Write an argumentative essay on ocean conservation and management*”
  - **Description:** “*Describe the mechanism of underwater mineral enrichment*”
  - ...

# Performance of OceanGPT

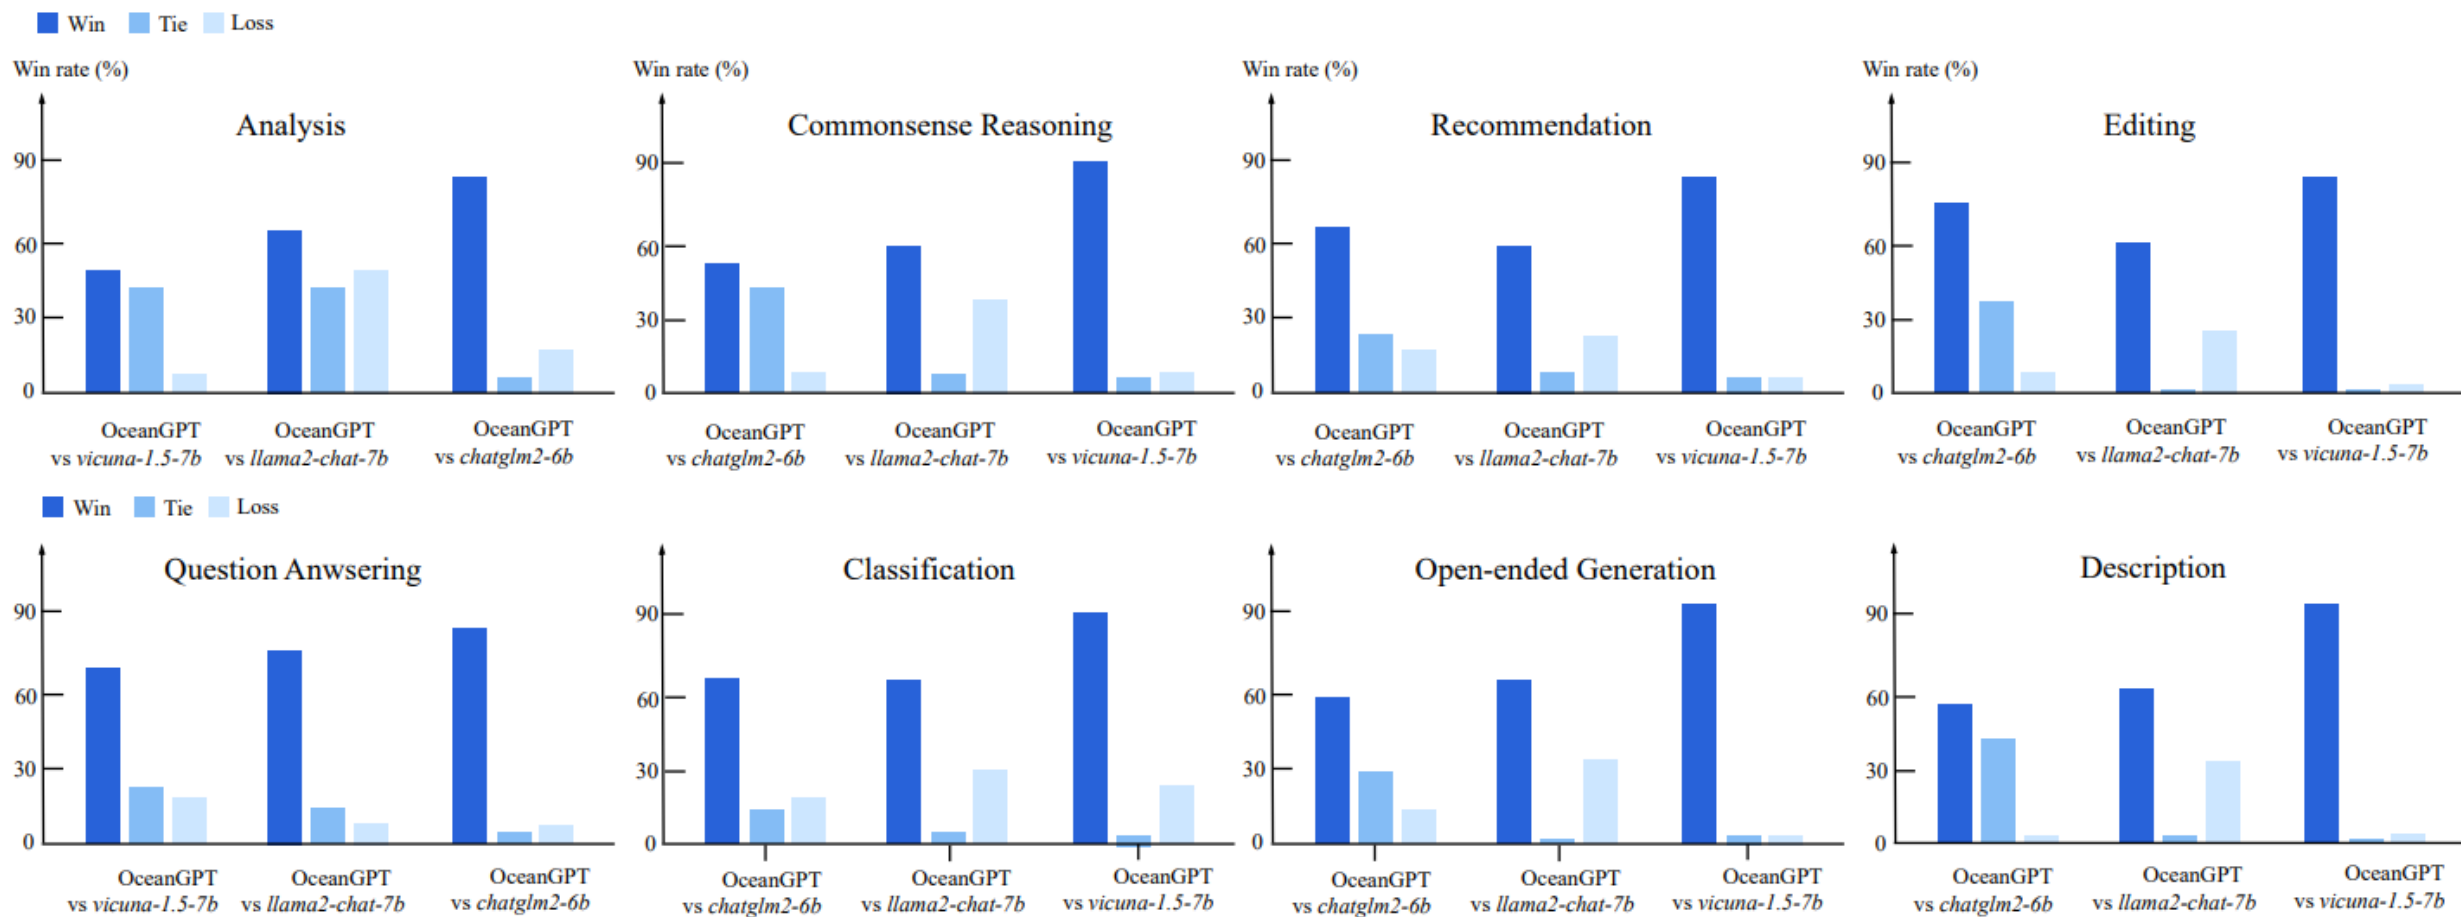


GPT-4 evaluation



Human evaluation

# Performance of OceanGPT



# Take-Away Messages

- Tuning LLMs to follow instructions enables them to deal with **unseen instructions without any examples** during inference (i.e., zero-shot generalization).
- Multiple ways to **harvest** instruction tuning data in the scientific domain:
  - Convert benchmark datasets to the instruction tuning format
  - Collect questions from textbooks, problem sets, etc.
  - **May not work for a new, fine-grained field!**
- Off-the-shelf powerful LLMs (e.g., GPT-4) can help the construction of instruction tuning data
  - Recover the chain-of-thought
  - Generate more instruction tuning data to complement human annotations



Thank You!

Course Website: <https://yuzhang-teaching.github.io/CSCE689-S25.html>