TEXAS A&M UNIVERSITY
Engineering

DNA/RNA/Single -Cell Language Models

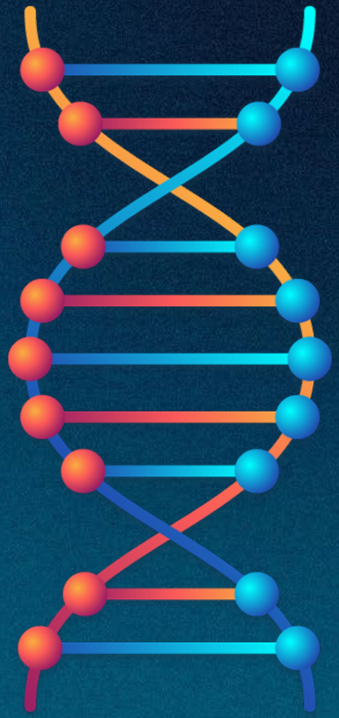

Student: Omnia Sarhan
Instructor: Yu Zhang

Table of contents

01.

DNABERT

Pre-trained Bidirectional Encoder Representations from Transformers Model for DNA-Language in Genome

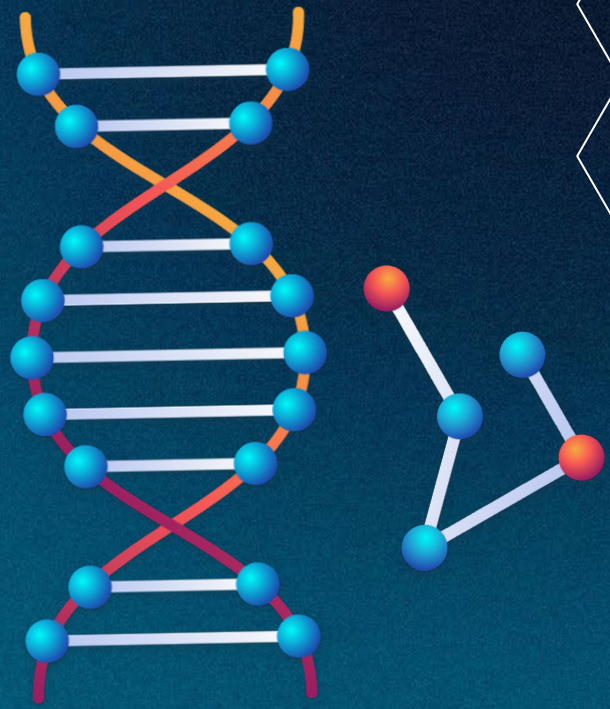

Introduction

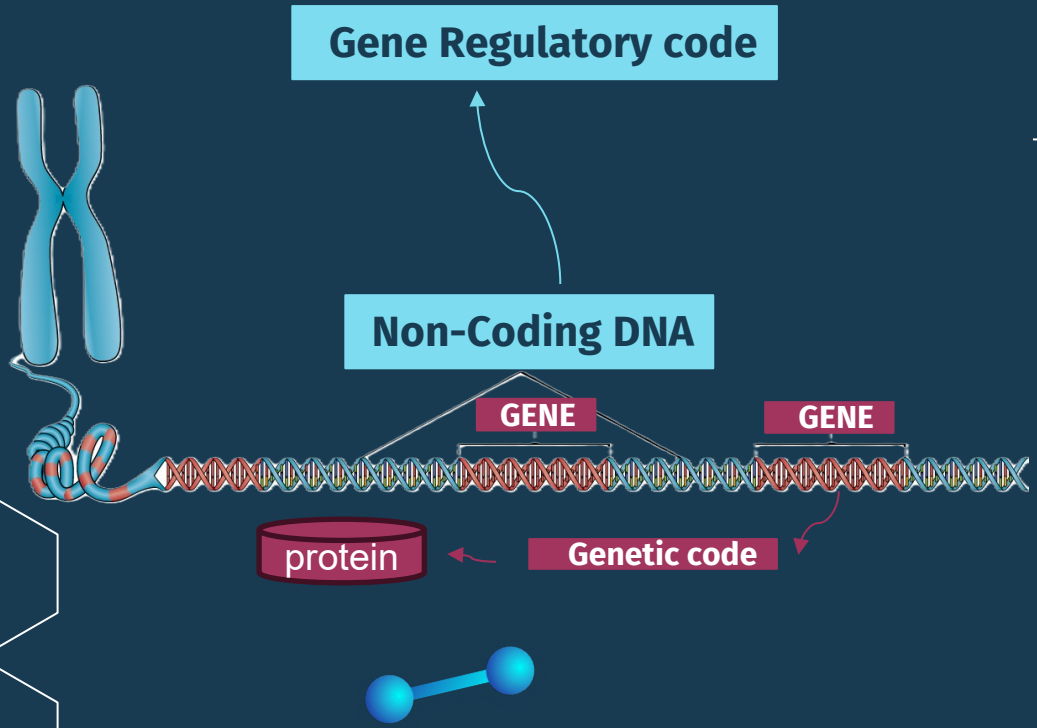

Problem Statement

- Deciphering Non-Coding DNA for hidden instructions is challenging.
- Traditional models fail to capture long-range dependencies and polysemous relationships within DNA sequences.

Objectives

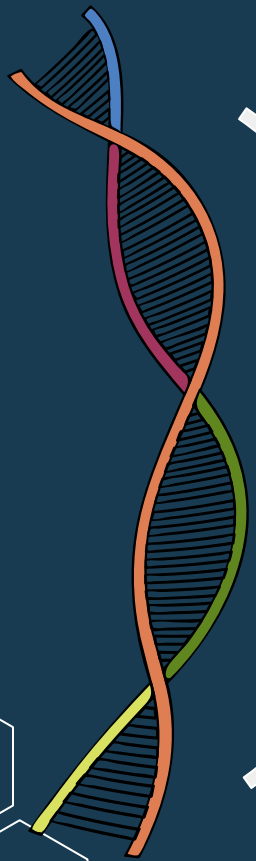

Capture global and transferable contextual information from DNA sequences

Outperform traditional deep learning models in various genomic tasks

Provide visualization mechanisms for interpretation of sequence motifs

Demonstrate cross-organism applicability

Facilitate fine-tuning on task-specific datasets

DNABERT Model

- BERT-based (same architecture)
- Adopts pre-training +fine-tuning

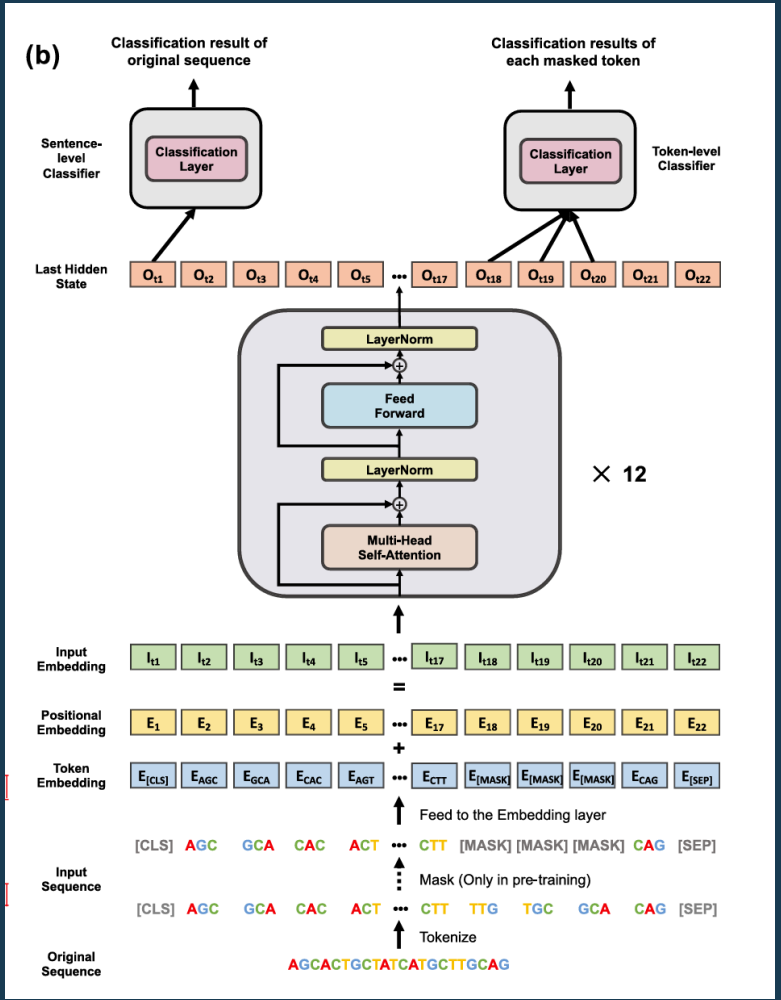

Methodoly

1

Tokenization

- k-mer representation instead of single nucleotides.
- Different values of k (3, 4, 5, 6)
- Added special tokens like [CLS], [PAD], [UNK], [SEP], and [MASK]

2

Pre-training

- masked language modeling (MLM) for random masking
- human genome (5-510 base pairs)
- 12 Transformer layers, 768 hidden units, and 12 attention heads

3

Fine-tuning

- task-specific datasets
- Long sequences exceeding 512 tokens are split and processed as DNABERT-XL.
- Best = DNABERT-6

Results

DNABERT

- Generalize over tasks
- Identifying functional genetic variants

DNABERT-Splice

Accurately recognizes canonical and non-canonical splice sites

DNABERT-viz

Allows visualization of important regions, contexts and sequence motifs

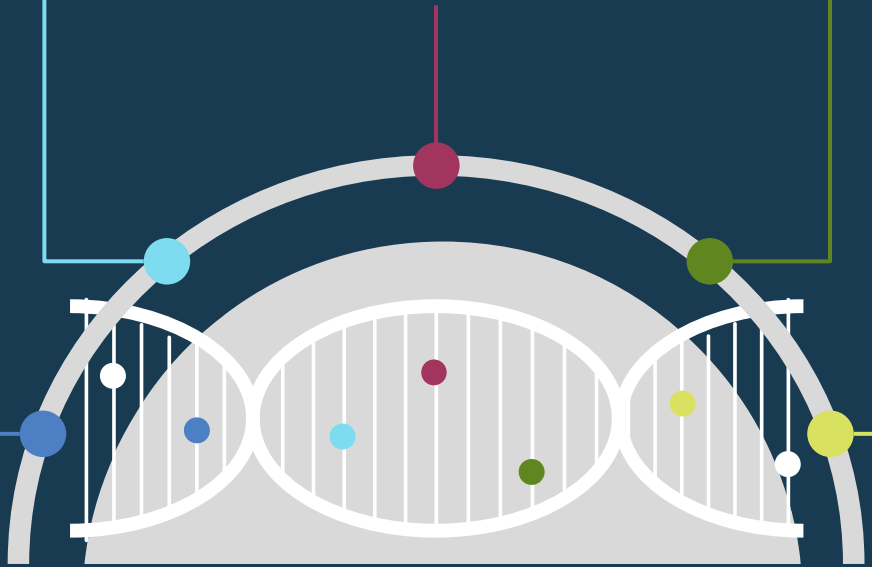

DNABERT-TF

Accurately identifies transcription factor binding sites

DNABERT-Prom

Effectively predicts proximal and core promoter regions

Results: (left to right) accuracy, F1 and MCC

Splice

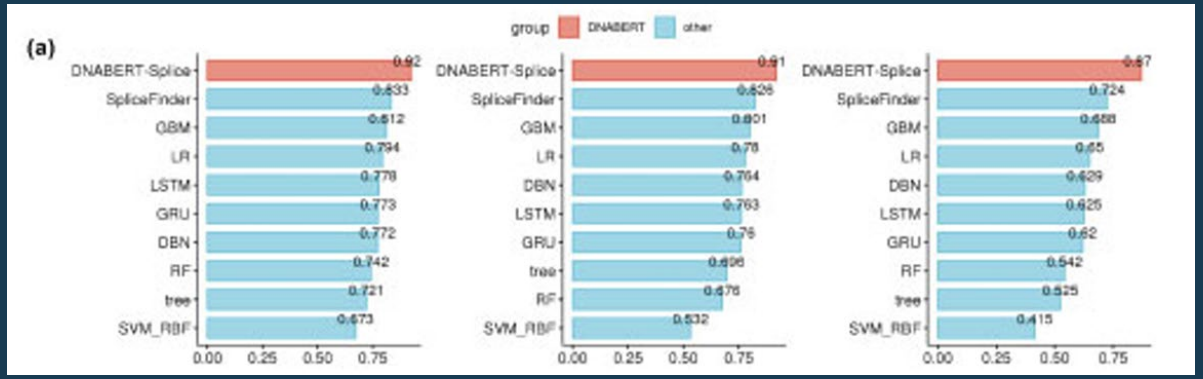

Prom

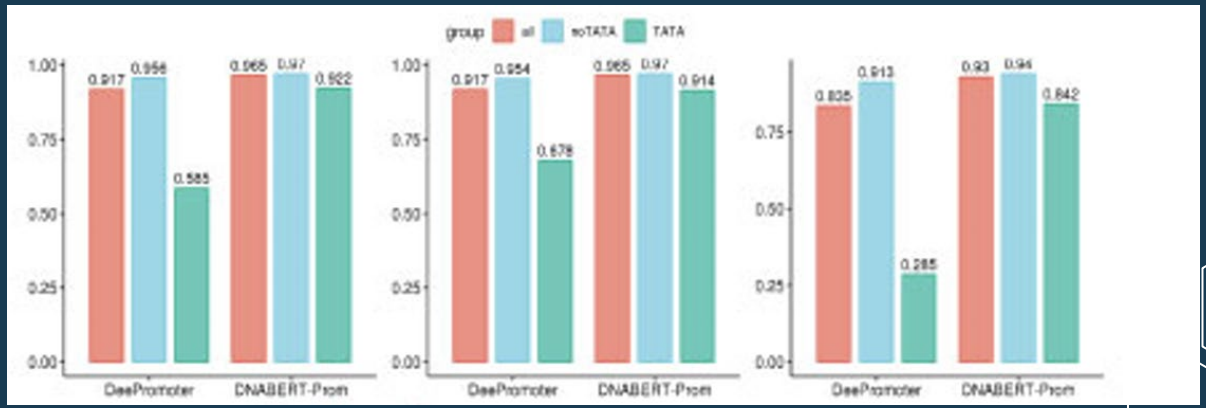

Results: (left to right) accuracy, F1 and MCC

TF

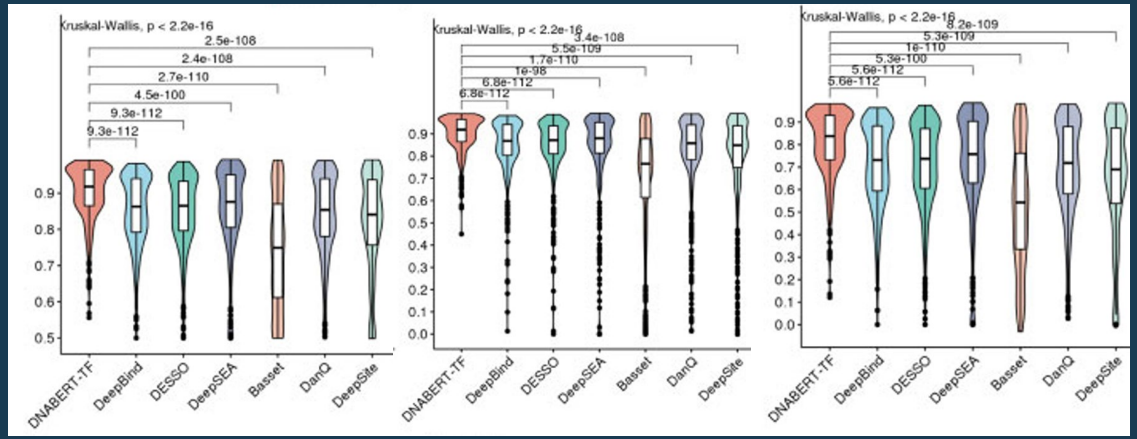

General

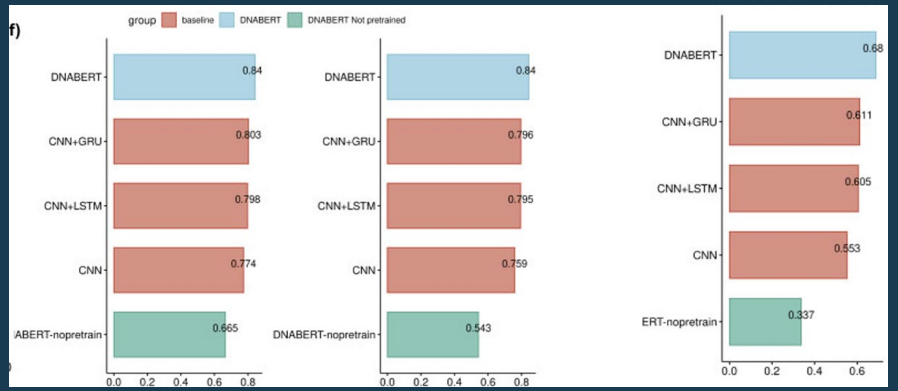

Future Work

1. Other sequence prediction tasks

Determining CREs and enhancer regions from ATAC-seqand DAP-seq

3. Direct machine translation on DNA

s

2. Prediction of binding preferences of RNA-binding proteins (RBPs)

s

02.

5' UTR

A Language Model for Decoding Untranslated Regions of mRNA and Function Predictions

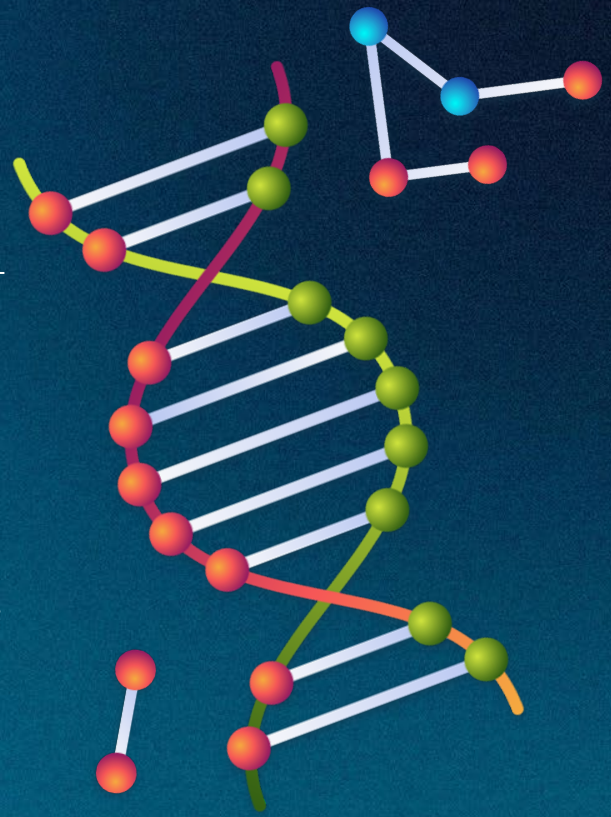

Background

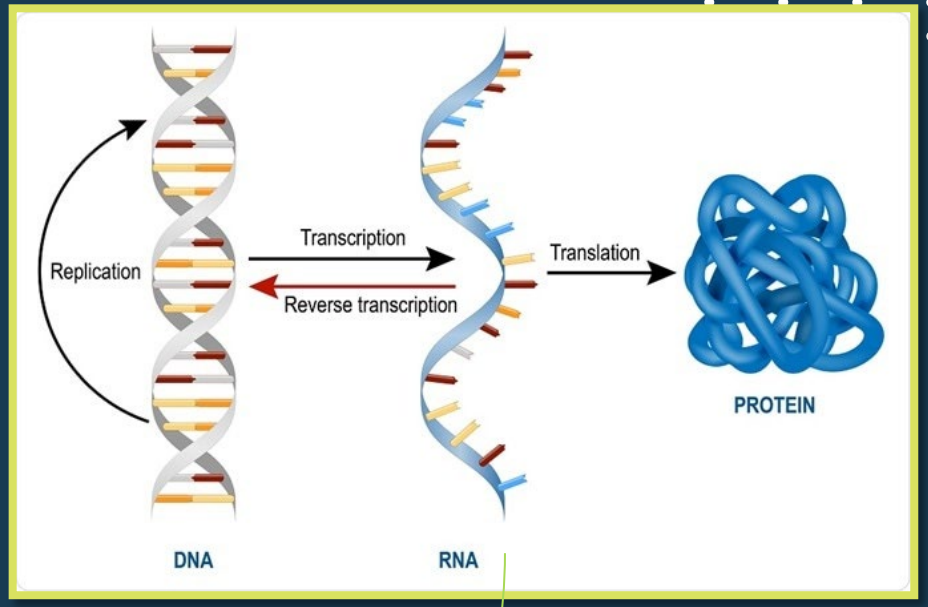
Replication
Transcription
Reverse transcription
Translation
PROTEIN
DNA
RNA


5′ Cap - plays a critical role in translational yield and nucleic acid stability *in vivo*
Poly(A) tail - protects the mRNA from nuclease degradation
5′ Cap
Poly(A)
5′ UTR
Open reading frame
3′ UTR
5′ untranslated region (UTR)
5′ UTR - regulates protein expression levels and translation initiation
3′ UTR - regulates protein expression levels and influences the stability and half-life of the mRNA

Introduction

Problem Statement

No unified foundation model to study function of 5’UTR

Objectives

Use Language model to Extract meaningful semantic representations from UTRs of raw mRNA sequences and map them to predict functions of interest.

5’UTR-LM Model Overview

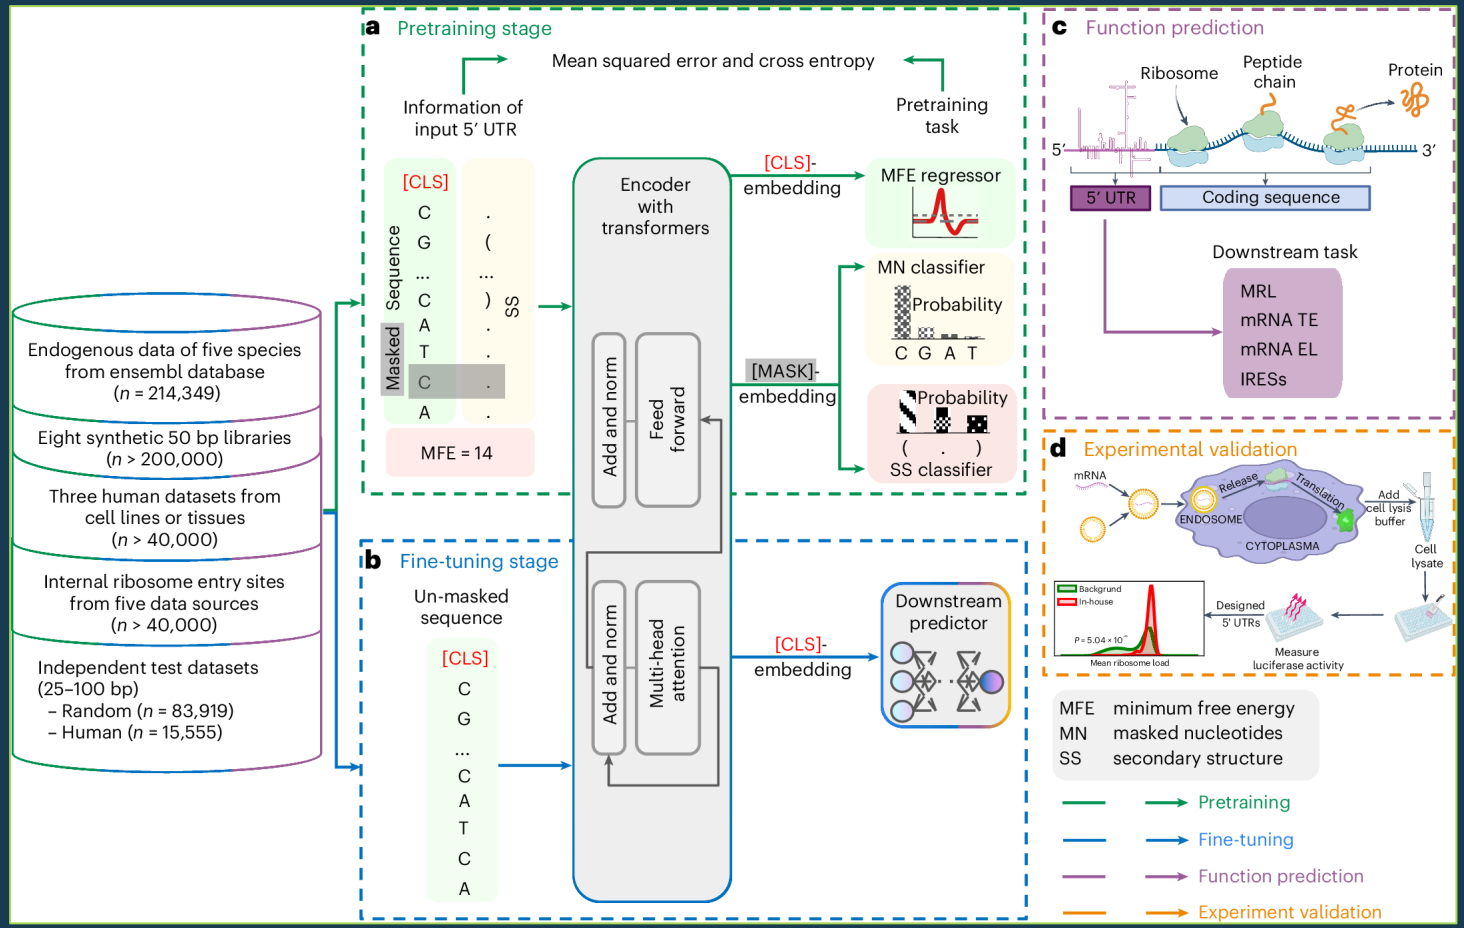

Results

UTR-LM predicts the ribosome loading

URR-LM identifies IRESs

UTR-LM predicts mRNA TE and expression

New designs validated in wet-lab experiments

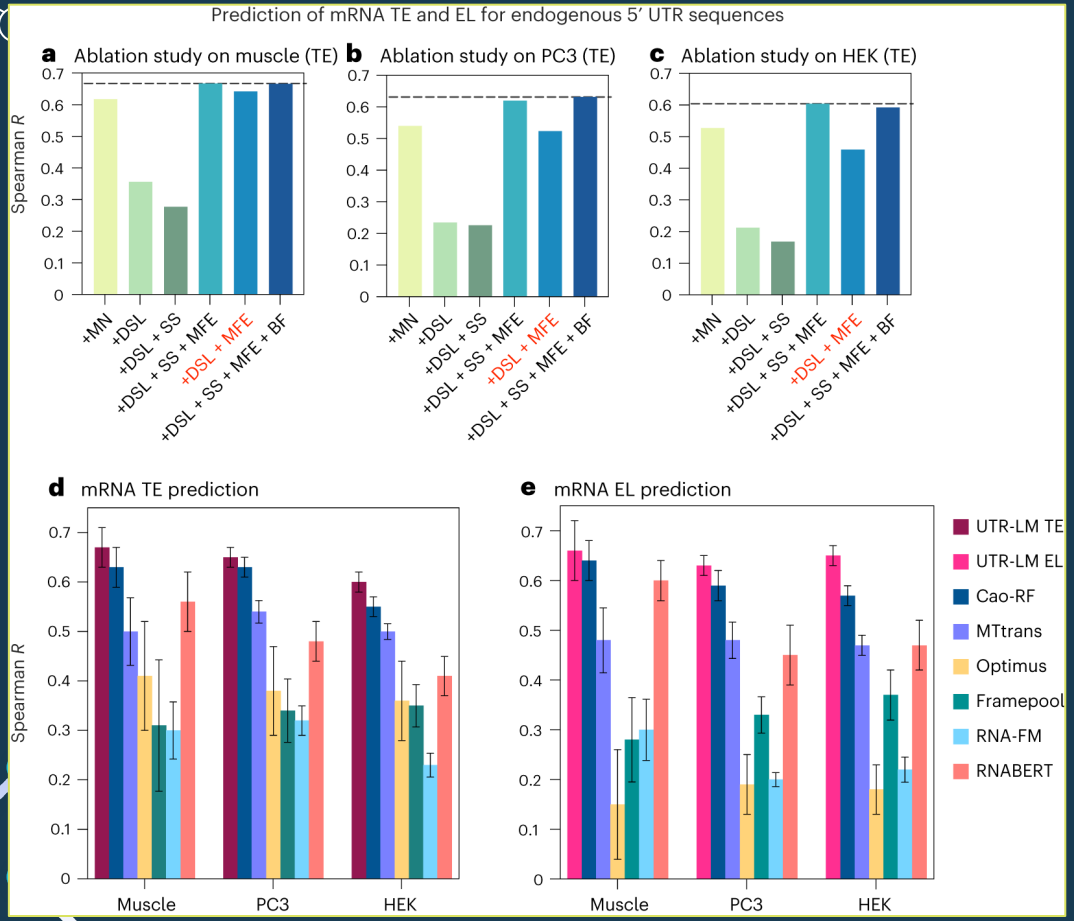

UTR-LM predicts mRNA TE and EL

URR-LM identifies IRESs

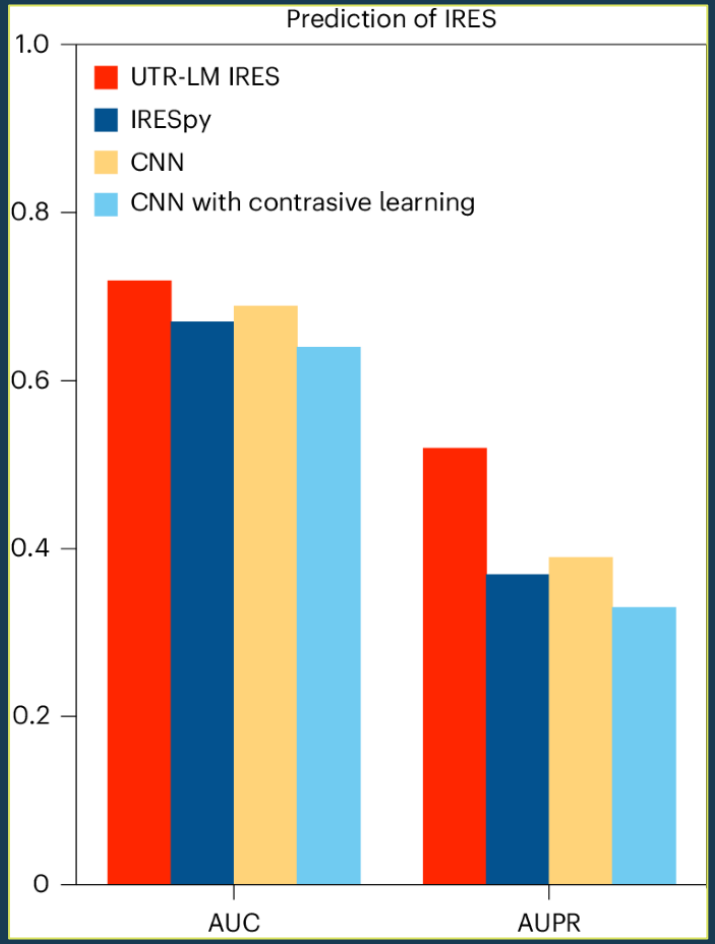

Conclusion

Conclusion

- Outperforms the best-known baseline in each task.
- Performance not limited by sequence length

limitations

- Computationally expensive

Future

sparse transformers for modelling longer RNA sequences and more complex biological functions

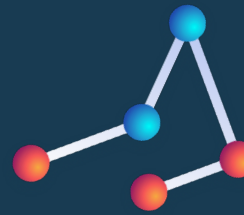

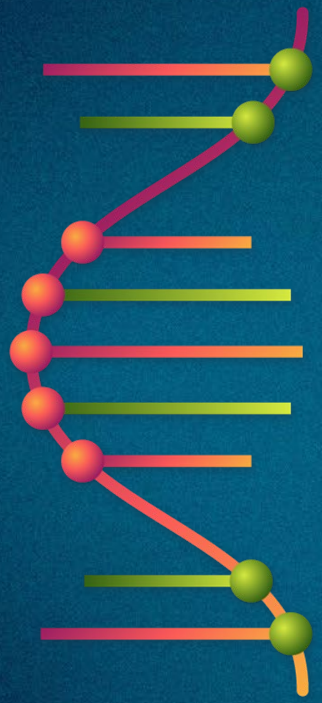

03.

scGPT

Towards Building a Foundation Model for Single-Cell Multi-omics using Generative AI

Single -cell RNA sequencing (scRNA-seq)

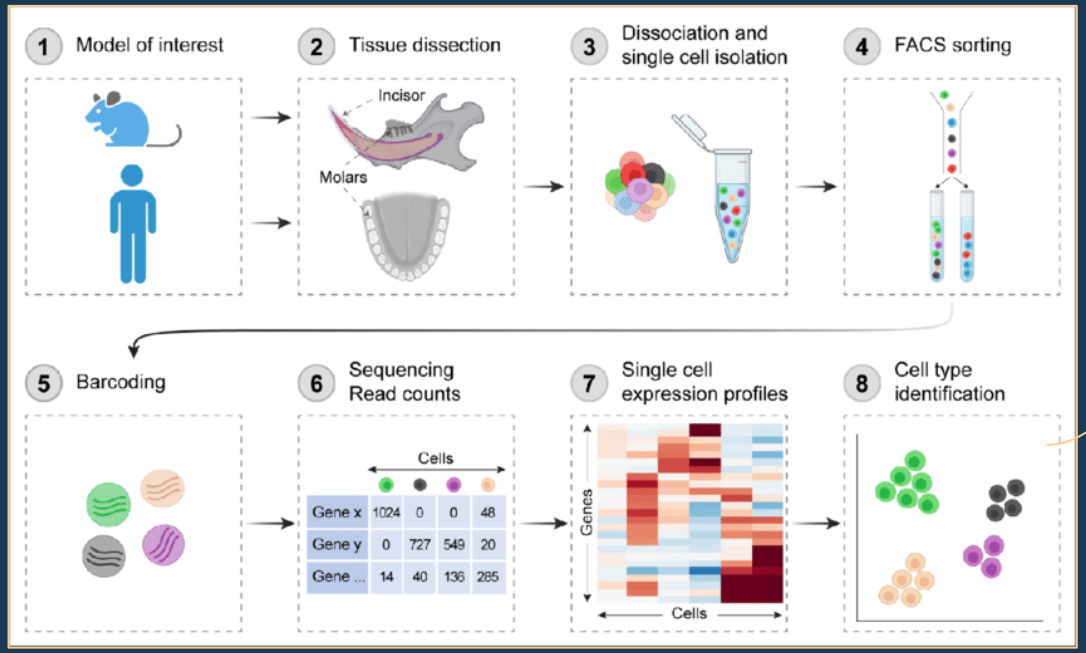

- personalized therapeutic strategies
- cellular heterogeneity exploration
- lineage tracking
- pathogenic mechanism elucidation

Introduction

Problem Statement

Current machine-learning-based methods in single-cell research are scattered

Objectives

- Foundation model pretrained on large-scale data
- comprehend the complex interactions between genes across diverse tissues.

scGPT Model Overview

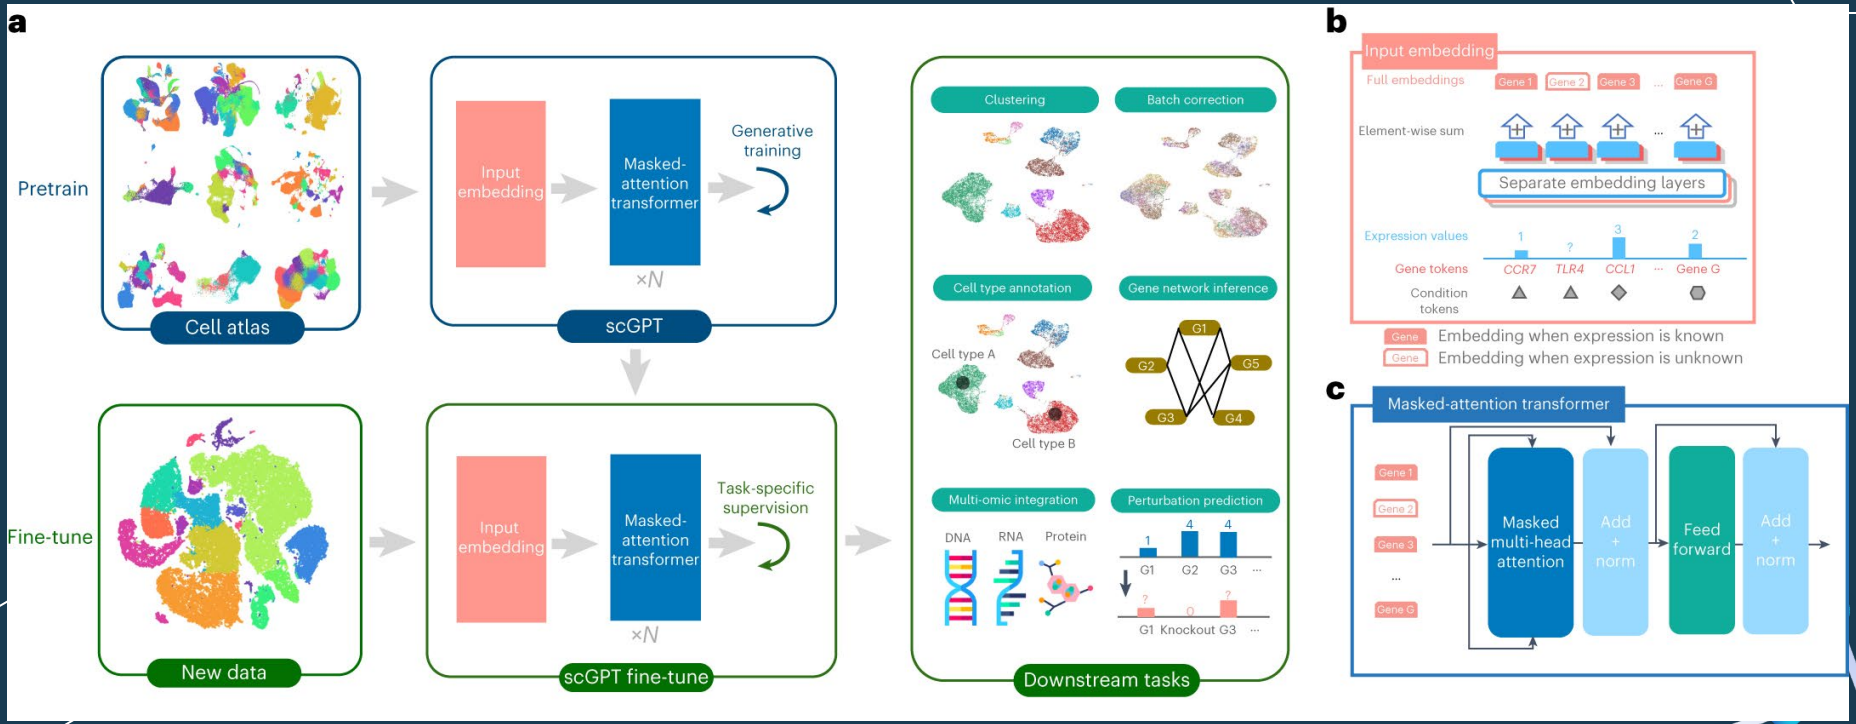

Results

Improves the precision of cell type annotation

multi-batch and multi-omic integration

Predicting Unseen Genetic Perturbation Responses

Uncovers gene networks for specific cell states

Cell Type Annotation

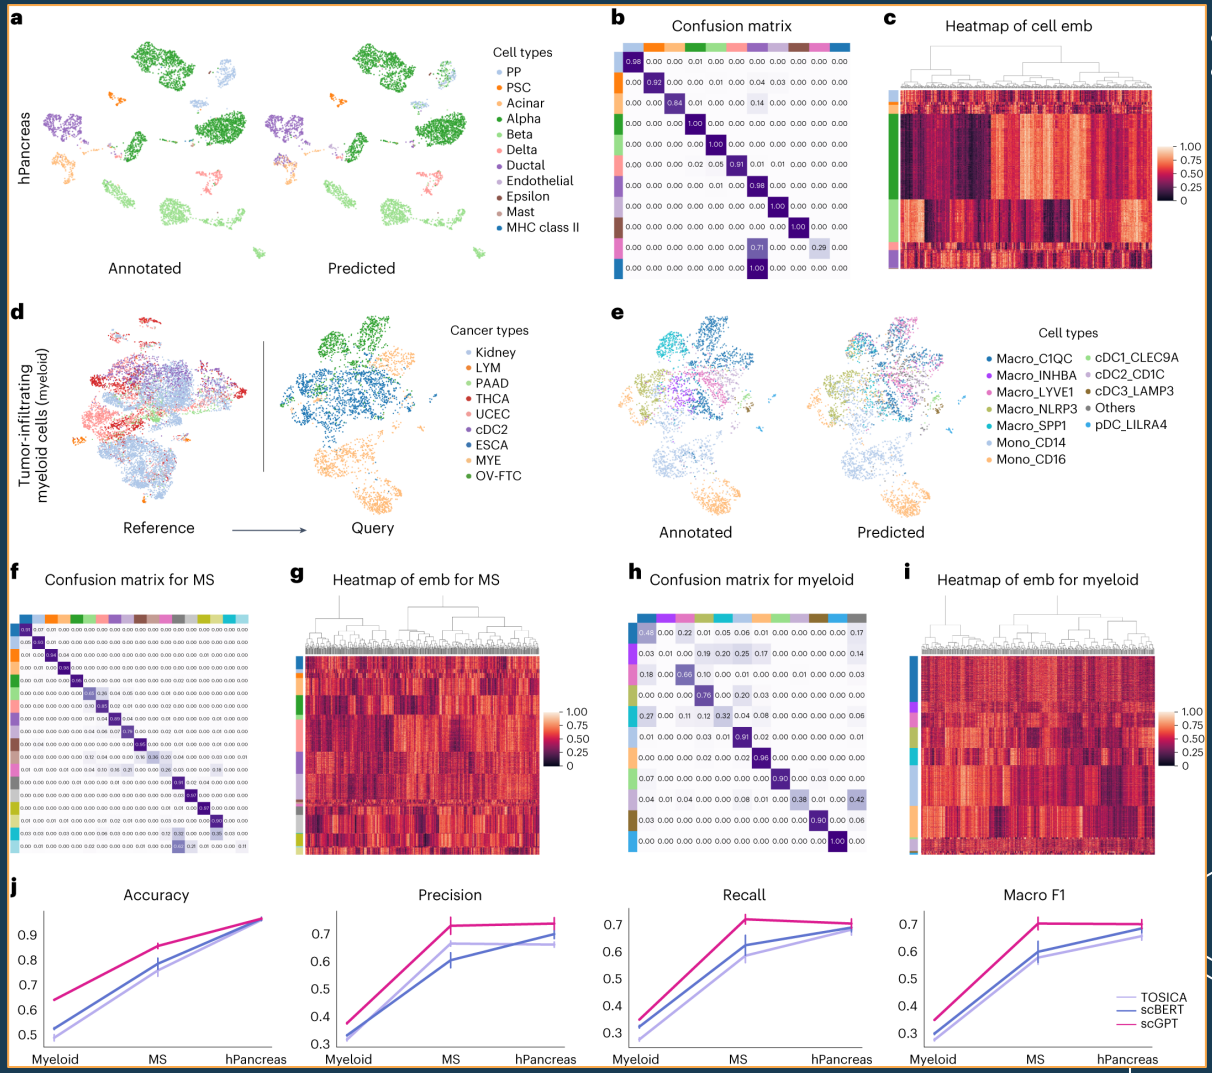

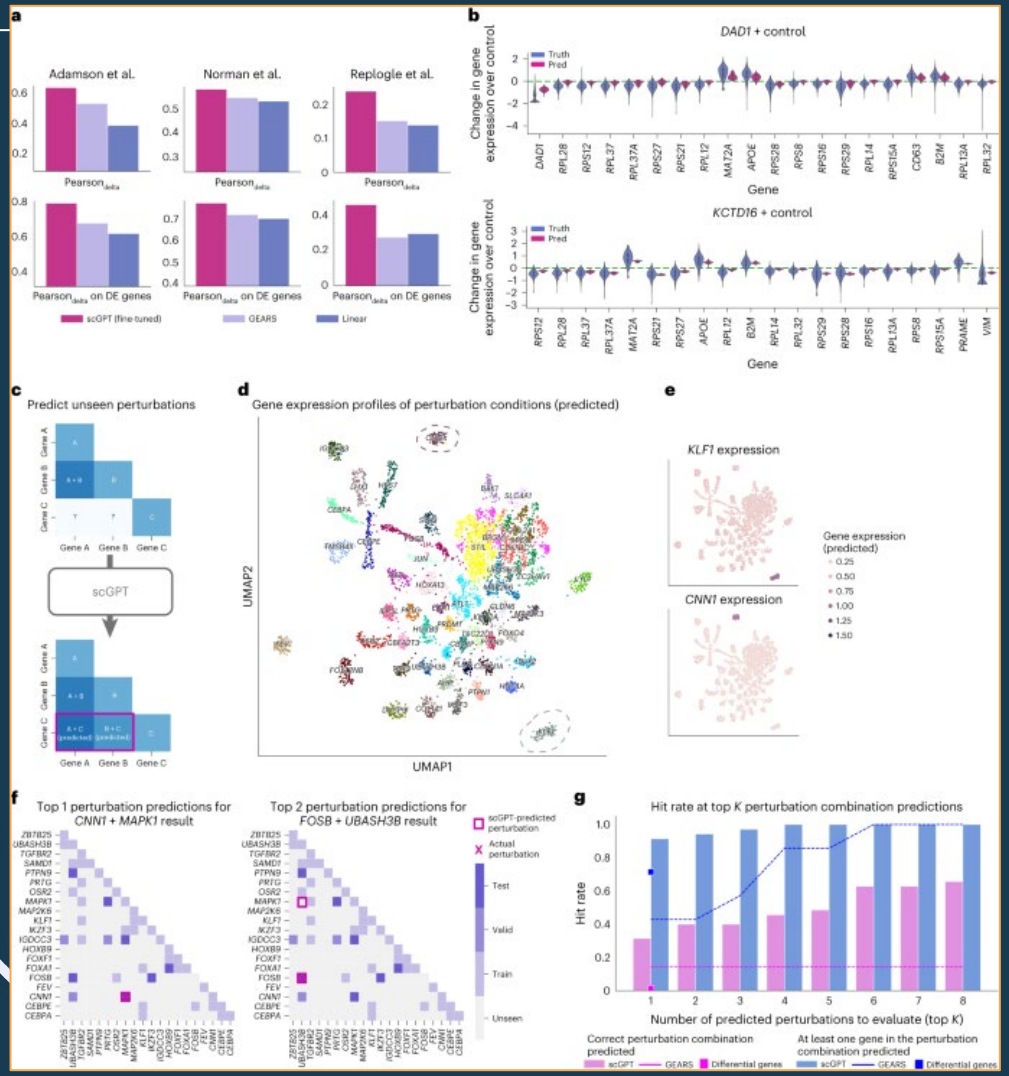

Predicting Unseen Genetic Perturbation Responses

Multi -Batch & Multi -Omic Integration

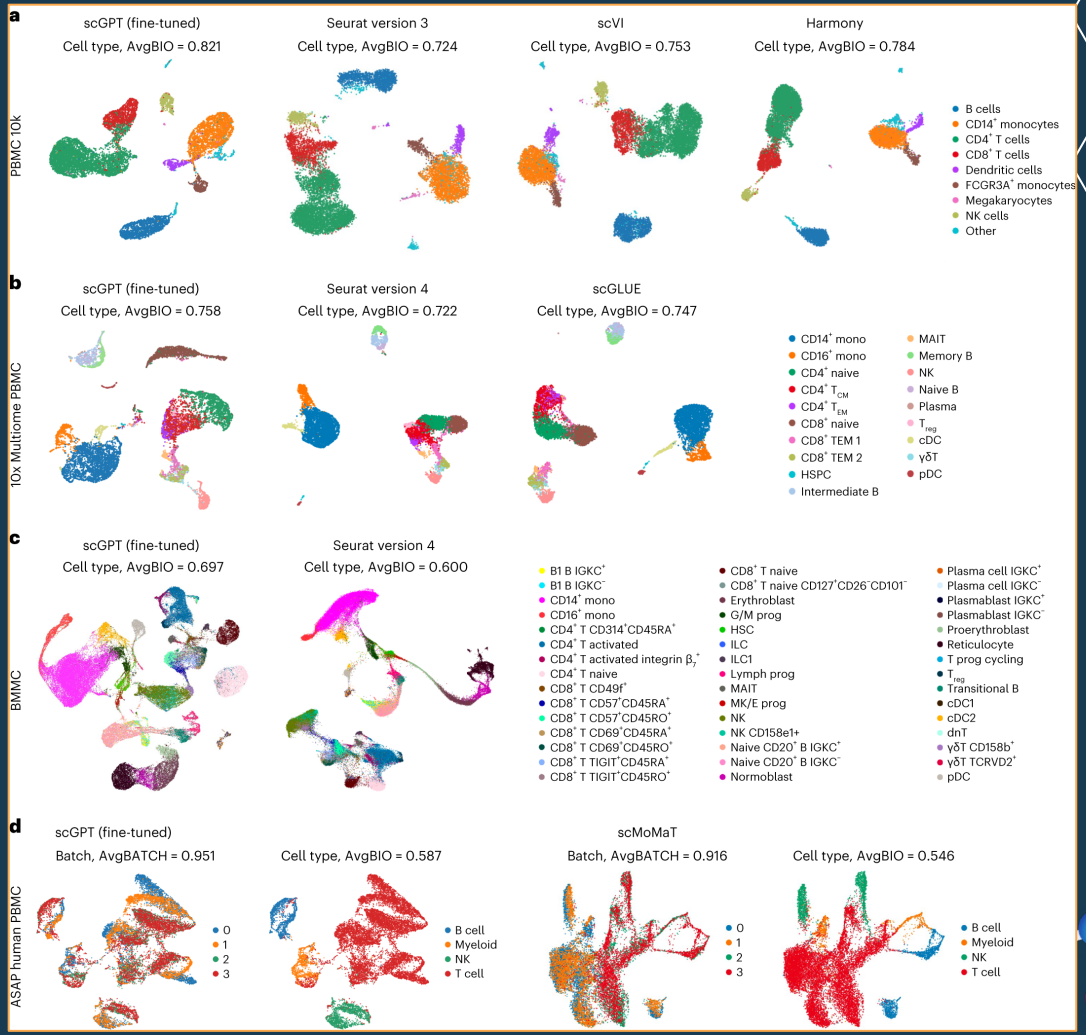

Conclusion

Limitations

- Pretraining does not mitigate batch effects.
- zero-shot performance could be constrained on datasets with technical variation
- Evaluating the model is also complex due to variation in data quality

Future Work

- pretrain on a larger-scale dataset with more diversity
- explore in-context instruction learning for single-cell data.

Summary

Conclusion & Questions

Model	DNABERT (BERT based)	UTR-LM	scGPT
Domain	DNA sequencing	5'UTR of mRNA	(scRNA-seq)
motive	Deciphering DNA sequences	Unified foundation model to study function of 5'UTR	Unified foundation model to study single-cell RNA functions
Method	• BERT architecture • Tokenization with k-mer (6) • Modify pre-training process • Fine-tuned on 3 specific tasks • Benchmark with current tools	• Transformer-based architecture • Masked nucleotide (MN) prediction • secondary structure (SS) • minimum free energy (MFE) • Fine-tuned on multiple downstream tasks	• Transformer-based architecture • Pretrained on a large corpus of single-cell RNA data • tokenization of gene expression profiles. • Multi-task learning approach
Results	• surpassing existing tools • Enhanced performance with limited data • No- separate training needed • Flexible learning of DNA in different situations	• outperforms the best-known baseline in each task. • Performance not limited by sequence length • Validated through wet-laboratory experiments • Zero shot generalization	• Pretrained model extrapolates to unseen datasets. • Outperform existing models • High accuracy in cell type annotation • strong scaling properties
limits	• Sequence Length Limitation • Dependence on k-mer Tokenization	• Computationally expensive	• Pretraining does not mitigate batch effects. • zero-shot performance could be constrained on datasets with technical variation

Questions

Paper	Question
Paper 1	
Paper 2	
Paper 3	

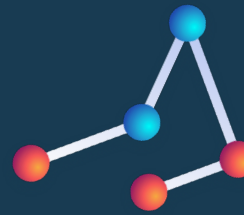

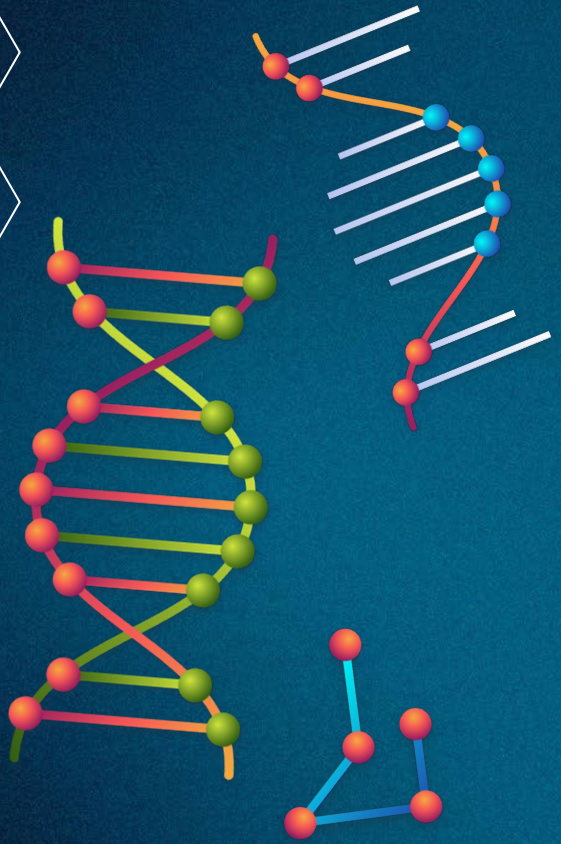

Thank you for listening

Any More Questions?